

Database Logical Design

Alessandro Bozzon

TI1506: Web and Database Technology

ti1506-ewi@tudelft.nl

Course overview [DB]

1. Introduction to Database Systems
2. The Relational Model
3. Introduction to SQL
4. Advanced SQL Topics
5. Introduction to NoSQL database systems
- 6. Conceptual Data Modelling with ER Diagrams**
- 7. Database Conceptual Design**
- 8. Database Logical Design**

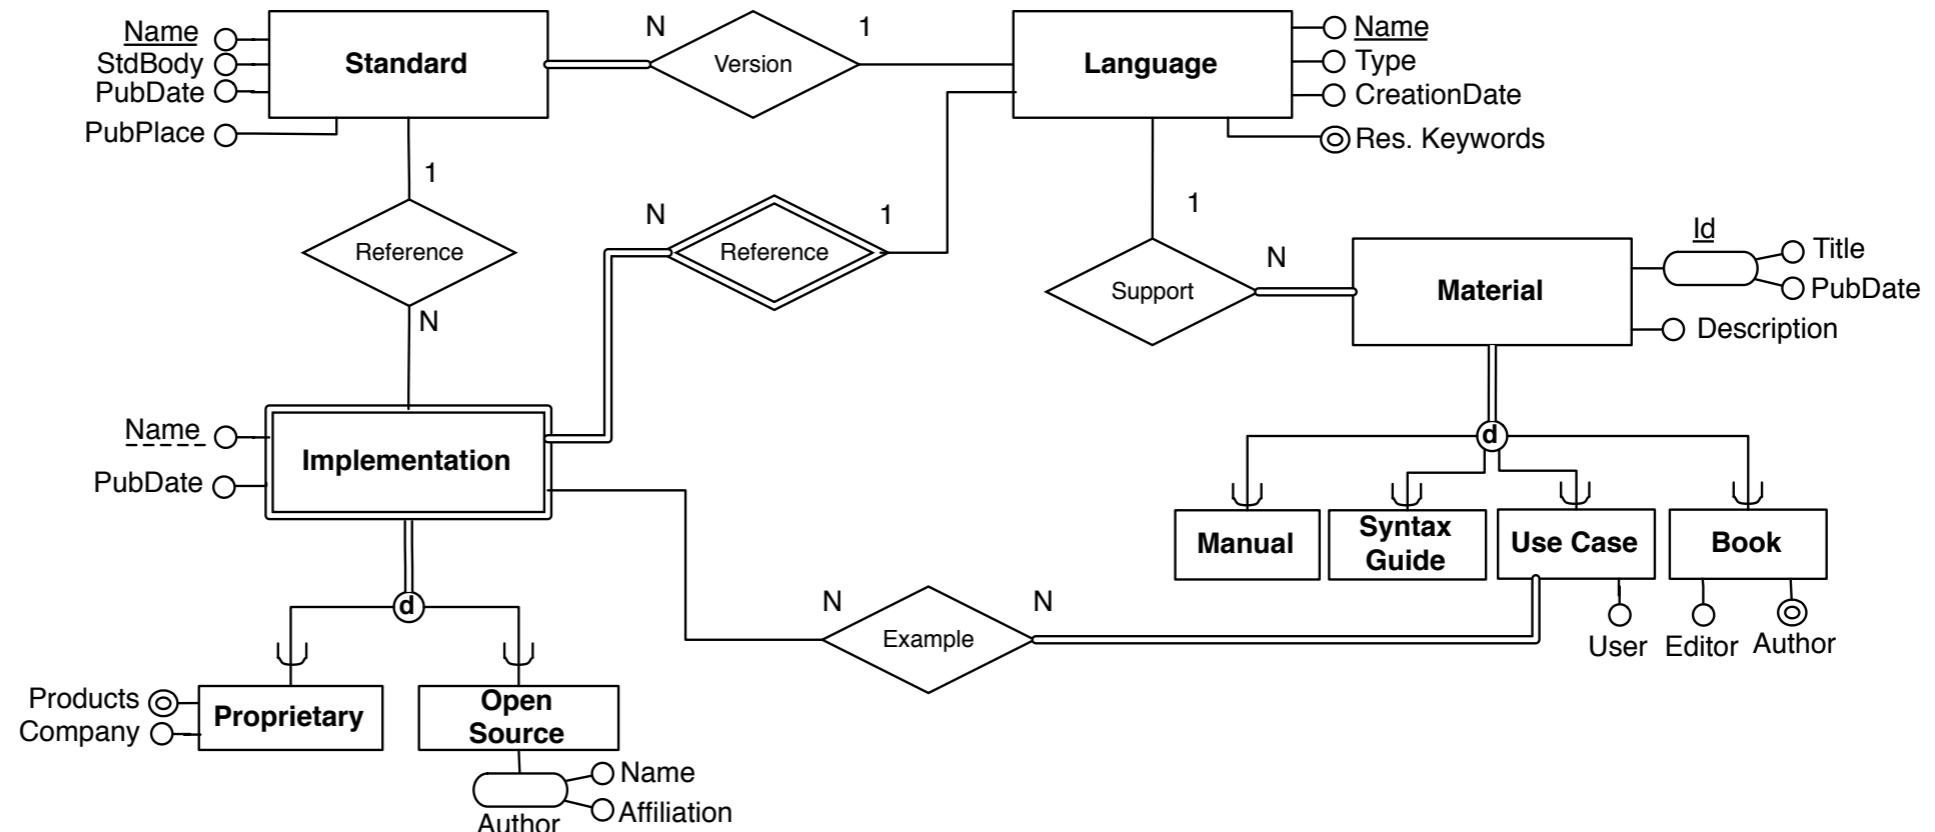
At the end of this lecture, you should be able to ...

- **Develop** and **refine** conceptual data models, using appropriate design notation
- **Derive** a logical design from a conceptual design, taking into account the requirements of both the adopted database technology, and the designed application
- **Develop** logical database schema, with principled design that enforce data integrity

ASQ Exercise (1 min)

A S Q

Consider the EER diagram in figure.
How many internal identifiers are defined?



1

2

2

3

3

4

4

5

EER Model

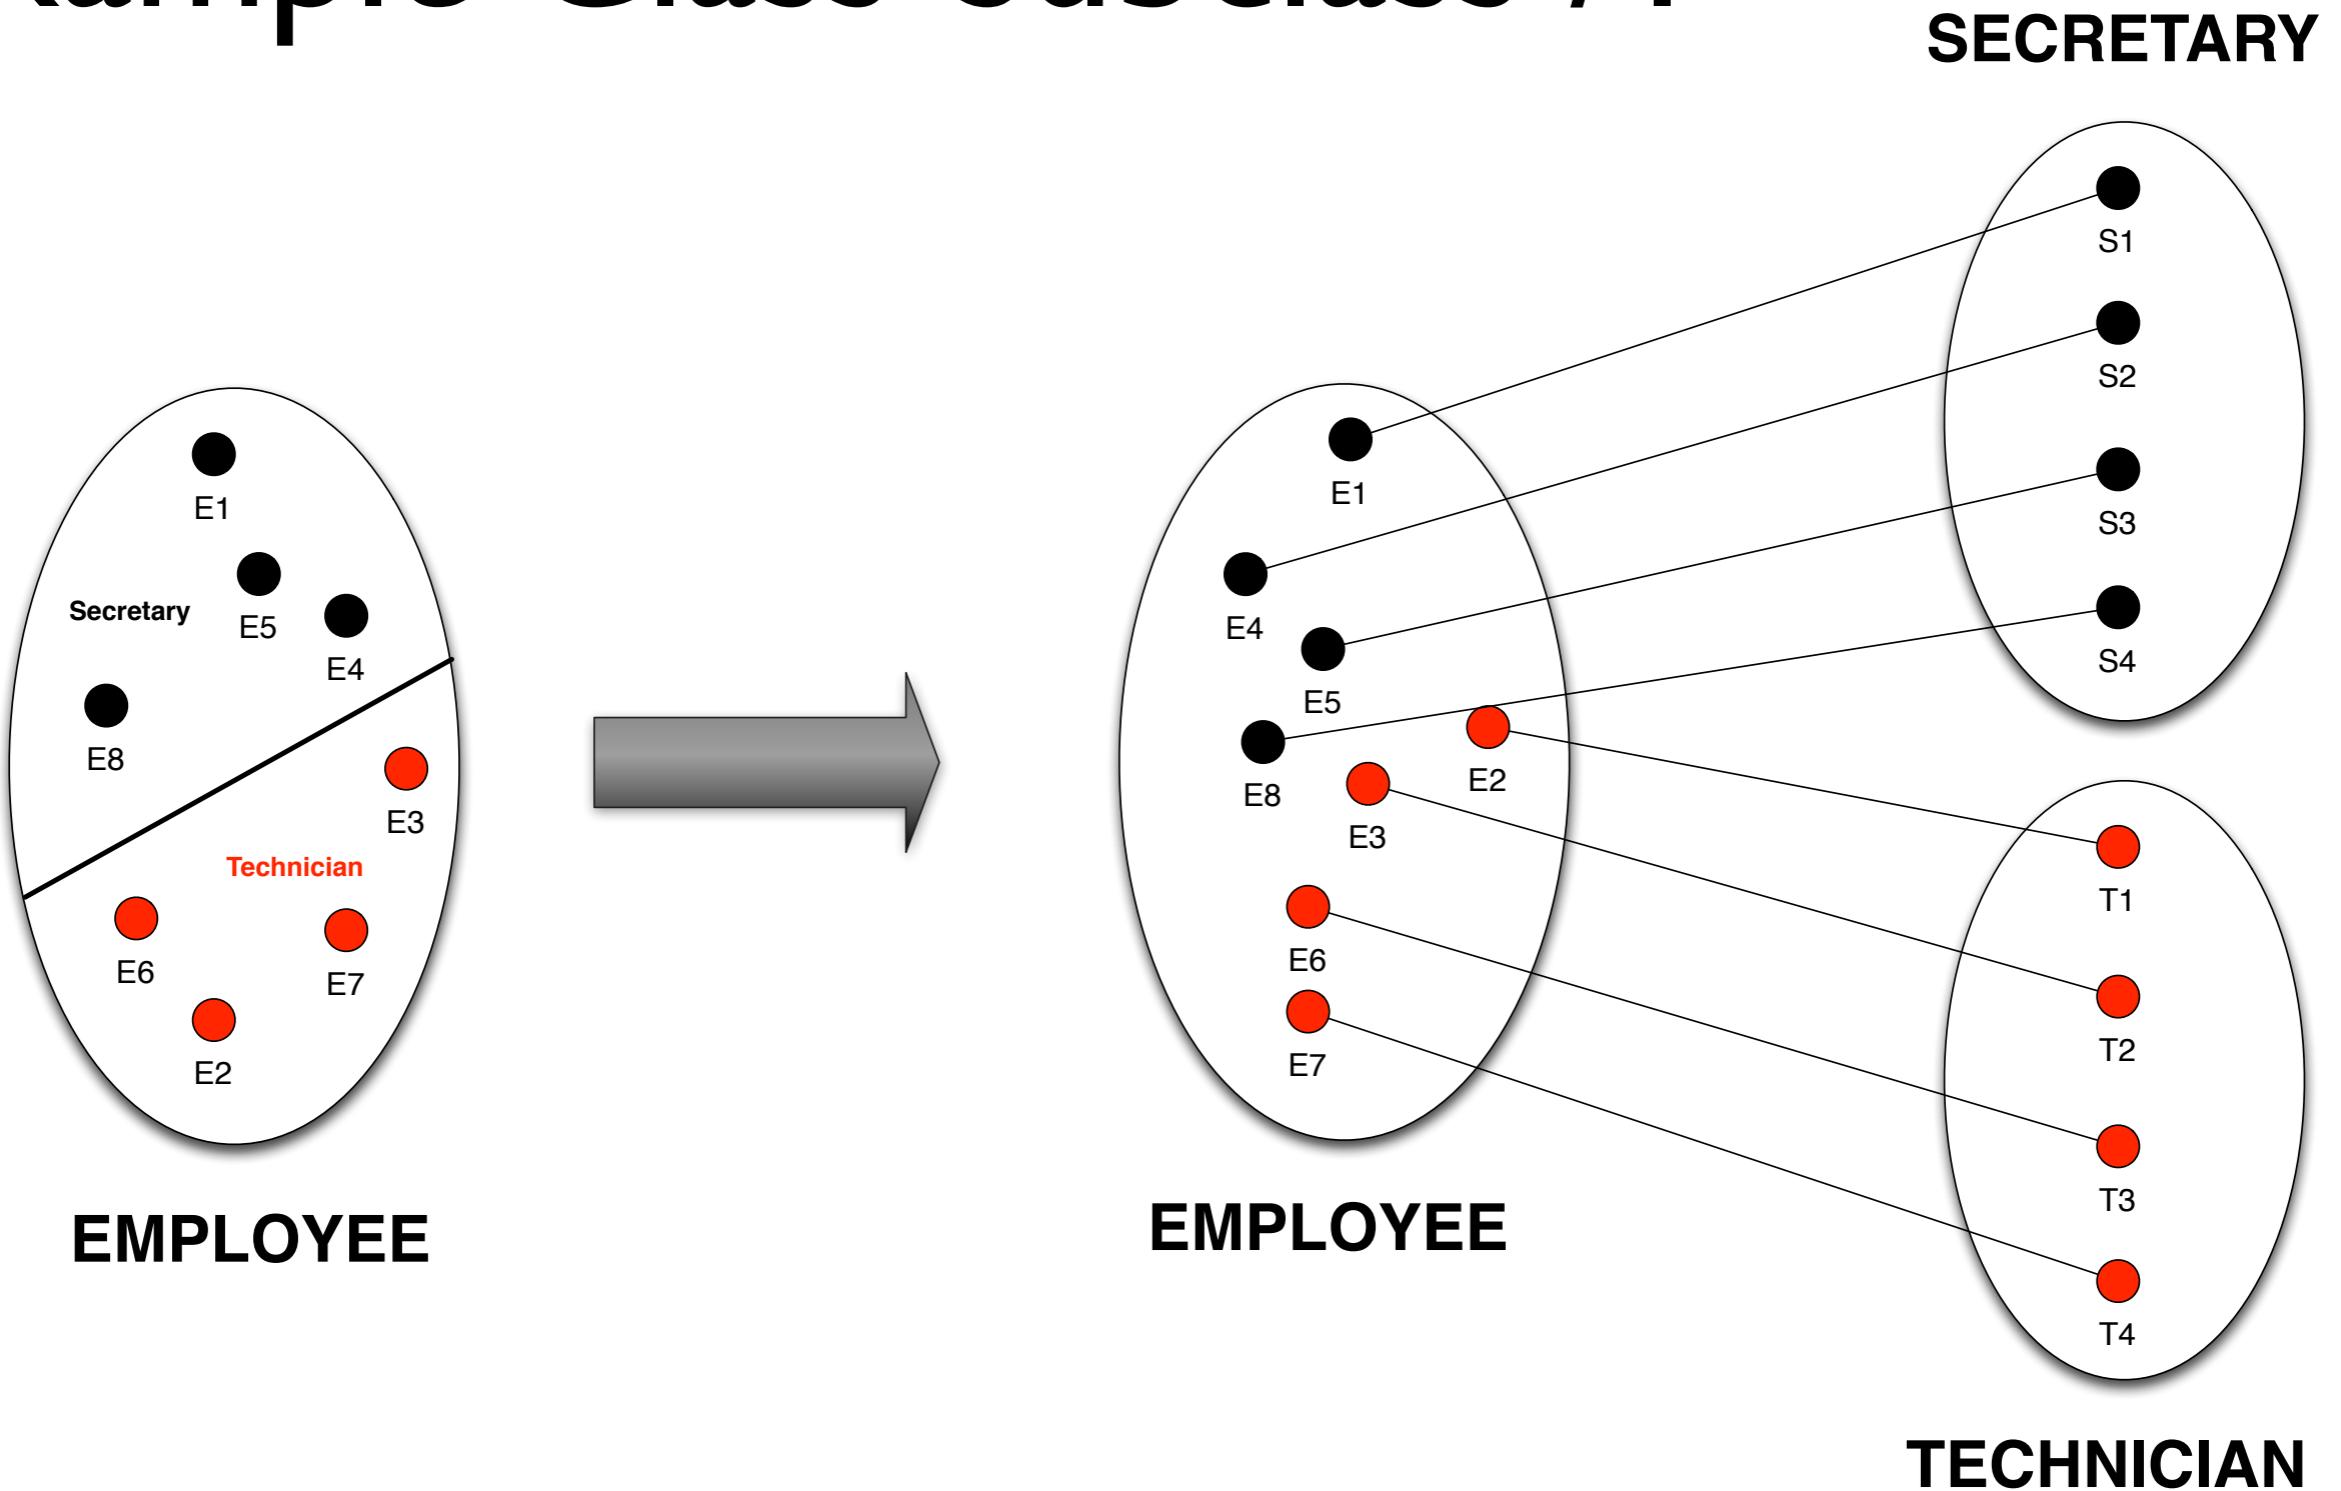
The Enhanced Entity-Relationship (EER) Model

- Created to design more accurate database schemas Reflect the data properties and constraints more precisely
- More complex requirements than traditional applications
- EER model includes all modelling concepts of the ER model.
- In addition, EER includes:
 - Subclasses and superclasses
 - Specialisation and generalisation
 - Category or union type (not covered)
 - Attribute and relationship inheritance (not covered)

Class / Subclass (Type / Subtype)

- An entity type is used to represent a type of instance (attribute and relationships)
 - EMPLOYEE (name, BSN, BirthDate, etc)
- But also the *collection of entities (entity set) of that type* that exist in the database
 - E.g. entities that represent specific types of employees
 - SECRETARY, TECHNICIANS, MANAGERSs, etc.
- The relationship that exists between a type of entity and its entity set is called a **class/ subclass** relationship
 - EMPLOYEE is the **super type** or **super class**
 - SECRETARY, TECHNICIAN etc. are **subtypes** or **subclasses** of EMPLOYEE

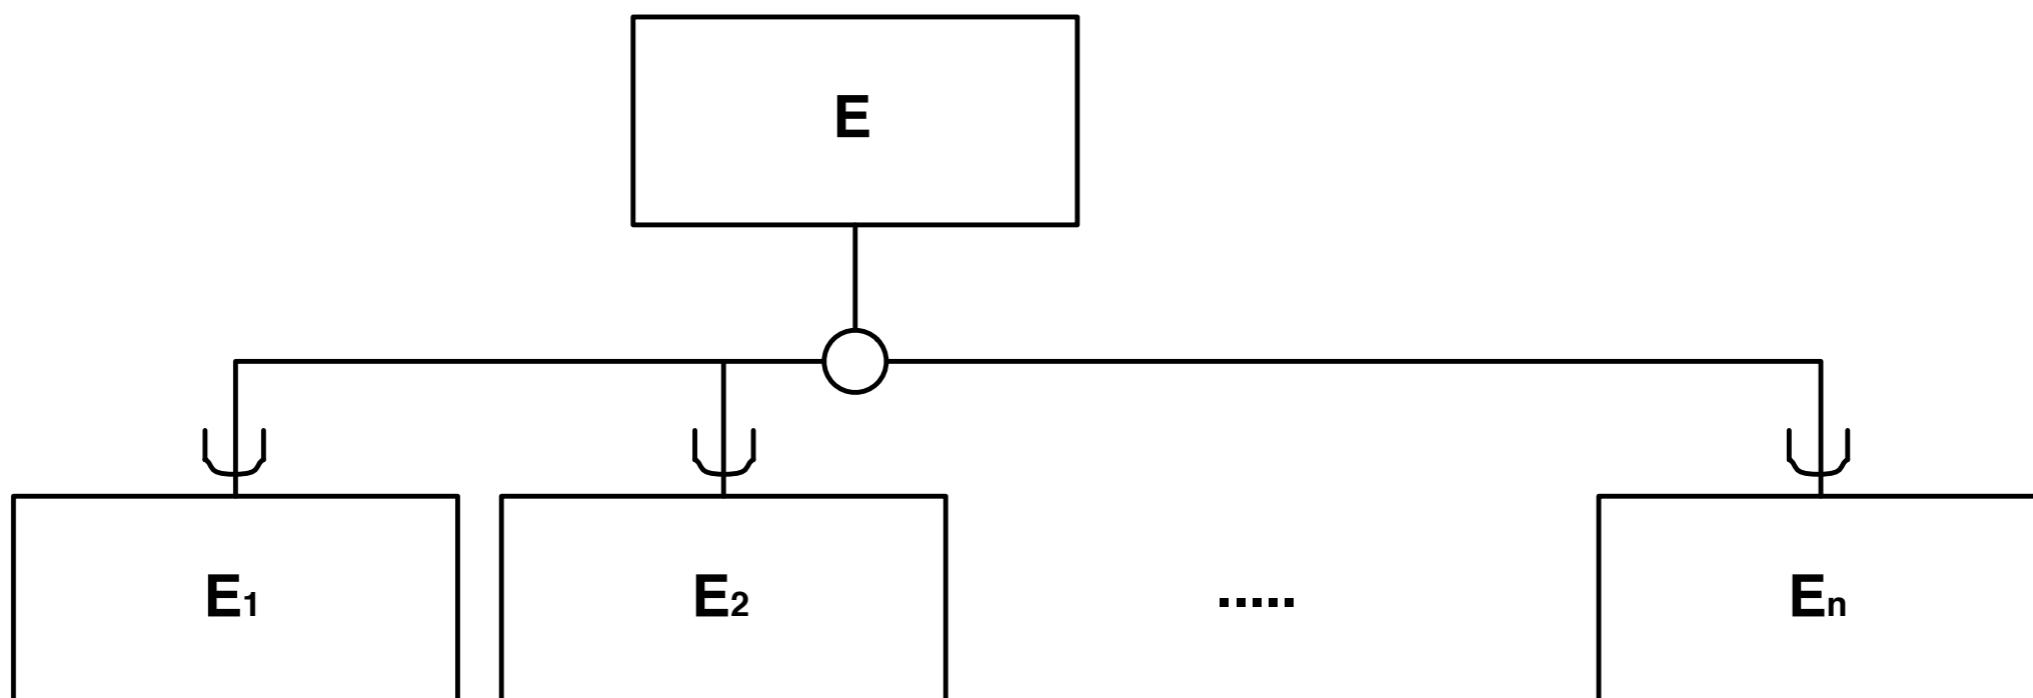
Example Class-Subclass / I



NOT a 1:1 relationship type, but a relationship
between OCCURRENCES

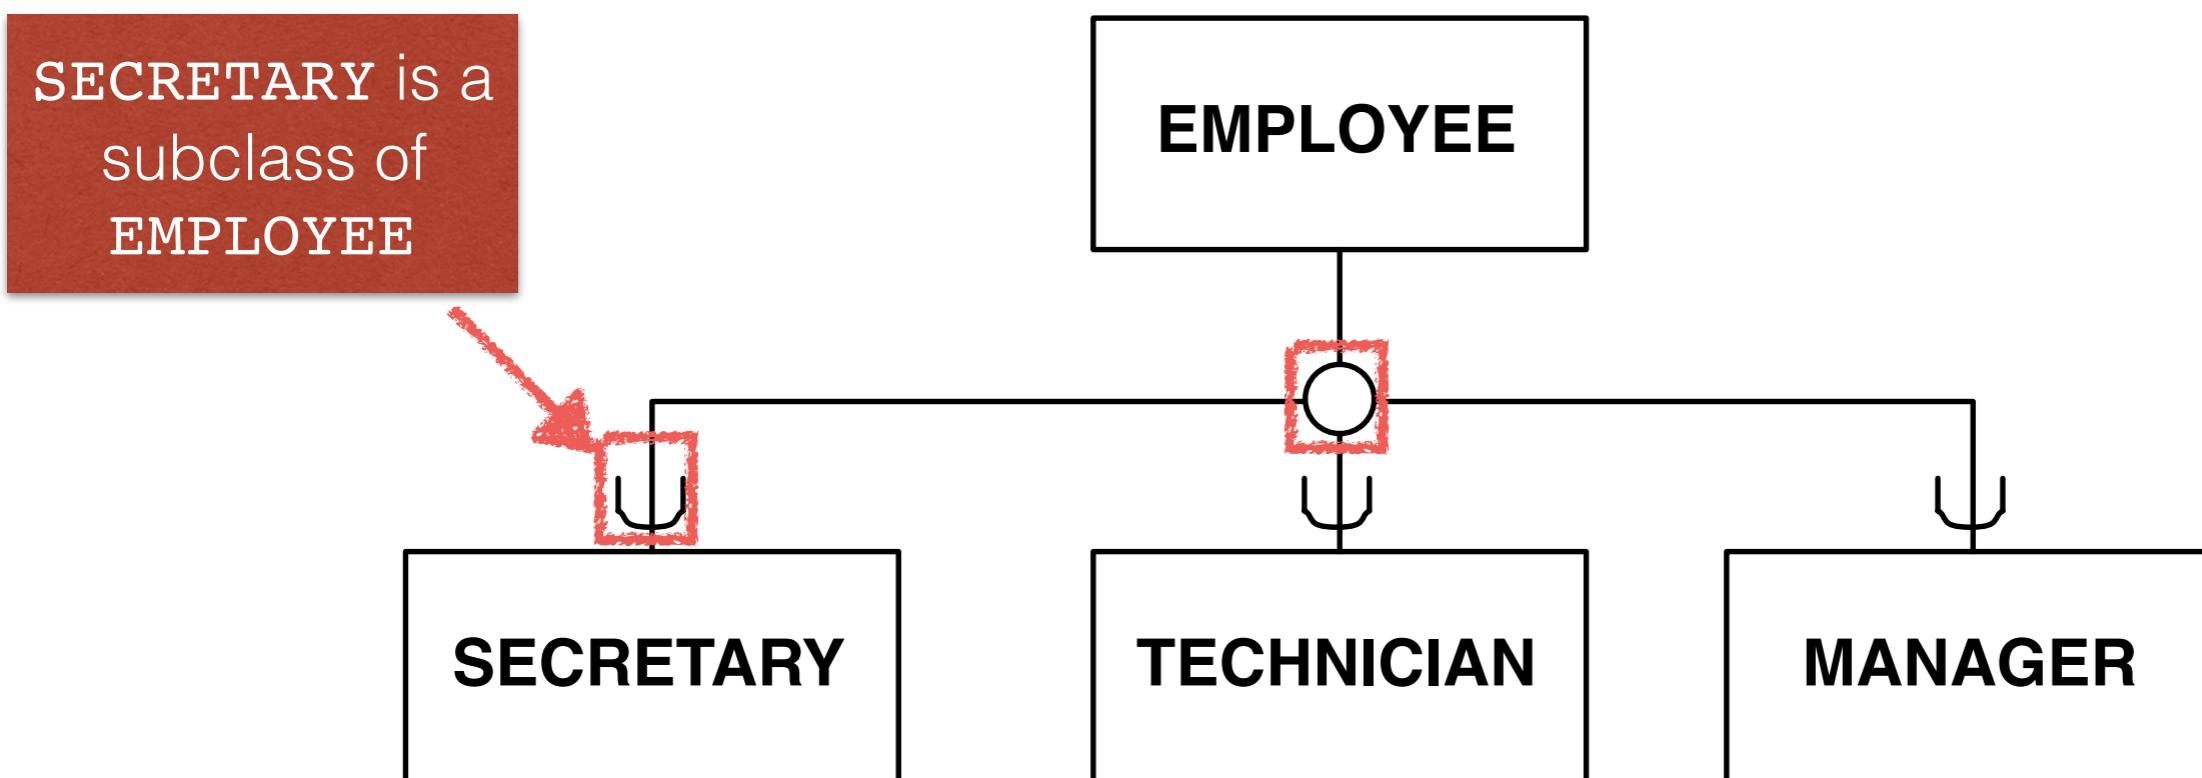
Class/ Subclass relationship

- Also called **generalisation** or **specialisation** relationship
- Represent logical links between an entity E , known as parent entity, and one or more entities E_1, \dots, E_n called child entities, of which E is more general, in the sense that it comprises them as a particular case.
- In this situation we say that E is a **generalisation** of E_1, \dots, E_n and that the entities E_1, \dots, E_n are specialisation of the E entity



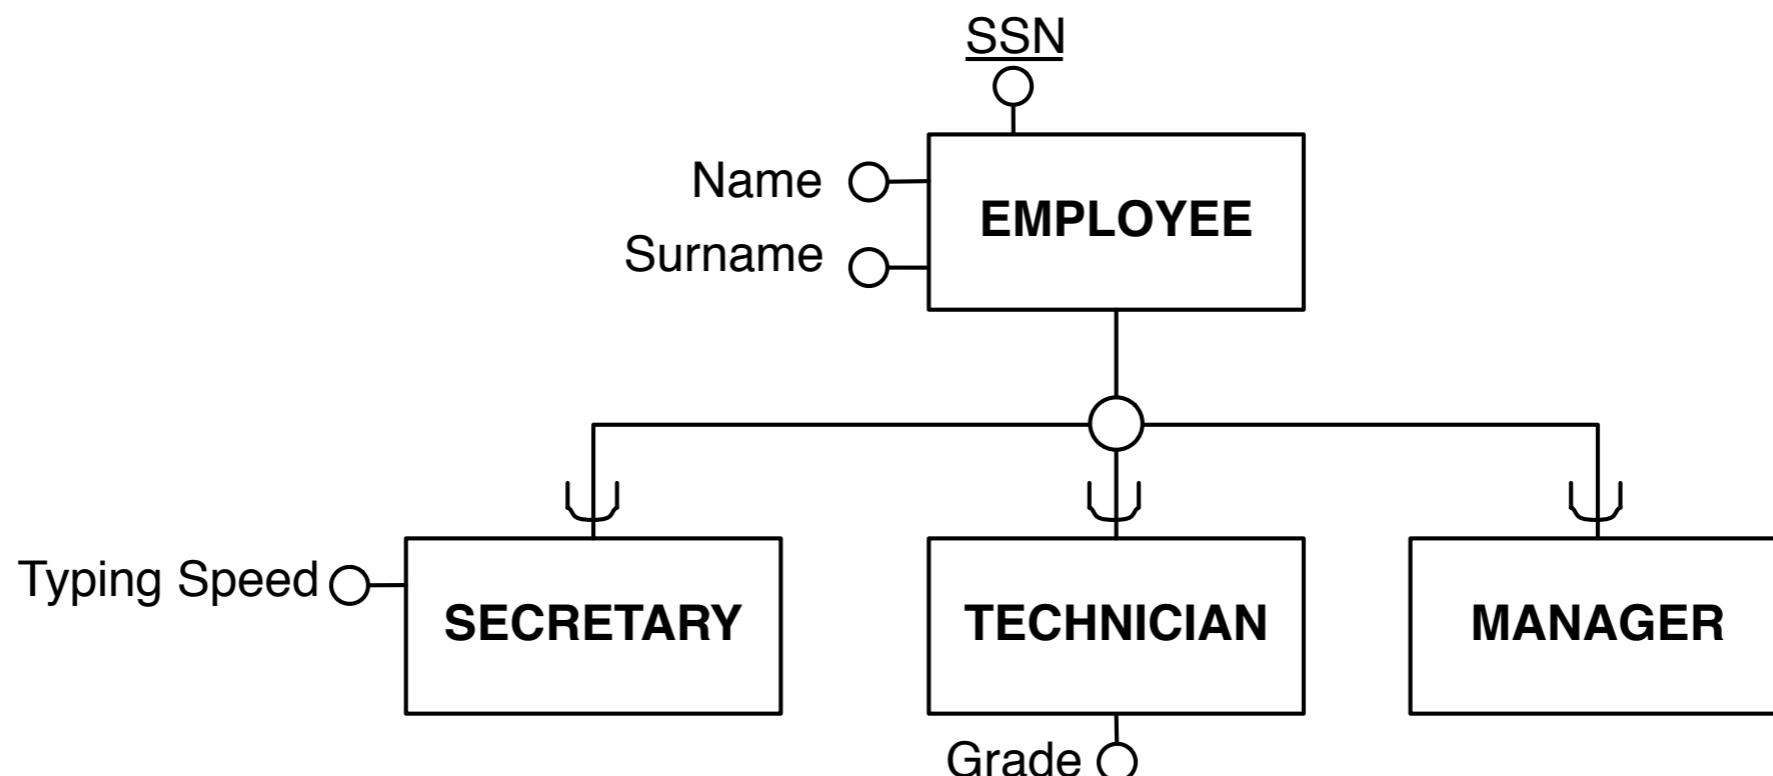
Example Class-Subclass /2

- The subset symbol indicates the *direction* of the **specialisation**
- The circle defines the entities involved in the specialisation



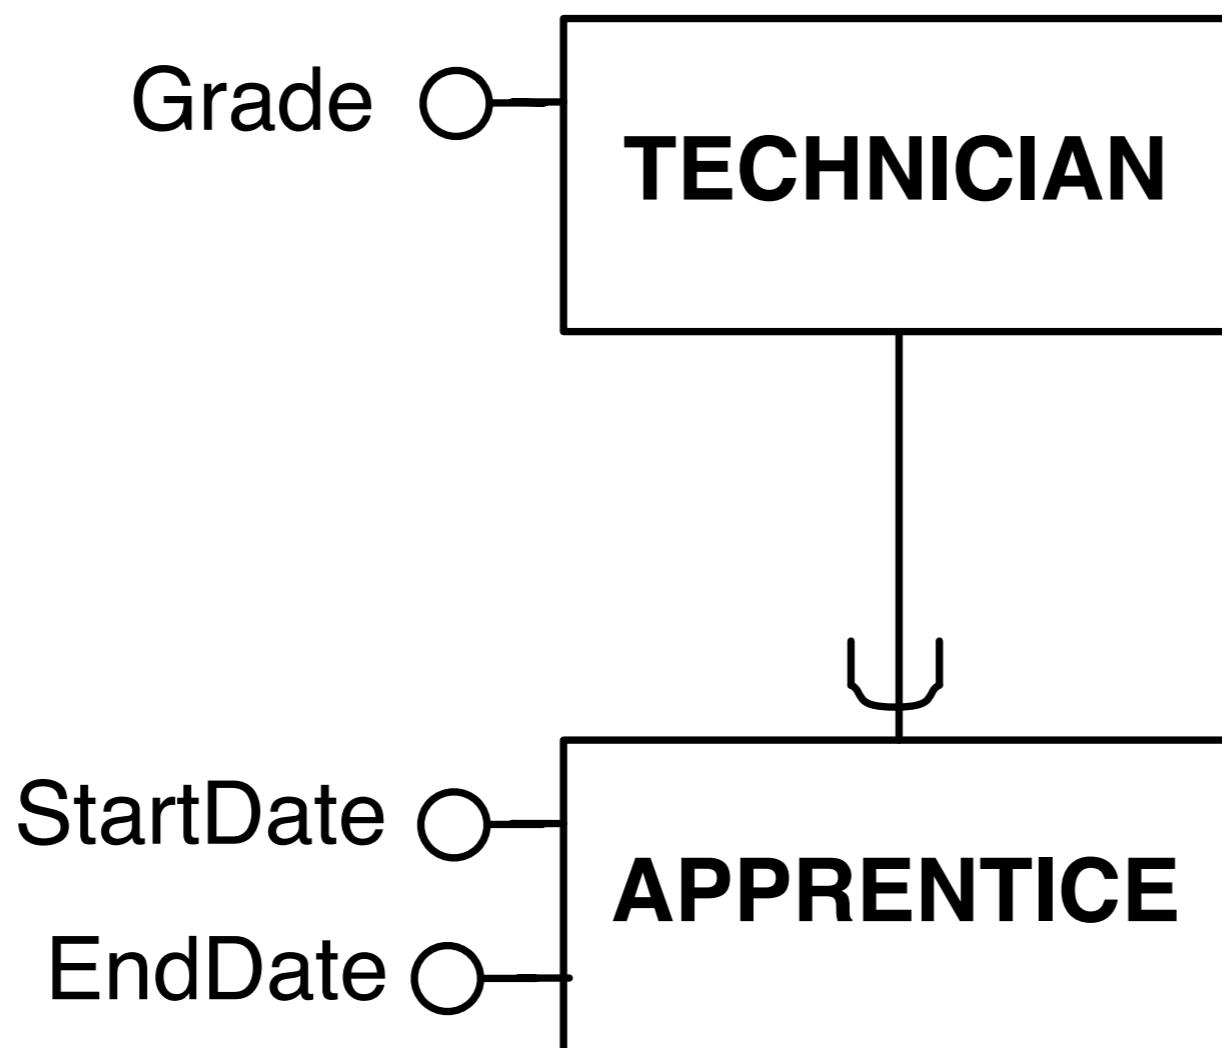
Properties of Specialisation

- Every occurrence of a child entity is also an occurrence of the parent entity
- **Inheritance**
 - Every property of the parent entity (attributes, identifiers, relationships and other generalisations) is also a property of a child entity.
 - Specialisation can define **specific attributes** and/or specific **relationship types**



Subset

- Special class of specialisation with a single child entity

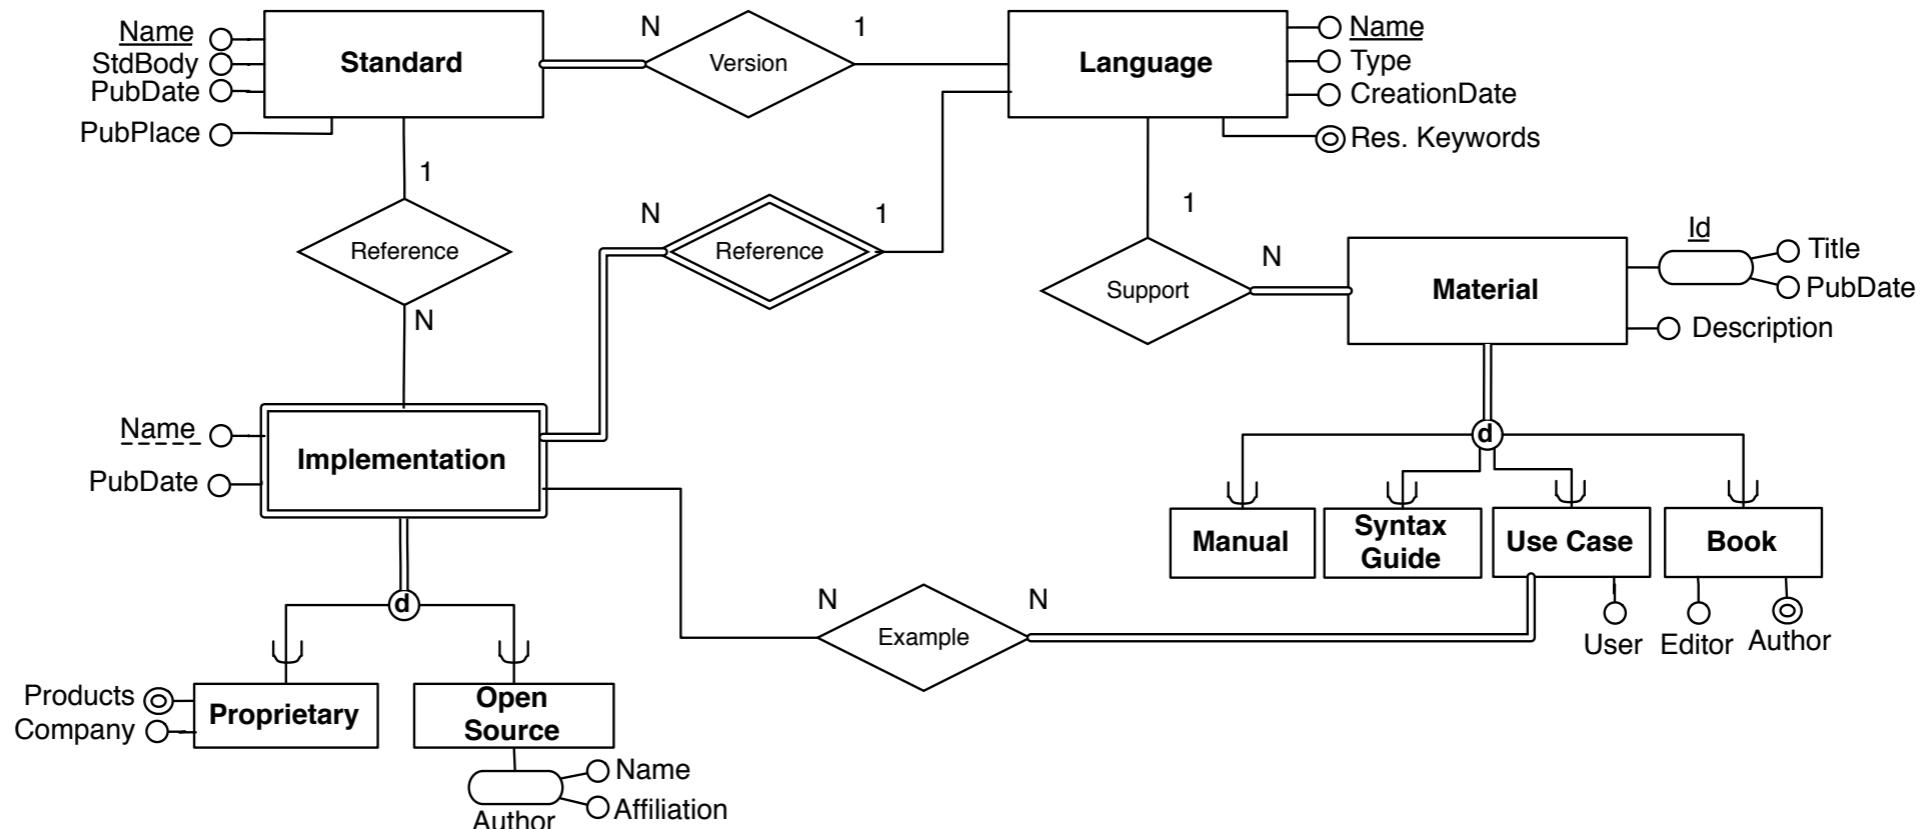


ASQ Exercise (1 min)



Consider the EER diagram in figure.

Which of the following statements is correct?



1

A *Use Case* is an example of *Implementation*

2

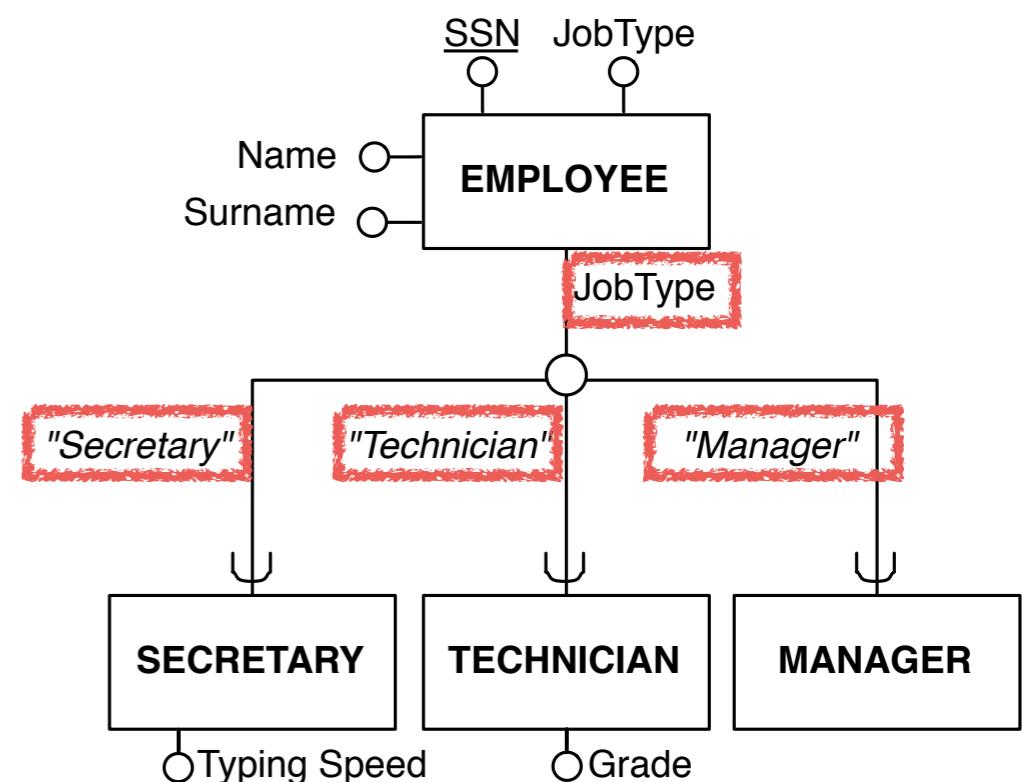
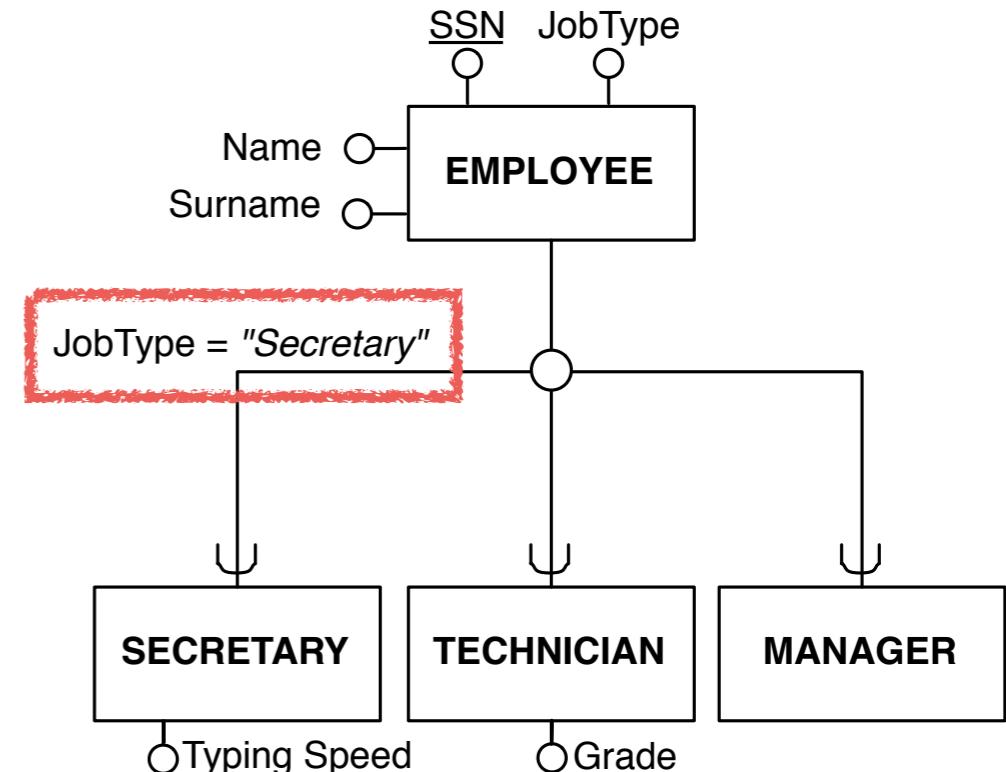
A *Use Case* is a type of *Material*

3

An *Implementation* generalises into *Proprietary Implementation*

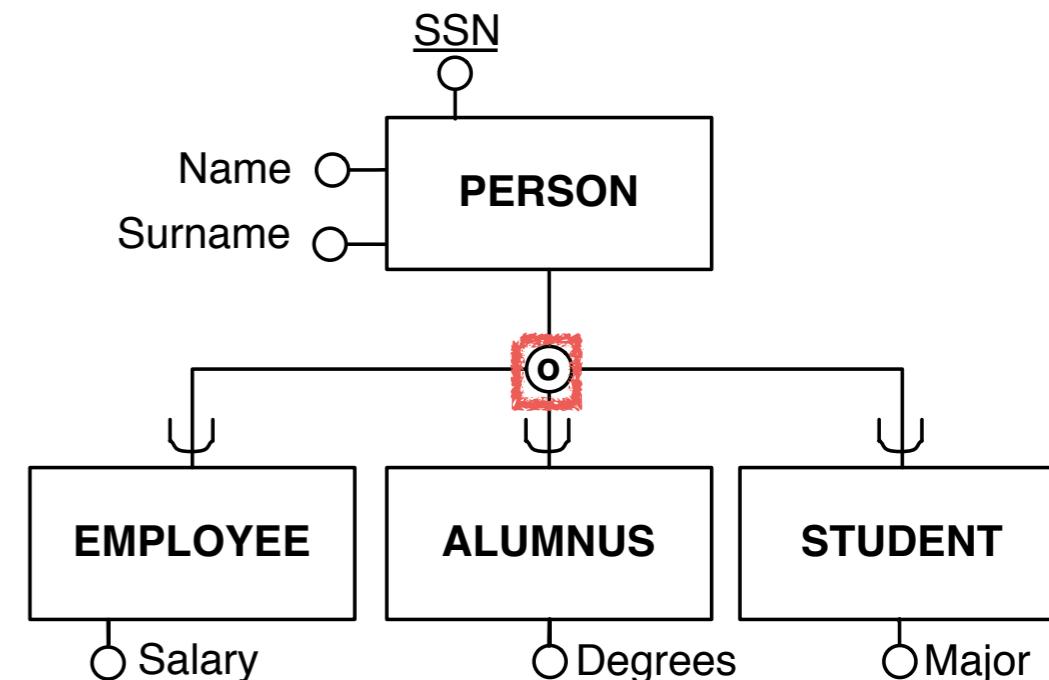
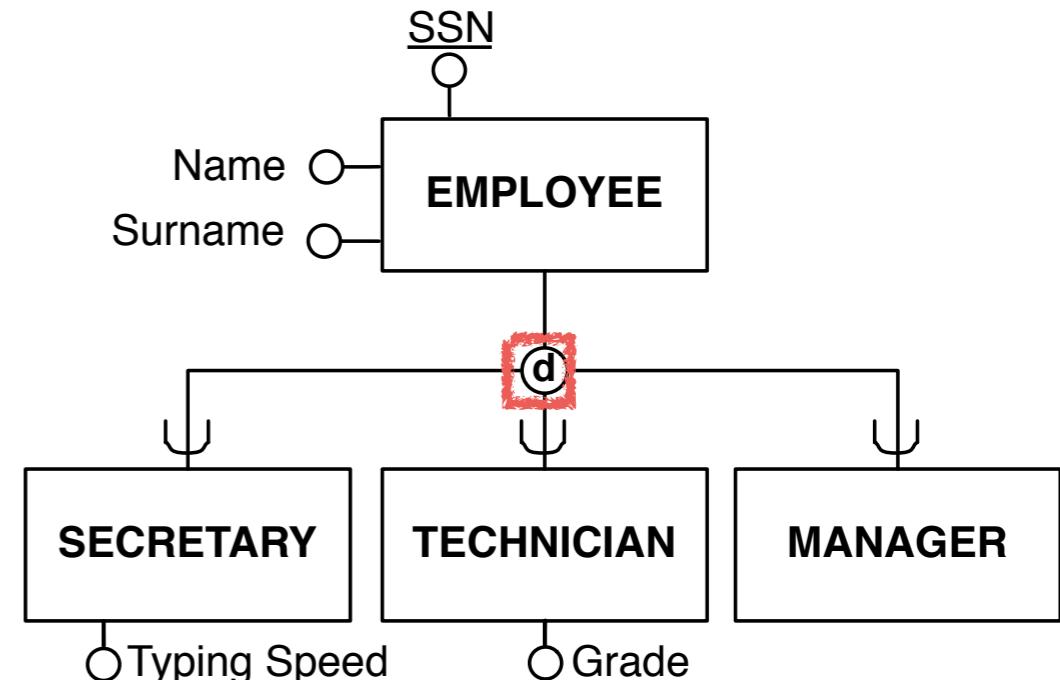
Constraints: Membership

- In some specialisation we can determine exactly the entities that will become members of each subclass
 - *Predicate-defined* (or condition-defined) subclasses
- All the subclasses in a specialisation have their membership condition on the same attribute
 - *Attribute-defined specialisation* (**Defining** attribute of the specialisation)
- No specific condition
 - **User-defined** subclass
 - specialisation defined by the database user, instance by instance



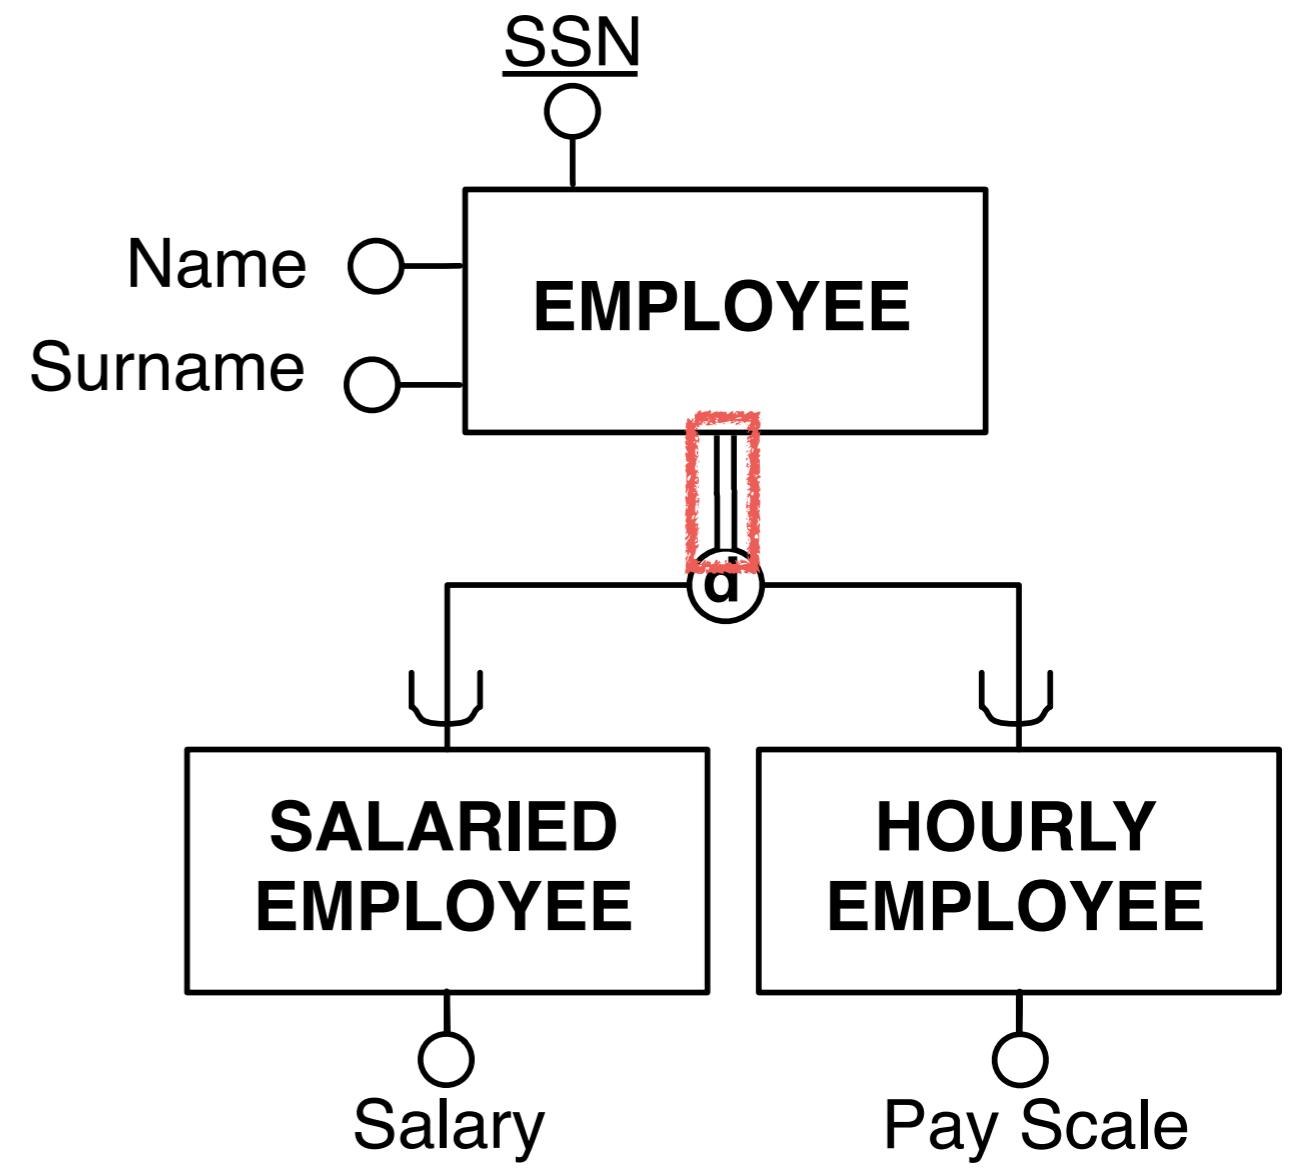
Constraints: Disjointness

- An instance can be a member of **at most** one of the subclasses of the specialisation
 - defined by a **d** in the graphical notations
 - An attribute-defined specialisation implies the **disjointness** constraint
- If the subclasses are not constrained to be disjoint, then their set of instances might be overlapping
 - defined by a **o** in the graphical notations



Constraints: Totalness

- A **Total** specialisation constraint specified that every instance of the superclass must be a member of **at least one** subclass in the specialisation
 - Defined by a **double line** in the graphical notations
- A **Partial** specialisation allows an instance not to belong to any of the subclasses



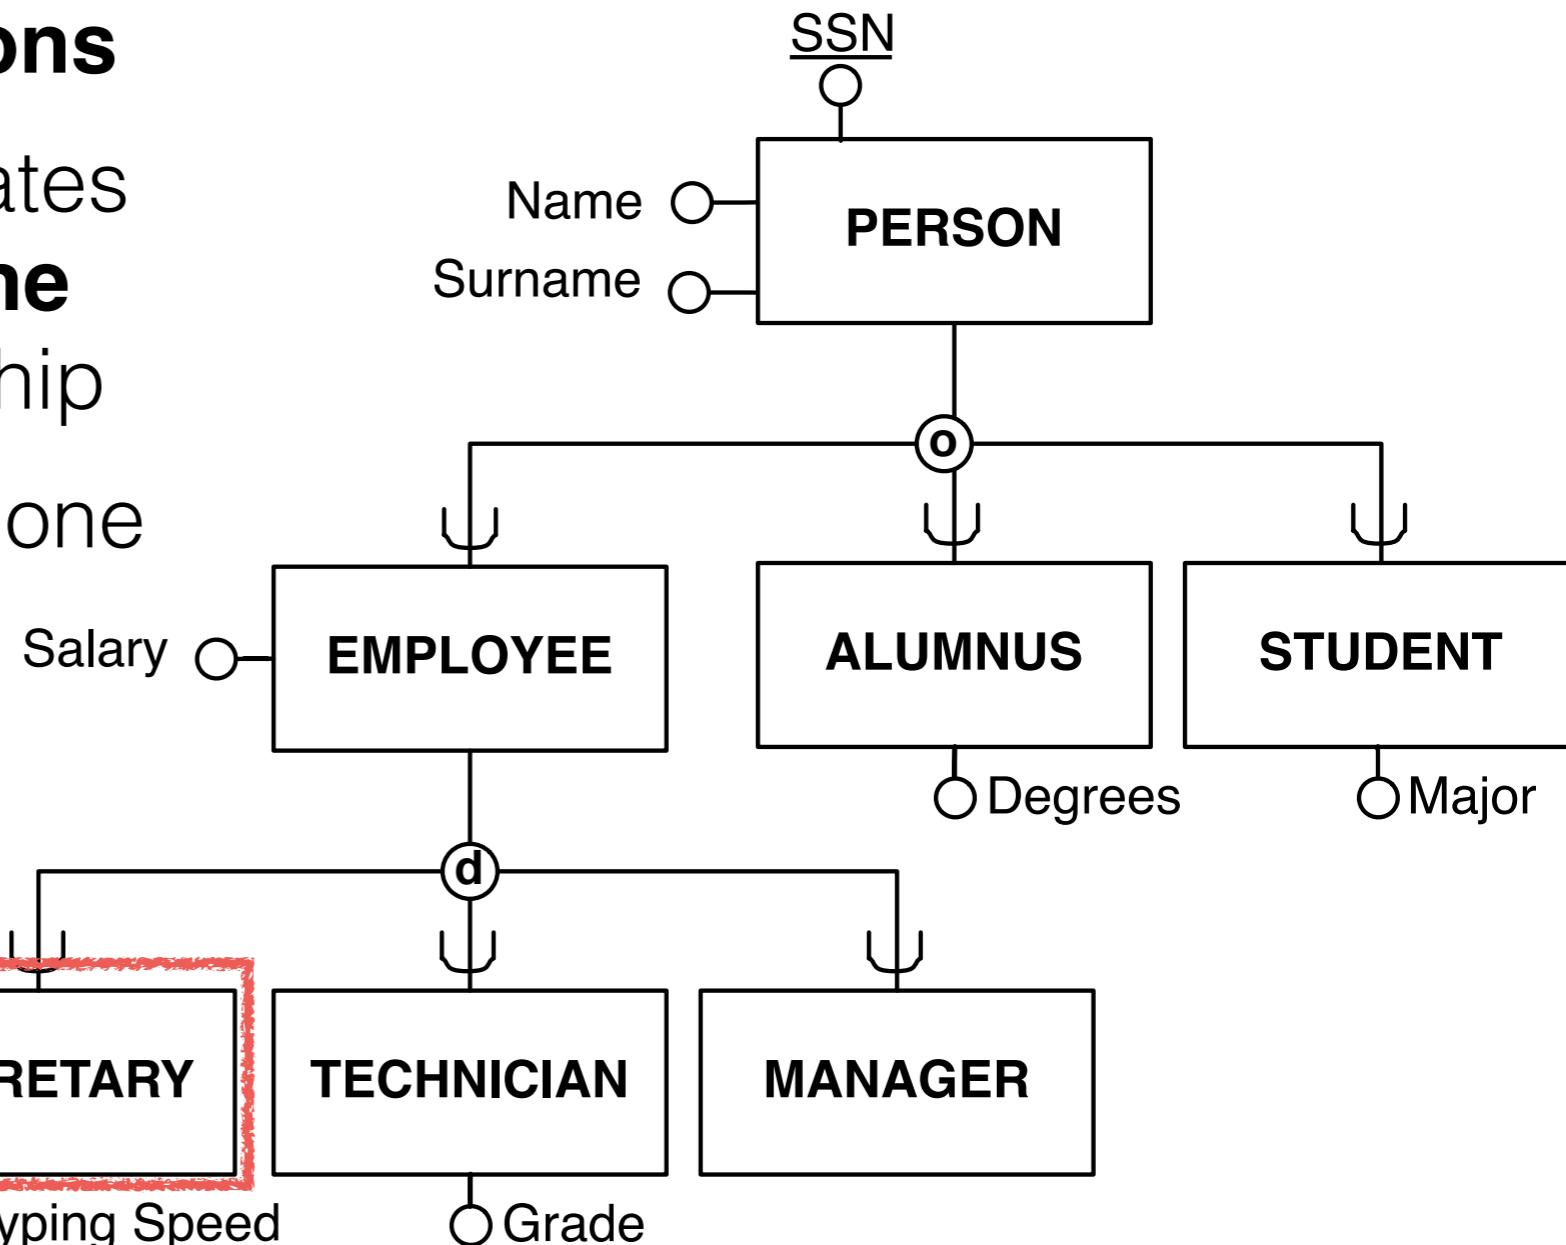
On Disjointness and Totalness

- **Disjointness** and **Totalness** are independent constraints
 - Disjoint, Total
 - Disjoint, Partial
 - Overlapping, Total
 - Overlapping, Partial
- Constraints define rules for instance insertion and deletion. For instance:
 - Deleting an entity from a superclass: deletion from all the subclasses (general)
 - Inserting in a superclass: **mandatory insertion** in a sub-class for predicate-defined or attribute-defined subclasses where the predicate evaluates to **TRUE**
 - Inserting in a superclass of a total specialisation: insertion in **at least one** of the subclasses is mandatory
 - ... (other rules left as exercise)

Specialisation Hierarchies

- A subclass might have further subclasses specified on it
- **Hierarchy of specialisations**
 - every subclass participates as a subclass in **only one** class/subclass relationship
 - Each subclass has only one parent

A PERSON and an EMPLOYEE

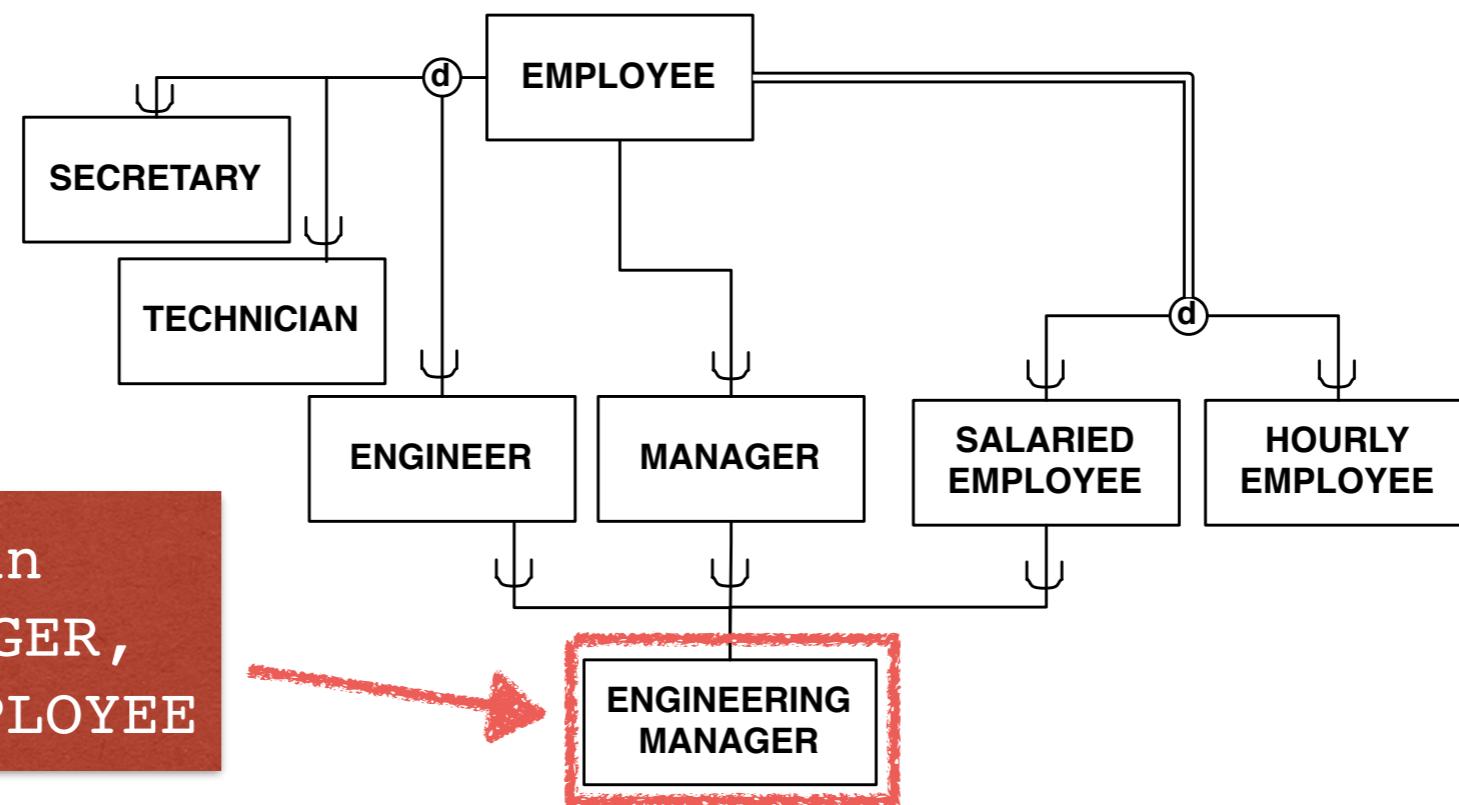


Specialisation Lattices

- A subclass might have further subclasses specified on it

- **Lattice of specialisations**

- a subclass can participate as a subclass in more than one class/subclass relationship
- A subclass with more than one superclass is called **shared subclass**
- **Multiple inheritance** of attributes from entities in the lattice but also from multiple subclasses
- If attribute (or relationship) is inherited more than once via different paths in lattice, then it is included only once in the same shared subclass



An **EMPLOYEE**, an
ENGINEER, a **MANAGER**,
and a **SALARIED EMPLOYEE**

ASQ Exercise (1 min)

A S Q

Consider the EER diagram in figure.

Which of the following statements is correct?

1

An occurrence of the *Monument* entity cannot be an occurrence of the *Place* entity

2

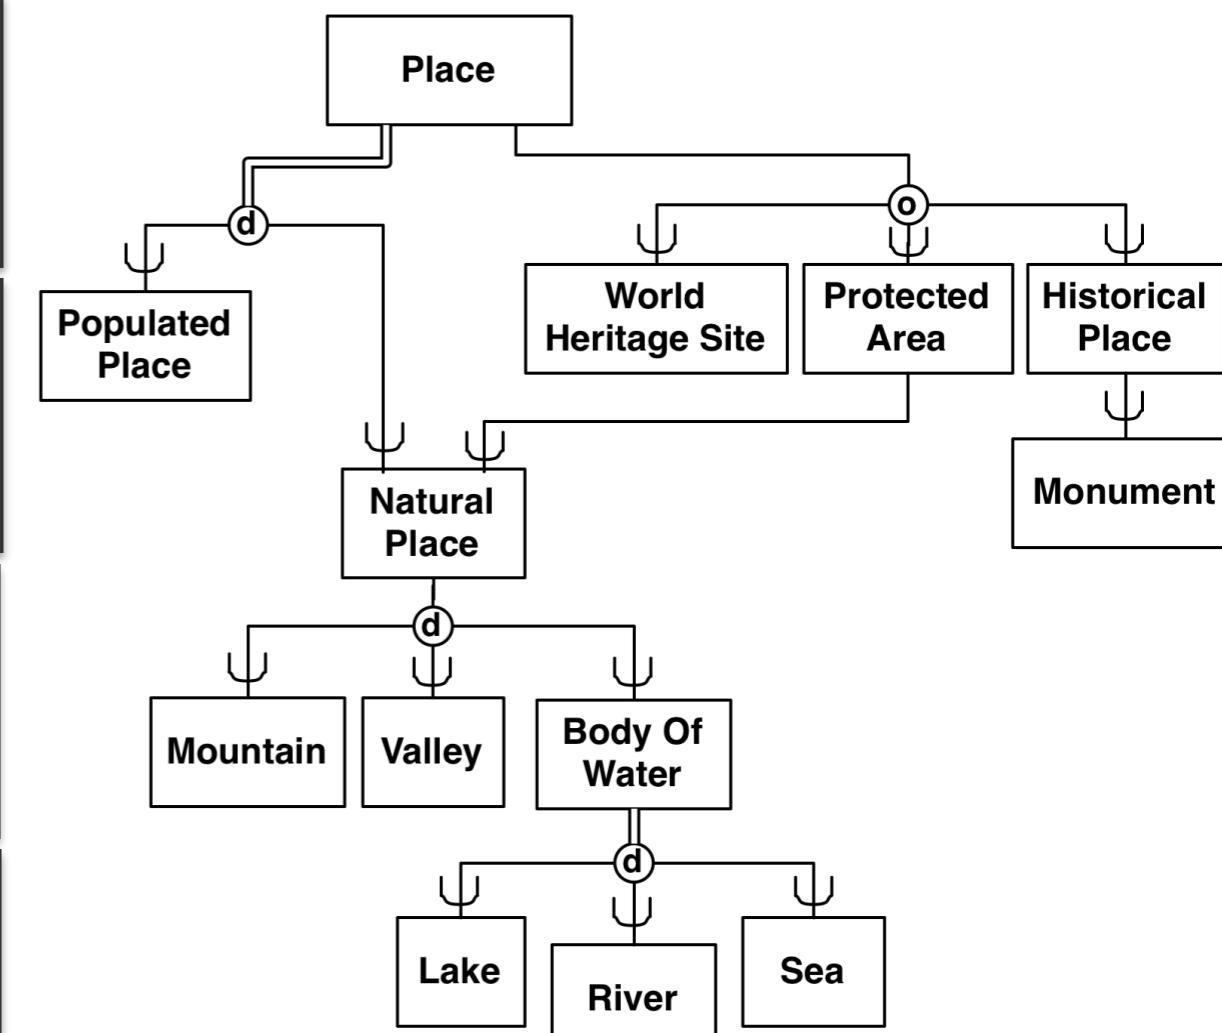
An occurrence of the *Valley* entity can be an occurrence of the *World Heritage Site* entity

3

An occurrence of the *Valley* entity must be an occurrence of the *Protected Area* entity

4

An occurrence of the *River* entity cannot be an occurrence of the *Populated Place* entity



Database Logical Design

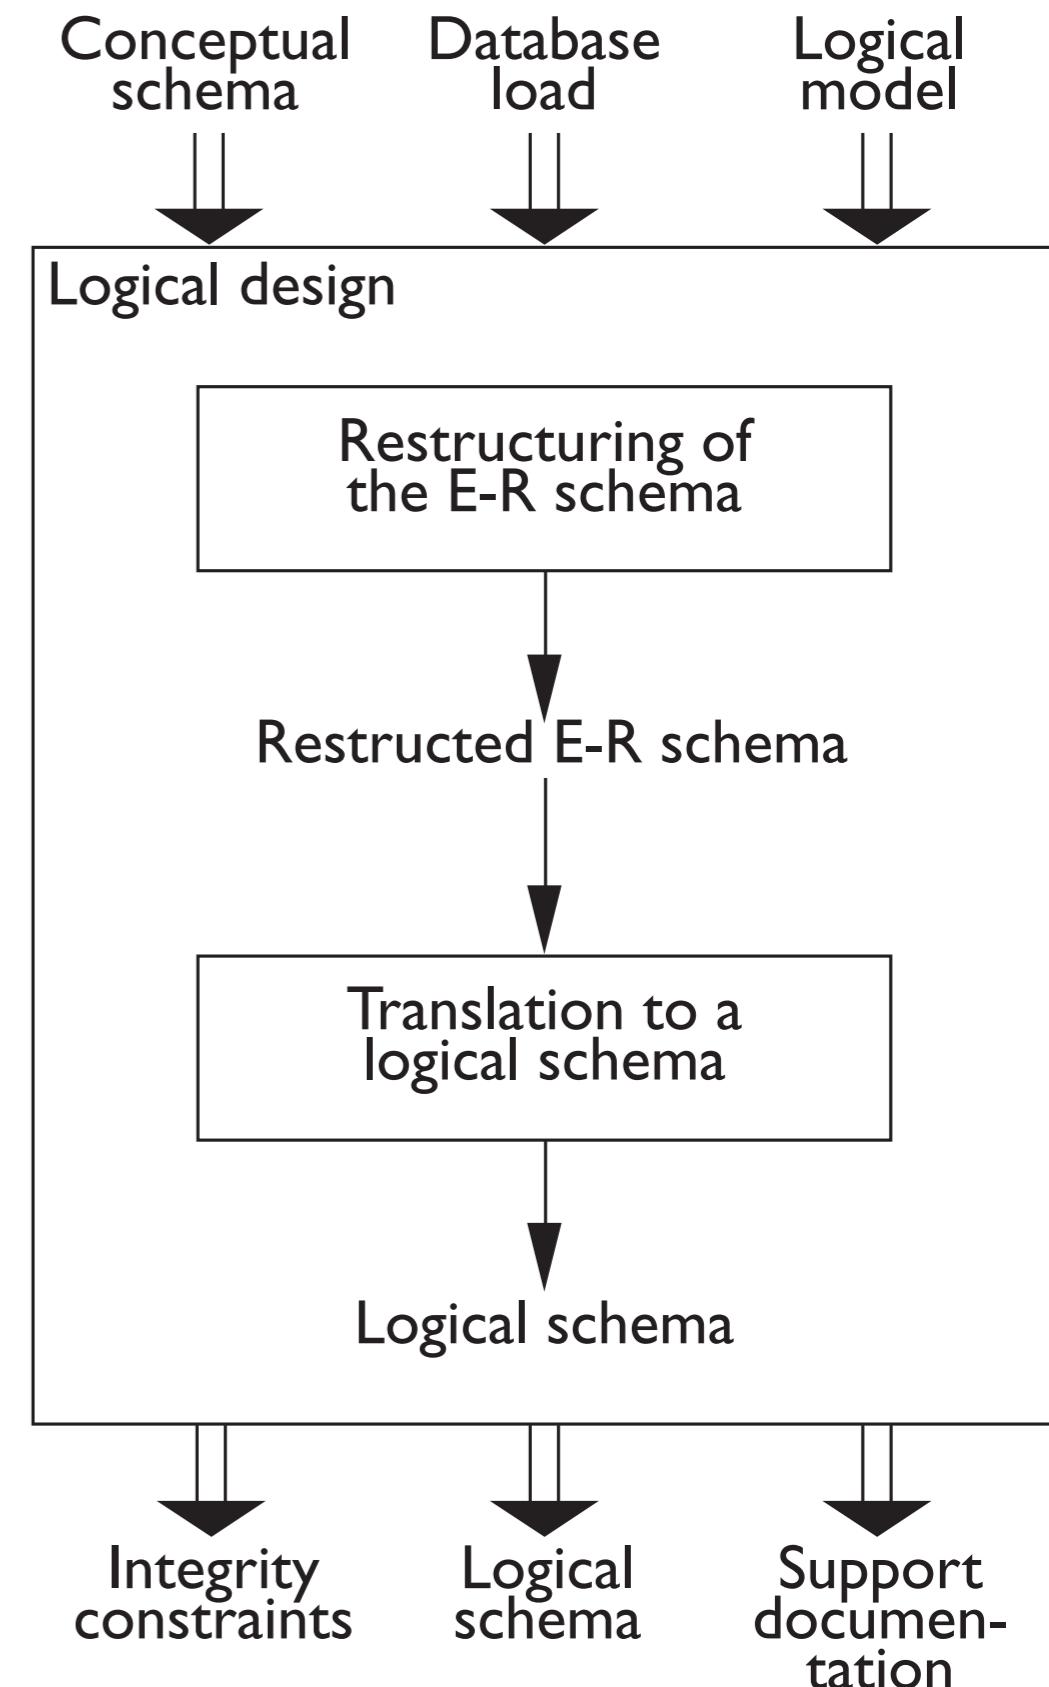
Logical Design

GOAL: construct a relational schema that represents all of the information described in an Entity-Relationship schema

- **correctly**
- **efficiently**
- Not just a simple model translation:
 - **not all** the constructs of the Entity-Relationship model can be translated naturally into the relational model
 - the ER schema **must be restructured** to make the execution of the projected operations as **efficient** as possible

Logical Design Steps

1. **Restructure** the Entity-Relationship schema
 - schema optimisation
 - simplification of following step
2. **Translate** into the logical model
 - based on the features of the logical model
 - in our case, the relational model



Restructuring the ER Schema

Restructuring of an E-R schema

- The restructured ER schema takes into account some implementation aspects (and constraints) coming from the logical model of choice
 - **Not a conceptual model anymore**
- **Goals**
 - Remove constructs for which there is no direct mapping in the relational model
 - Transformations aimed at increasing data access operations

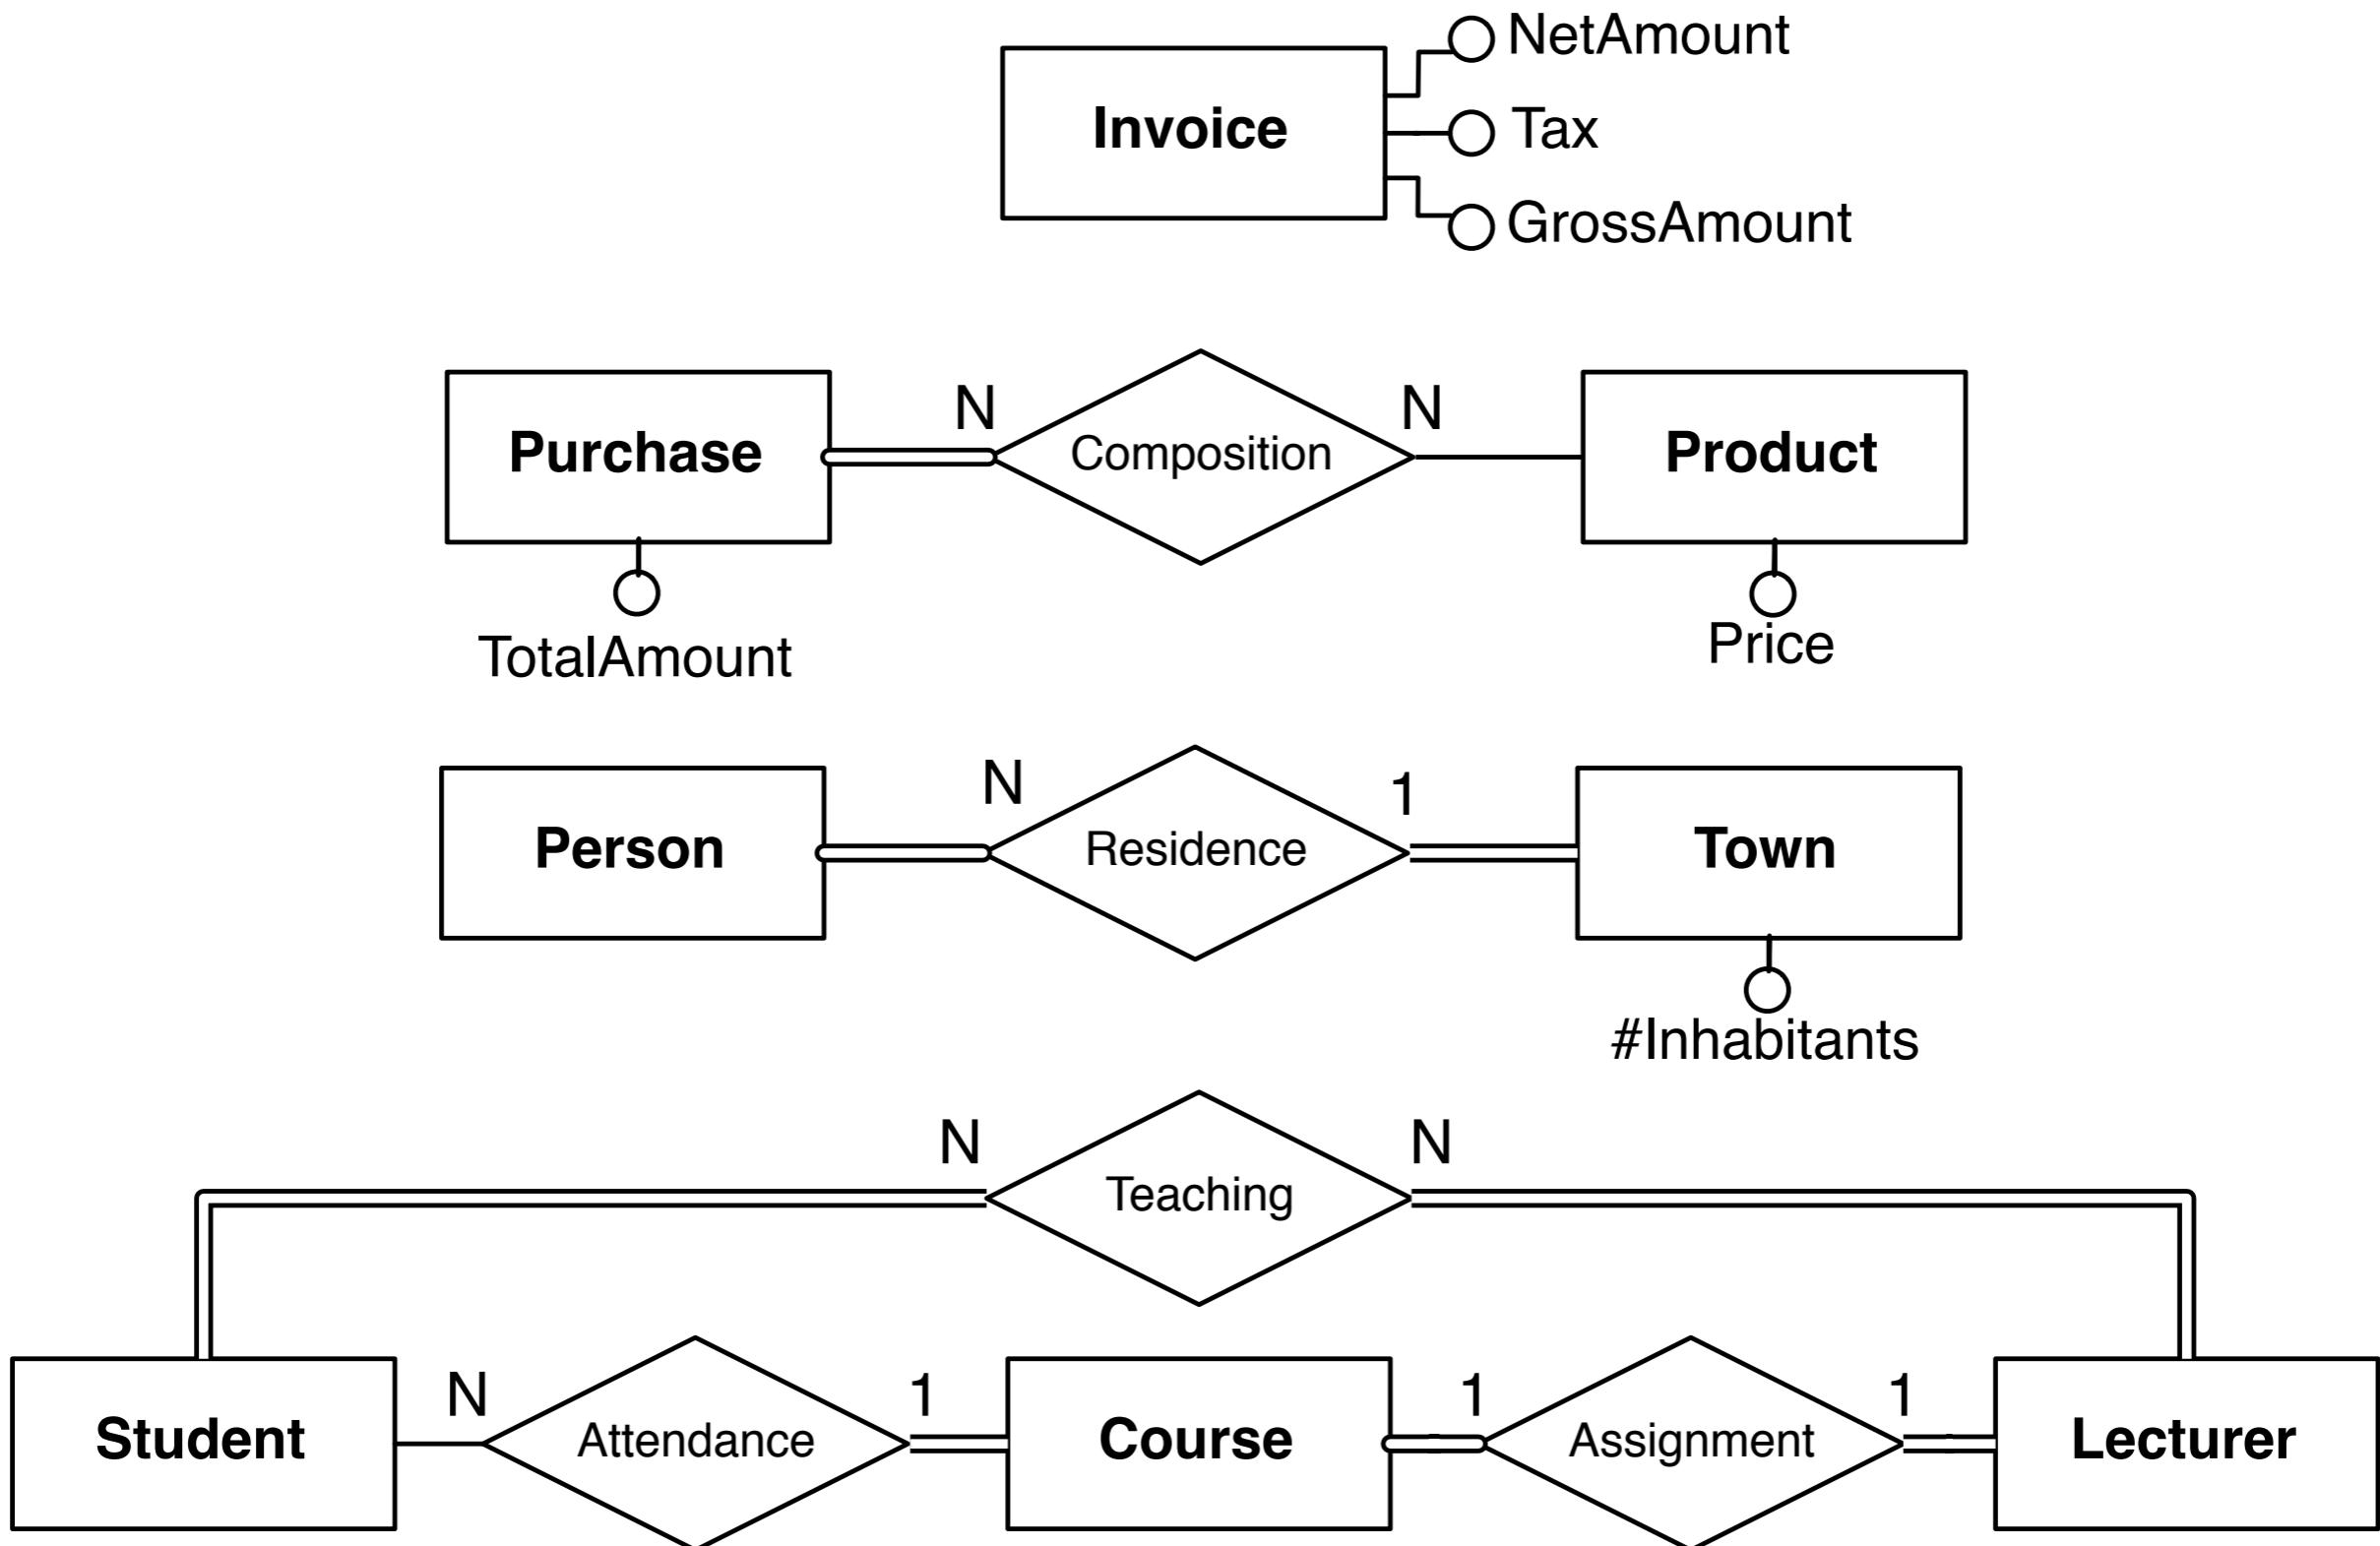
Restructuring Activities

- Analysis of redundancies
- Removal of generalisations
- Partitioning / aggregation of entities and relationships
- Selection of primary keys

Analysis of Redundancies

- Redundancies: relevant information, but derivable from other concepts
 - An Entity-Relationship schema can contain various forms of redundancy
 - Keep, remove, add?
- Effect of redundancies on the logical schema
 - Simpler or faster queries
 - More complex or slower updates
 - Disk space

Example Of Schemas With Redundancies



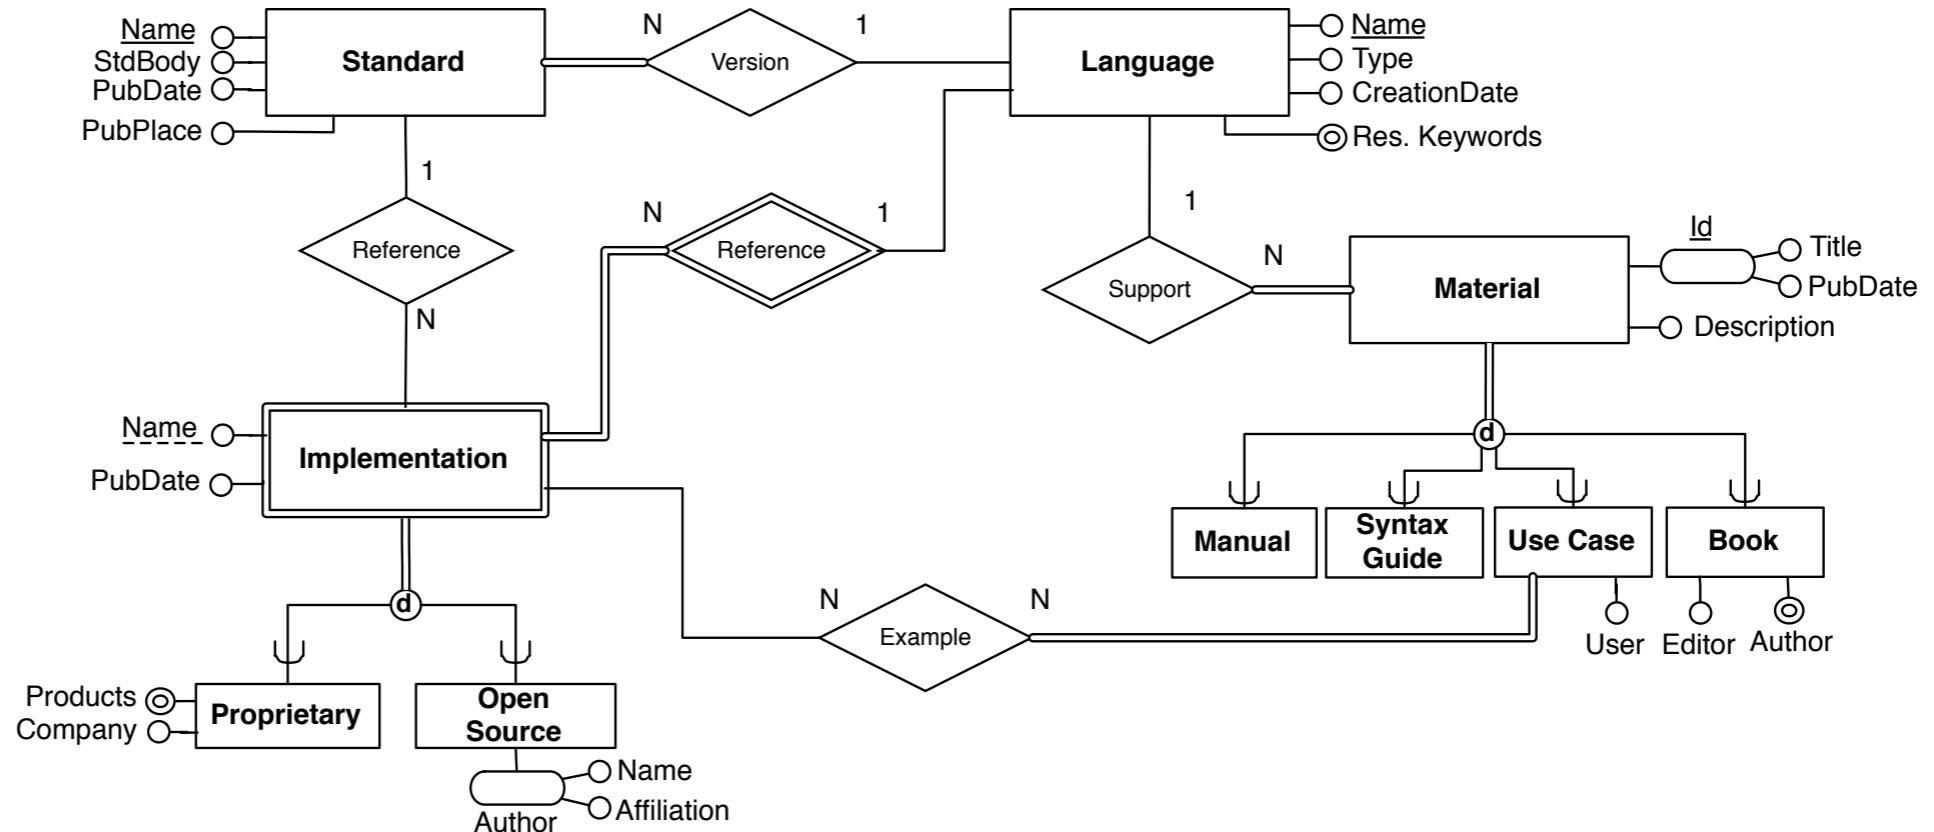
Taking A Decision About Redundancies

- Redundant information in a database can be
 - **An advantage:** a reduction in the number of accesses necessary to obtain the derived information
 - **A disadvantage:** a larger storage requirement (which is often a negligible cost) and the necessity for carrying out additional operations in order to keep the derived data up to date
- Final decision: cost of operations Vs. needed storage

ASQ Exercise (1 min)



Consider the restructuring of the EER diagram in figure.
Which of the following relationships is redundant?



1

Reference, between Standard and Implementation

2

Reference, between Language and Implementation

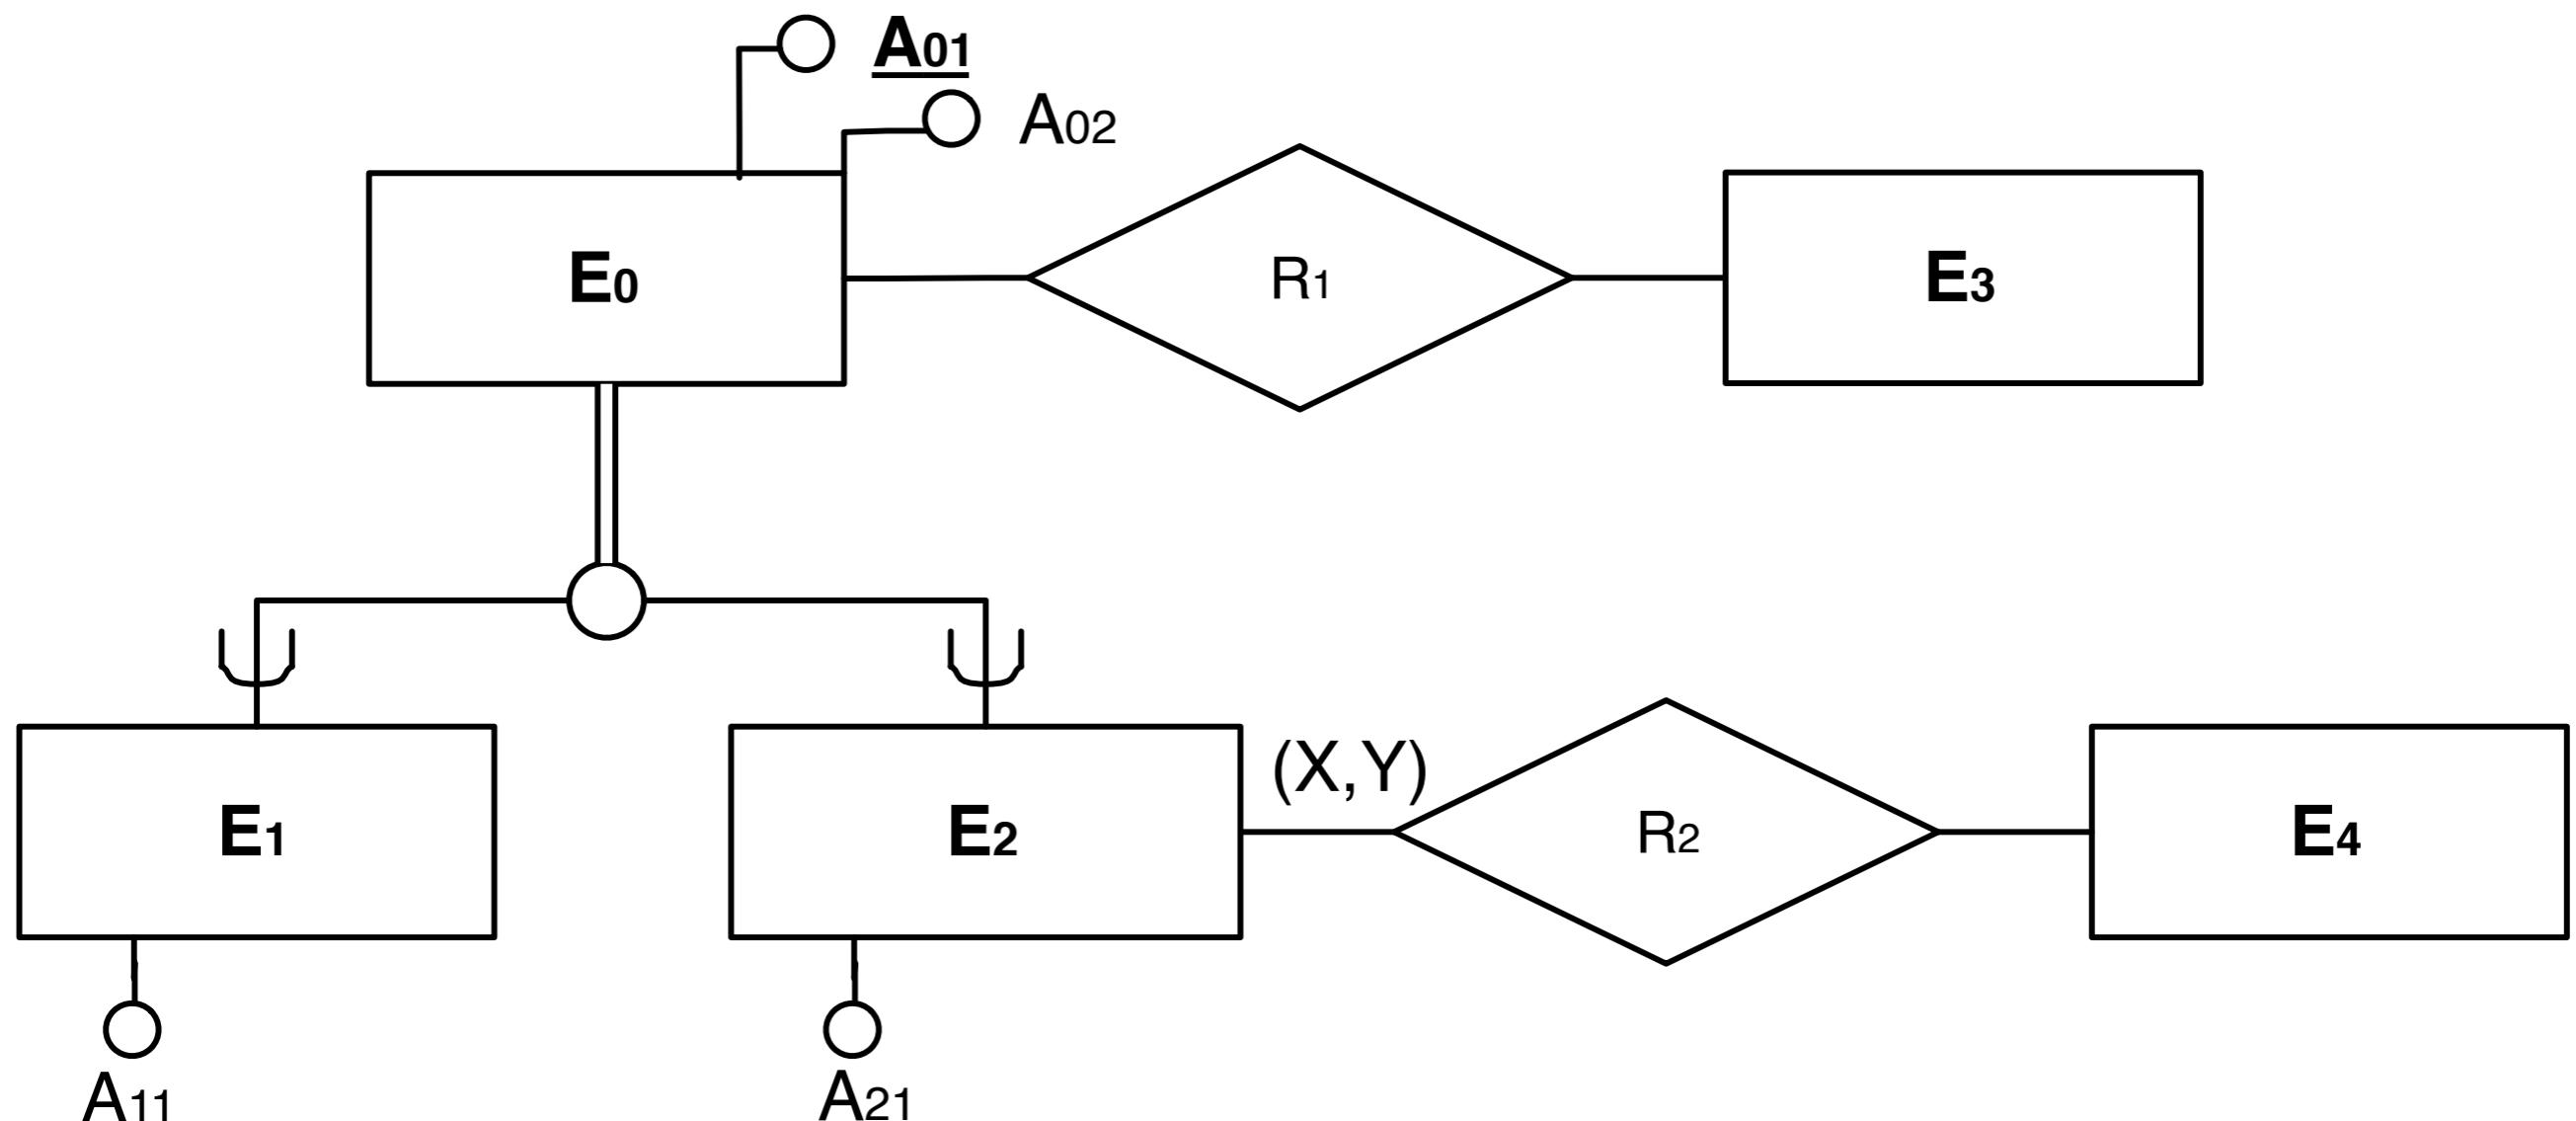
3

Example

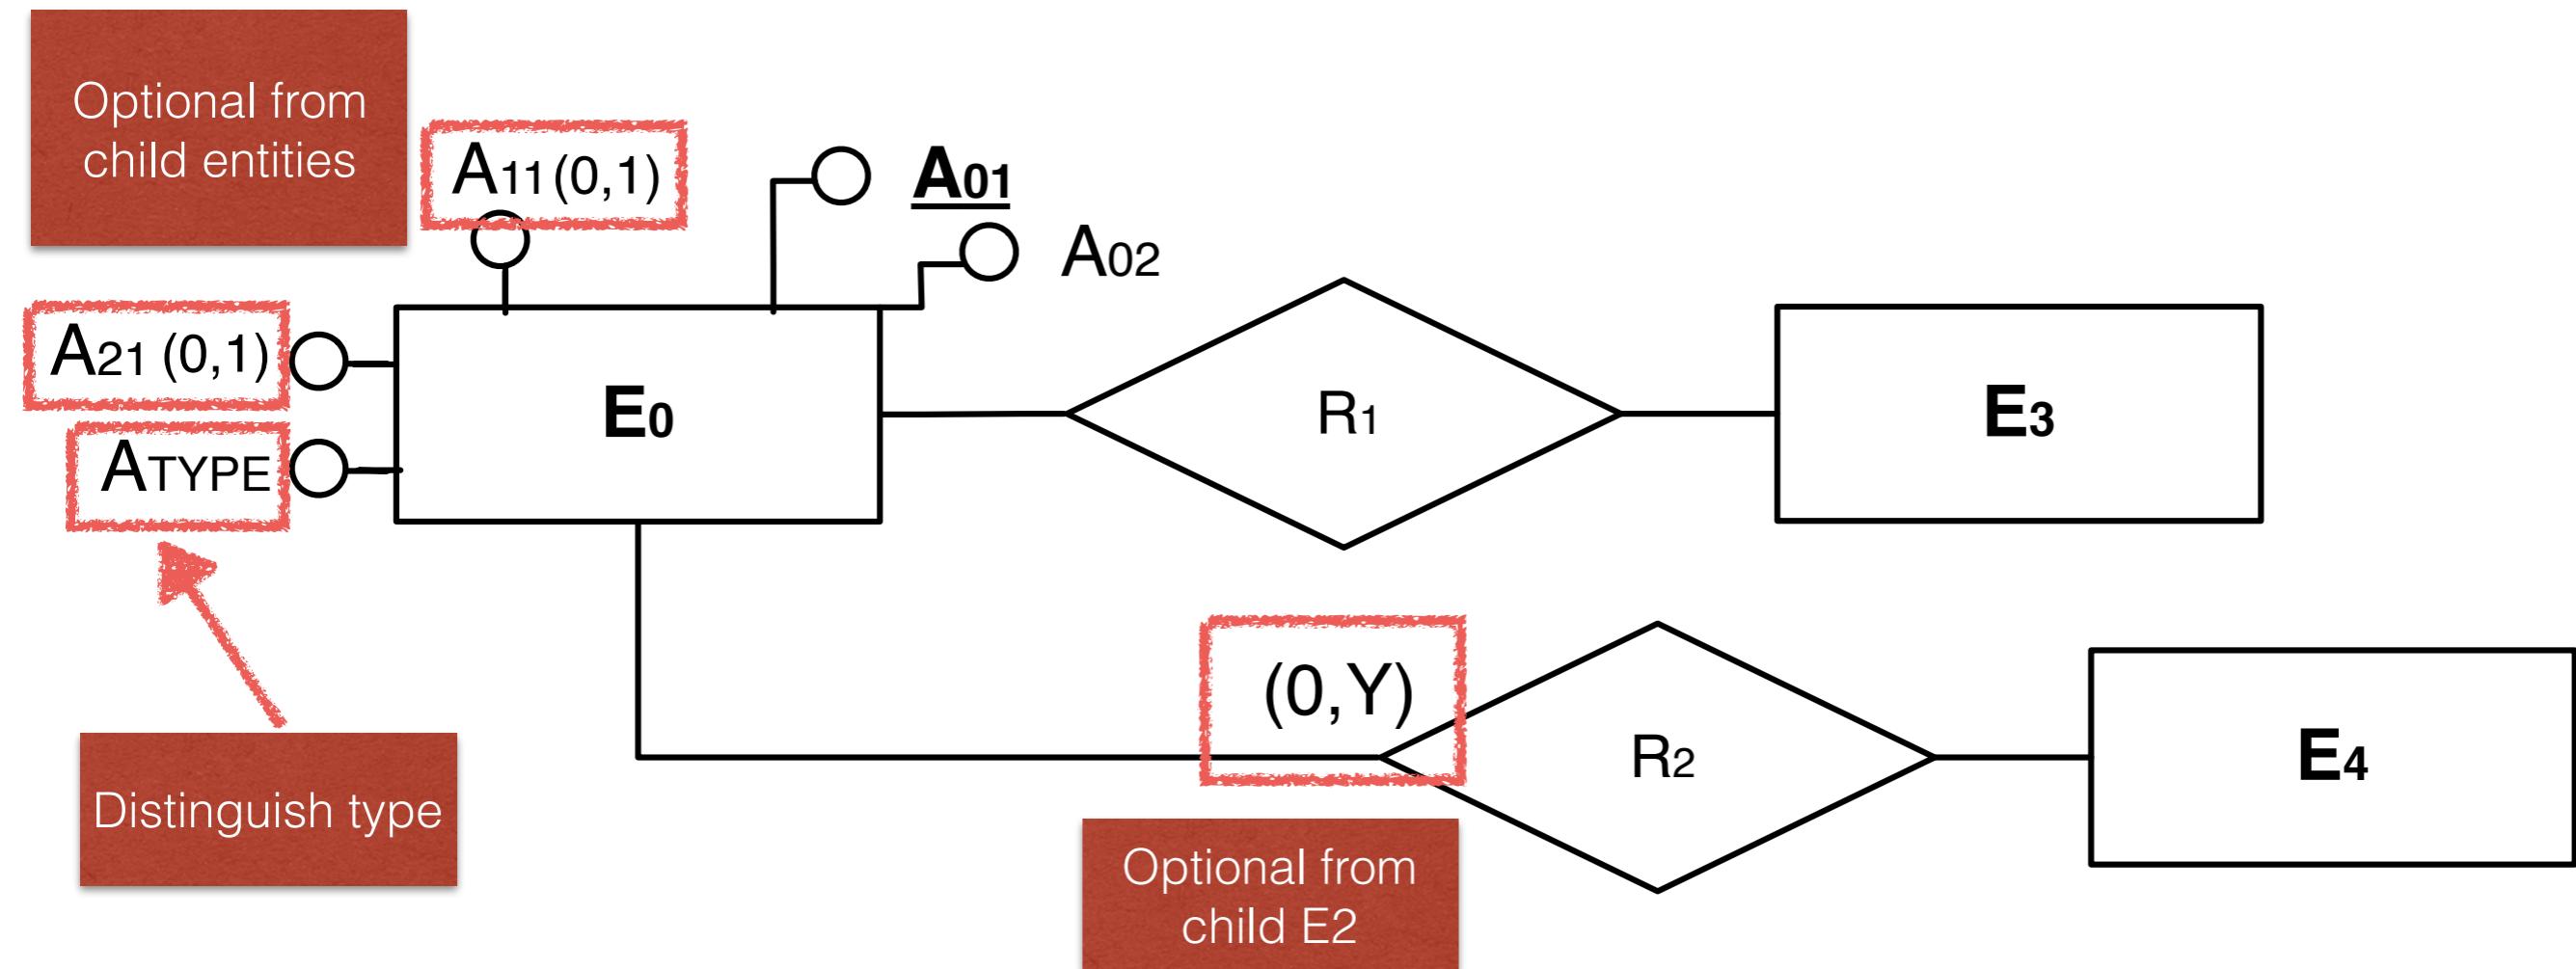
Removing Specialisations

- Specialisation/Generalisation **cannot be directly represented** in the relational model
- Substituted by entities and relationships
- Methods
 - 1.collapse child entities in the parent entity
 - 2.collapse parent entity on child entities
- **3.substitute the specialisation/generalisation with relationships**
 - **This is our standard way**
- Hierarchies: apply approach, **starting from the deeper levels**

Reference Generalisation Schema

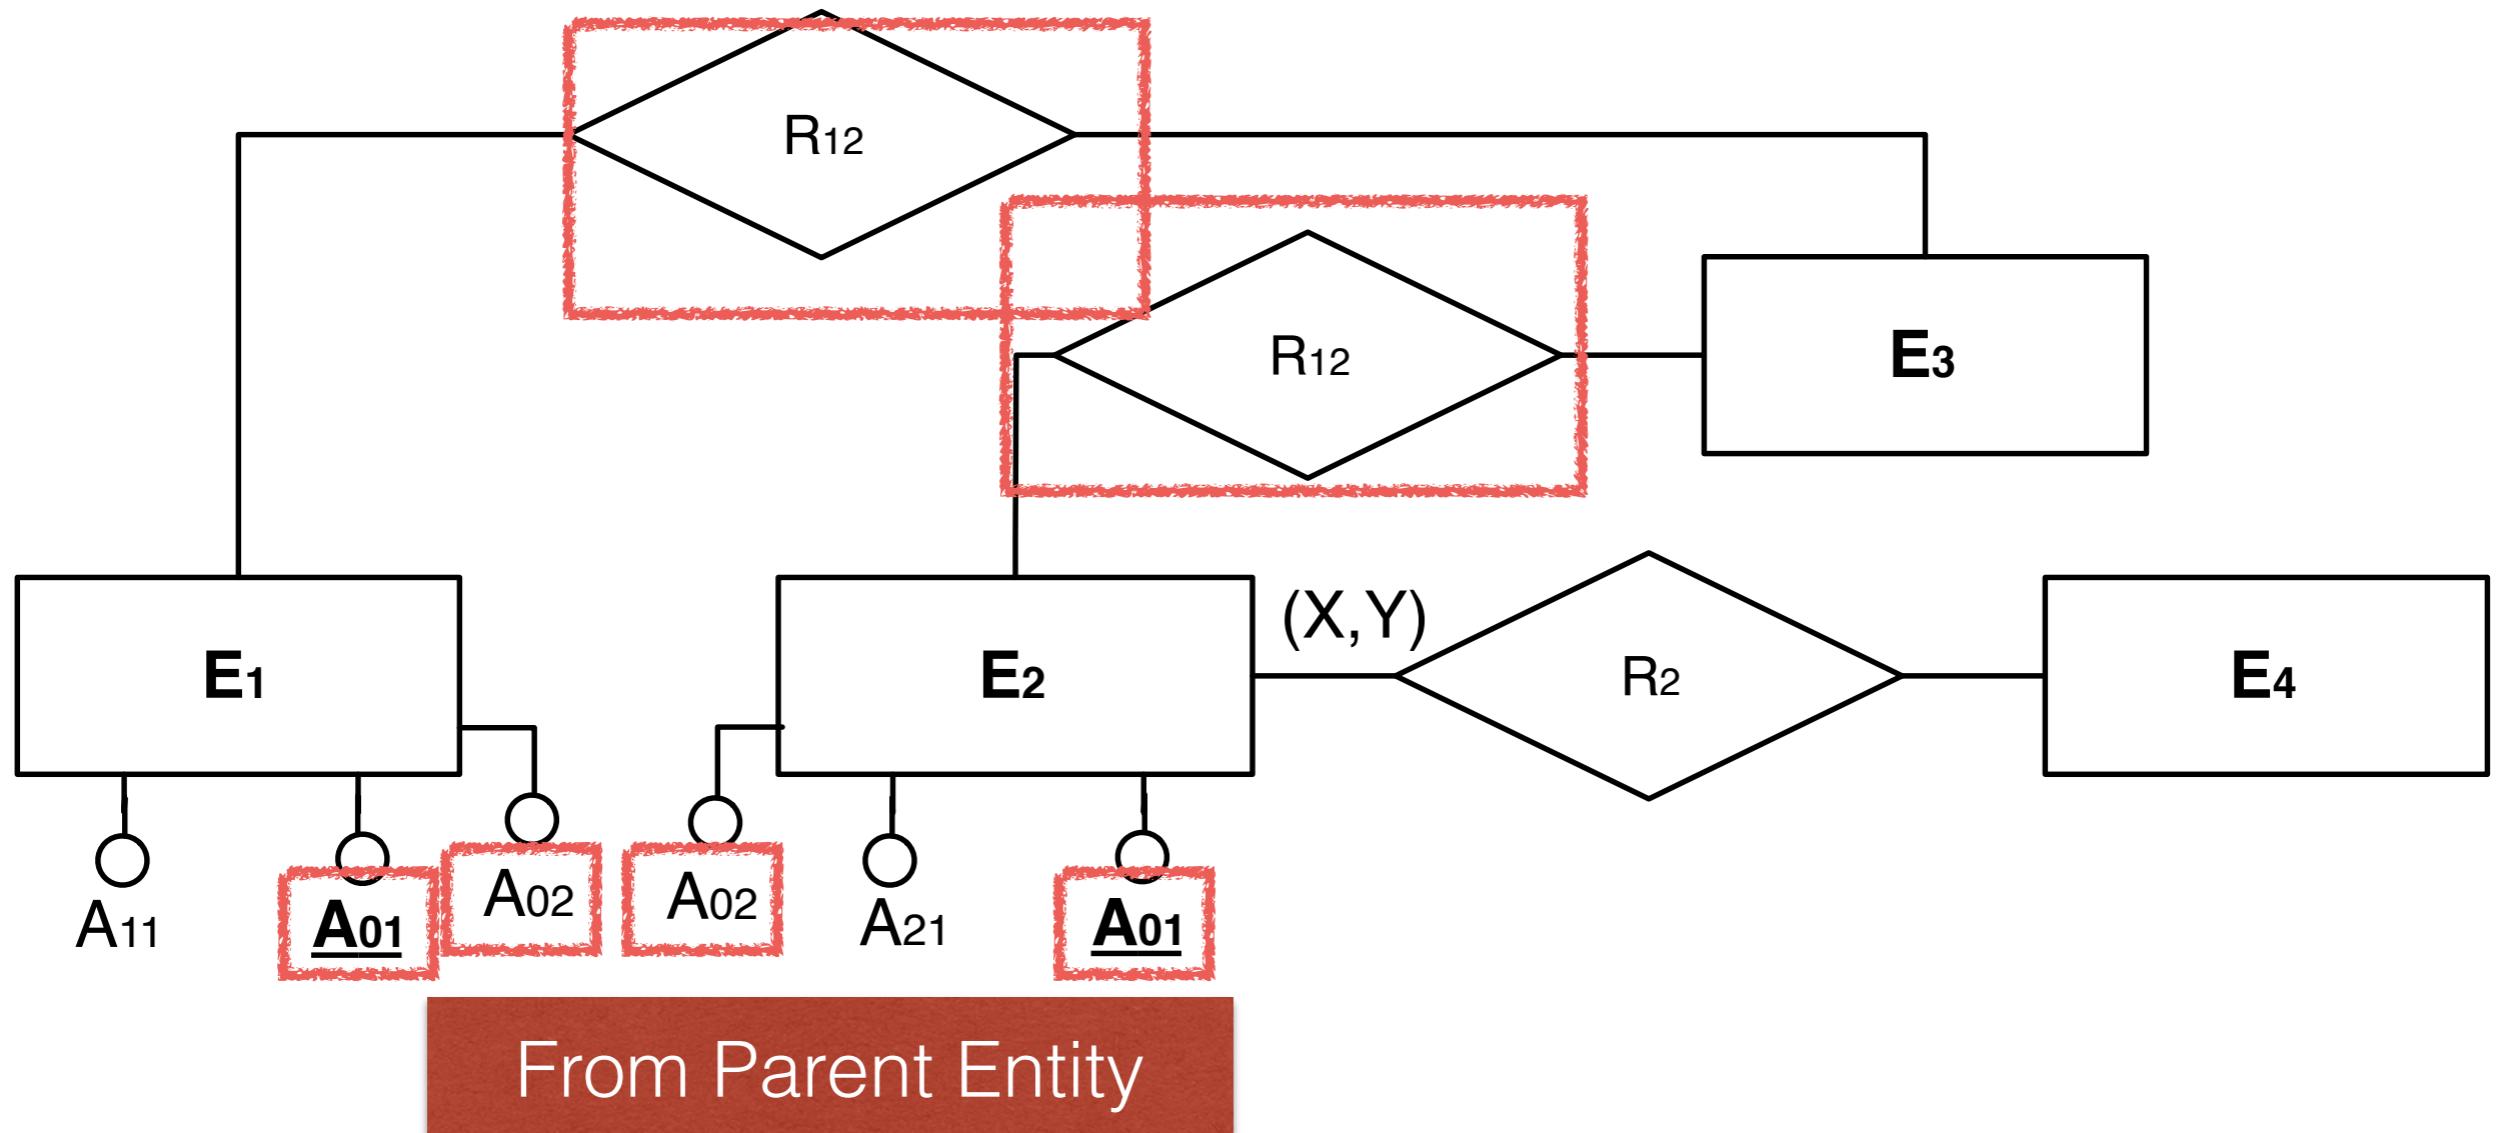


Parent Collapsing



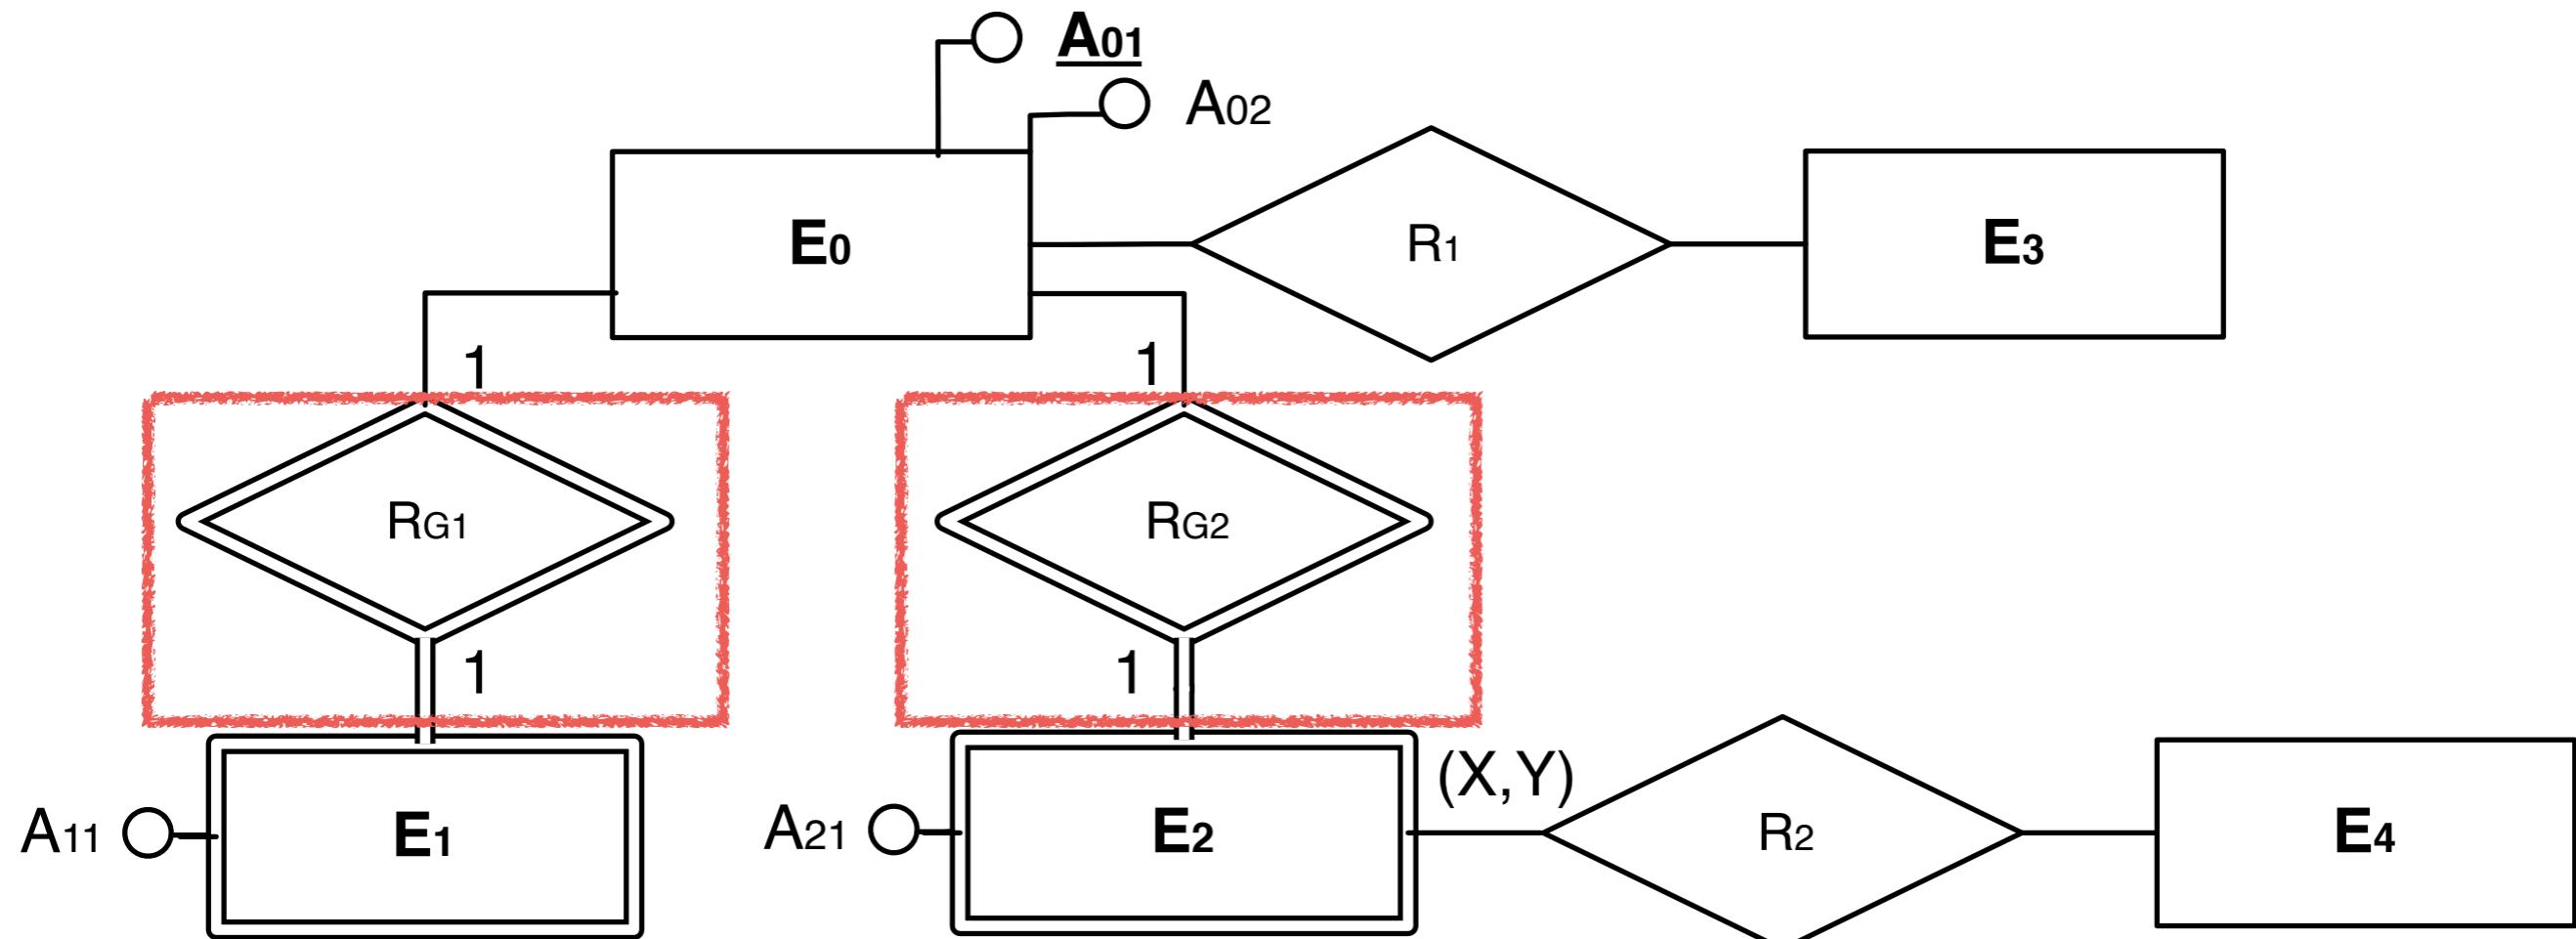
- More appropriate when
 - Child entities introduce little differences (so result in few NULL values)
 - Access operations do not need to distinguish between parent and child entities

Child Collapsing



- Possible only if the **specialisation** is **total**
- More appropriate when access operations need to **distinguish** between different child entities (more efficient)

Substitute with Relationships



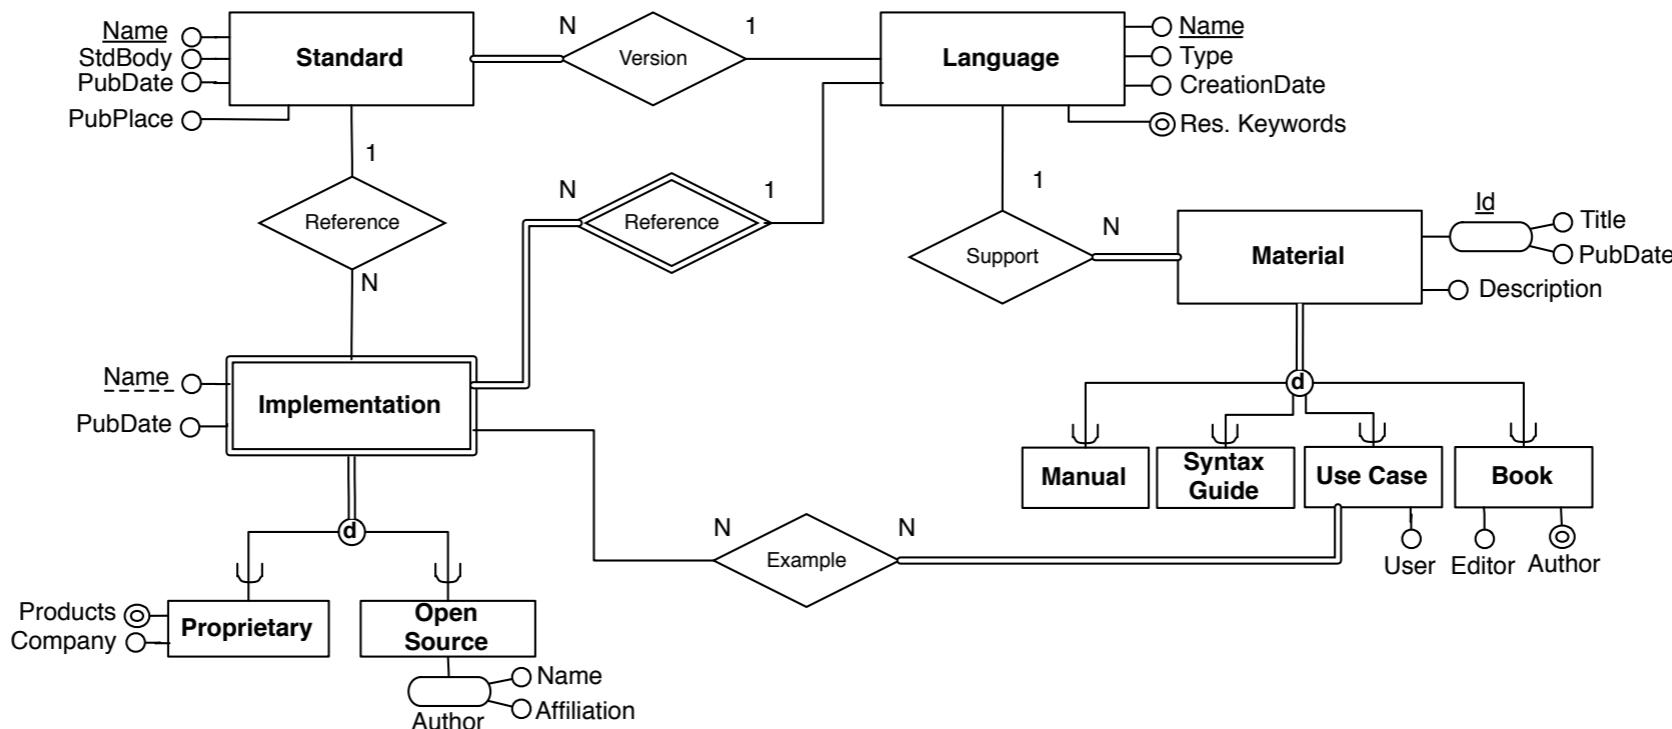
- **Most General Solution**, and always applicable
- But reconstructing the original information can be **expensive**

ASQ Exercise (3 min)



Consider the restructuring of the EER diagram in figure, performed with a *parent-collapsing* specialisation removal method

Which of the following statements is/are correct? There can be more than one



1

The Material entity participates in the Example relationship with a partial participation constraint.

2

The Implementation entity includes 5 attributes, but only the Author attribute becomes optional because it is multi-valued

3

The Implementation entity features 6 attributes, but only 3 become optional

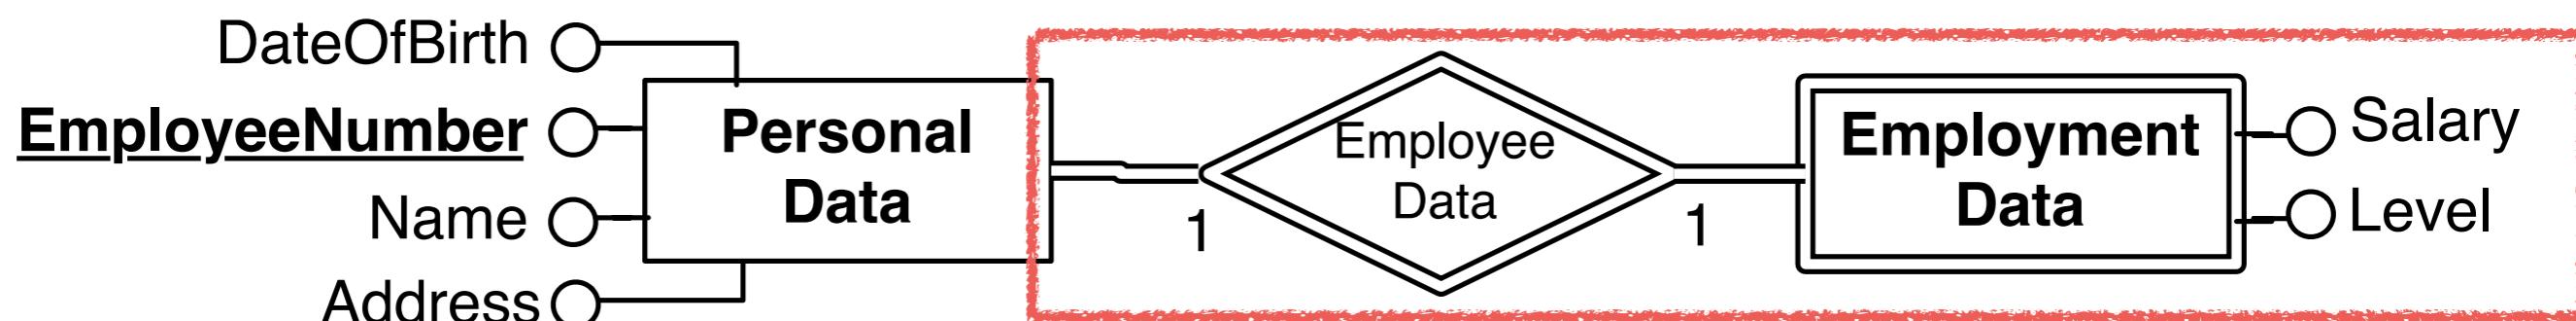
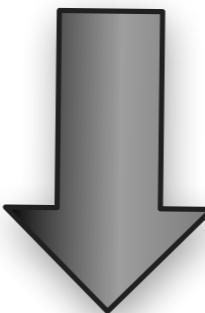
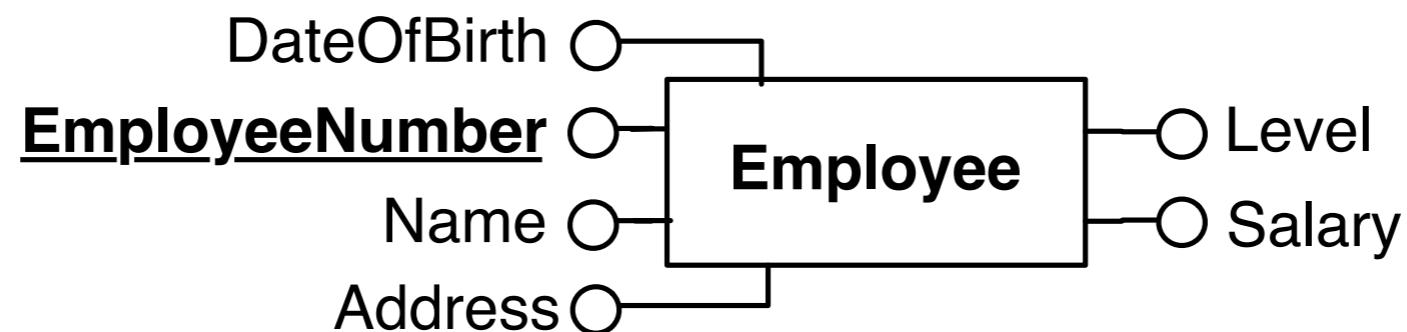
4

The Example relationship changes its cardinality from N:N to 1:N because not all the instances from the Material entities will be related to an instance of the Implementation entity

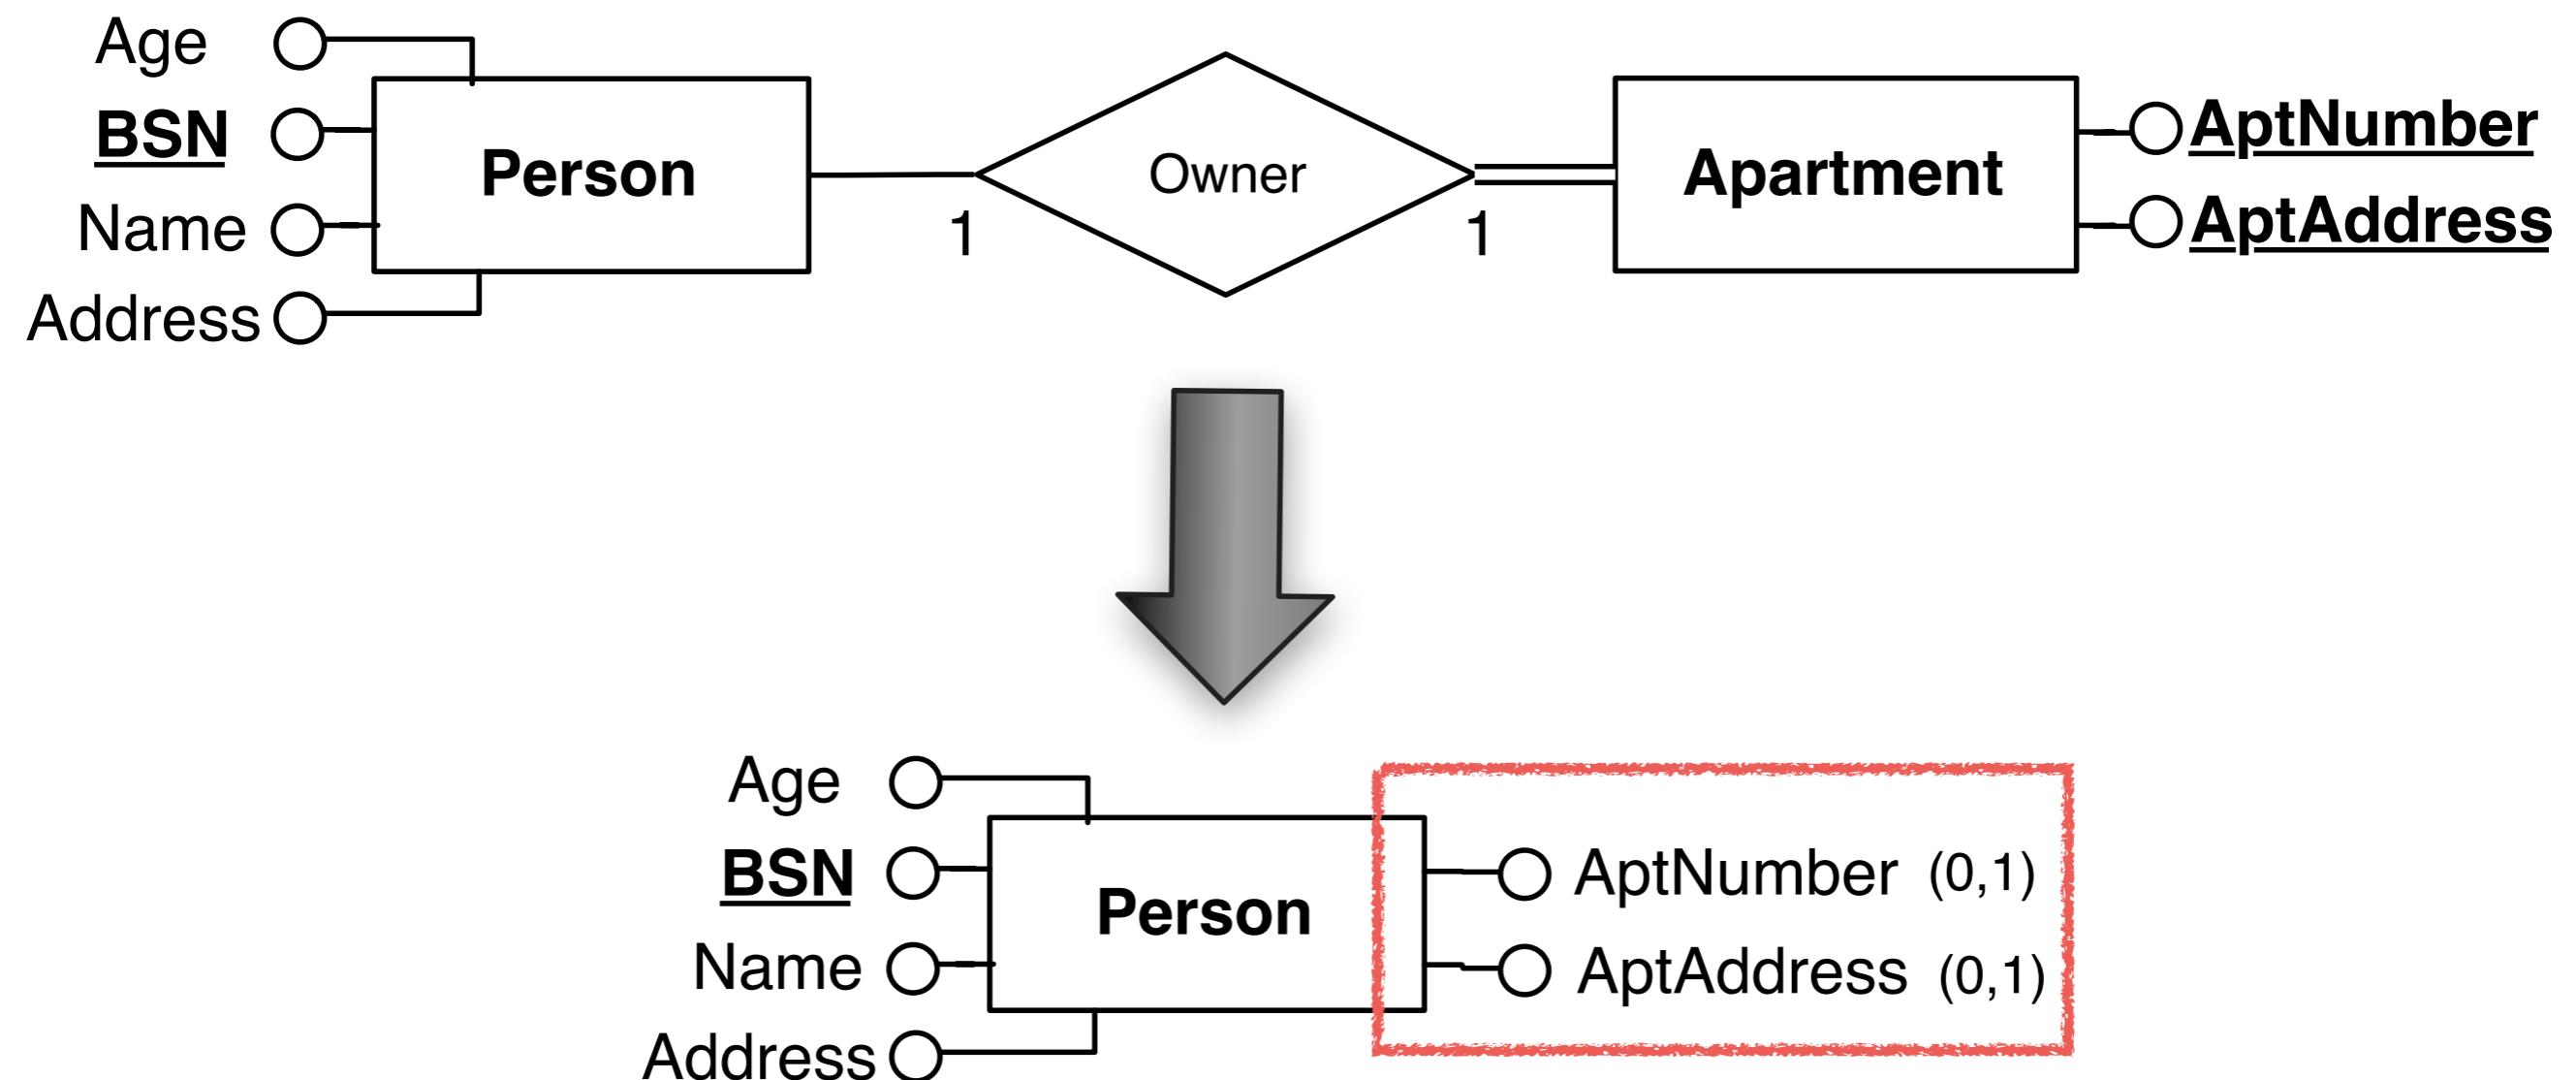
Partitioning and Merging of Concepts

- Entities and relationships of an E-R schema can be **partitioned** or **merged**
- Better representation of different concepts
- Improve the efficiency of operations. Accesses are reduced by:
 - Separating attributes of the same concept that are accessed by different operations
 - Merging attributes of different concepts that are accessed by the same operations
- The same criteria as for redundancies are valid in making a decision about this type of restructuring

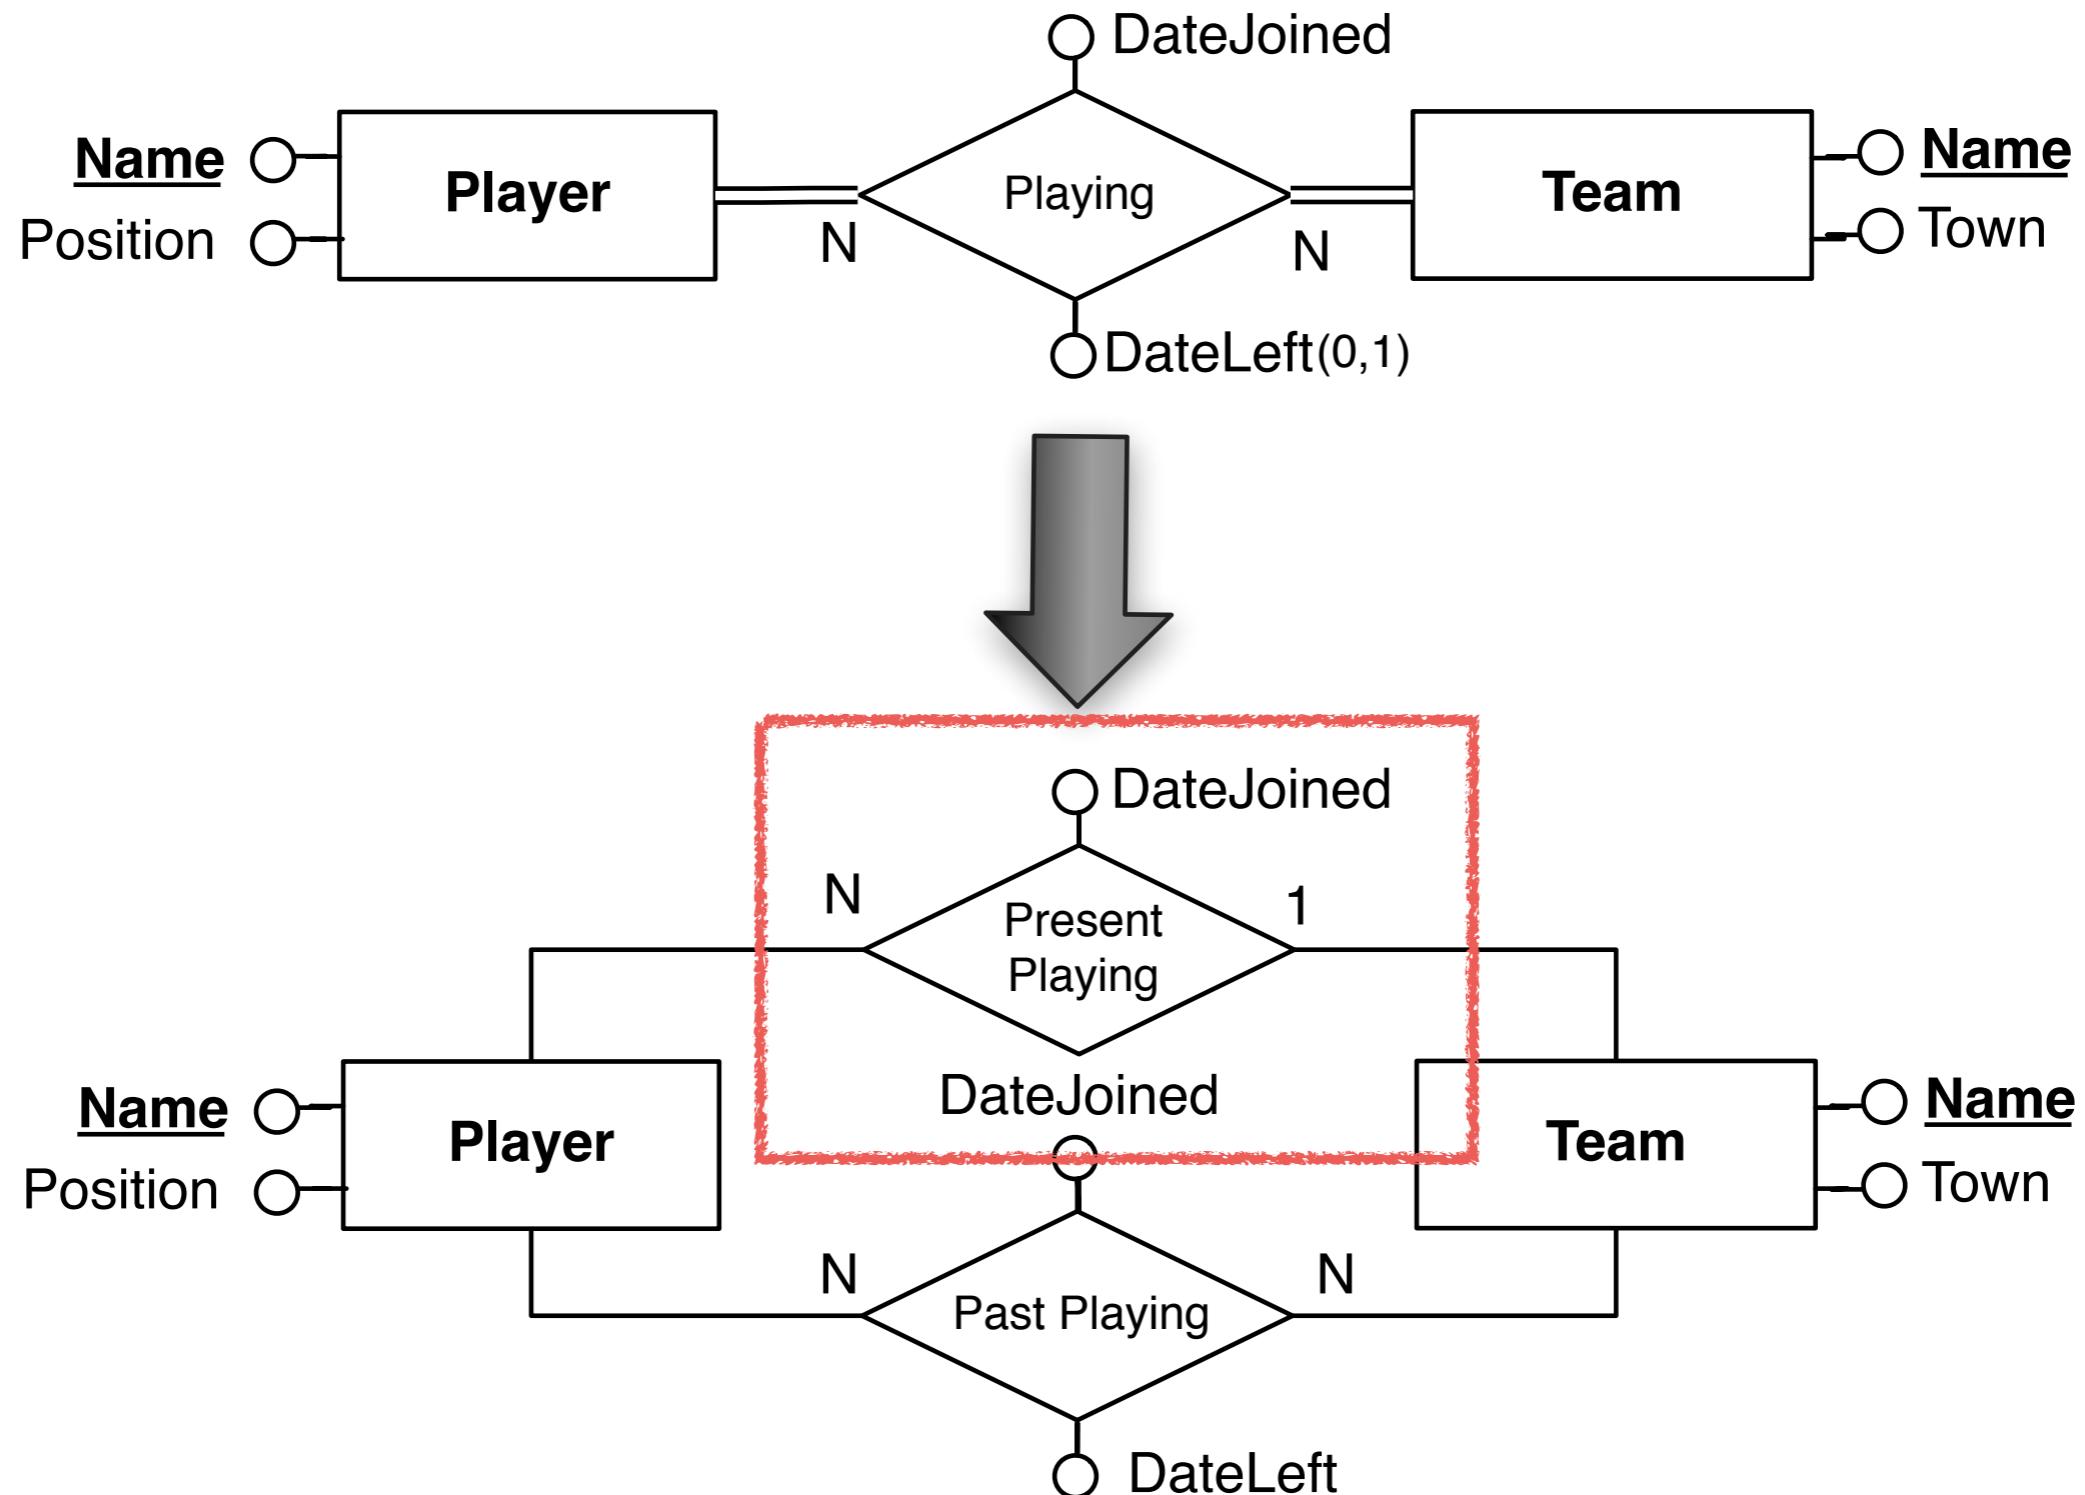
Example Schema: Entity Partitioning



Example Schema: Entity Merging



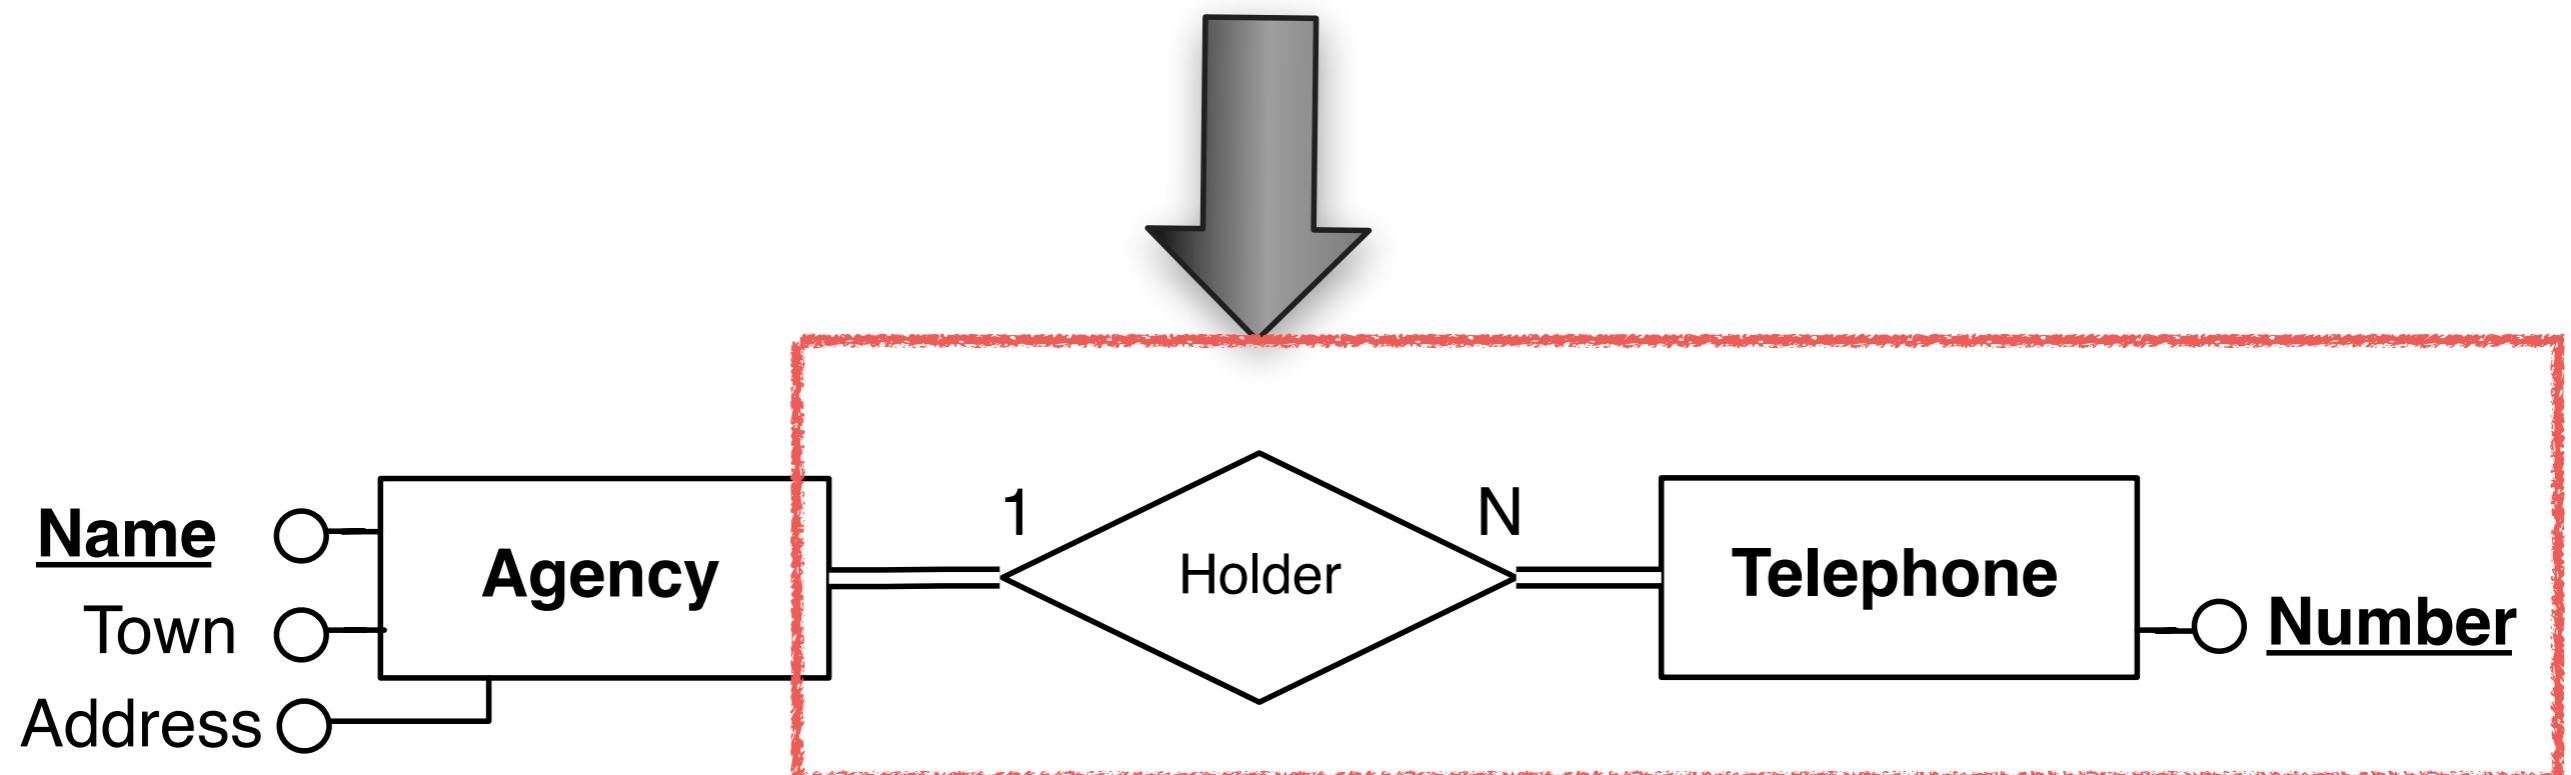
Example Schema: Relationship



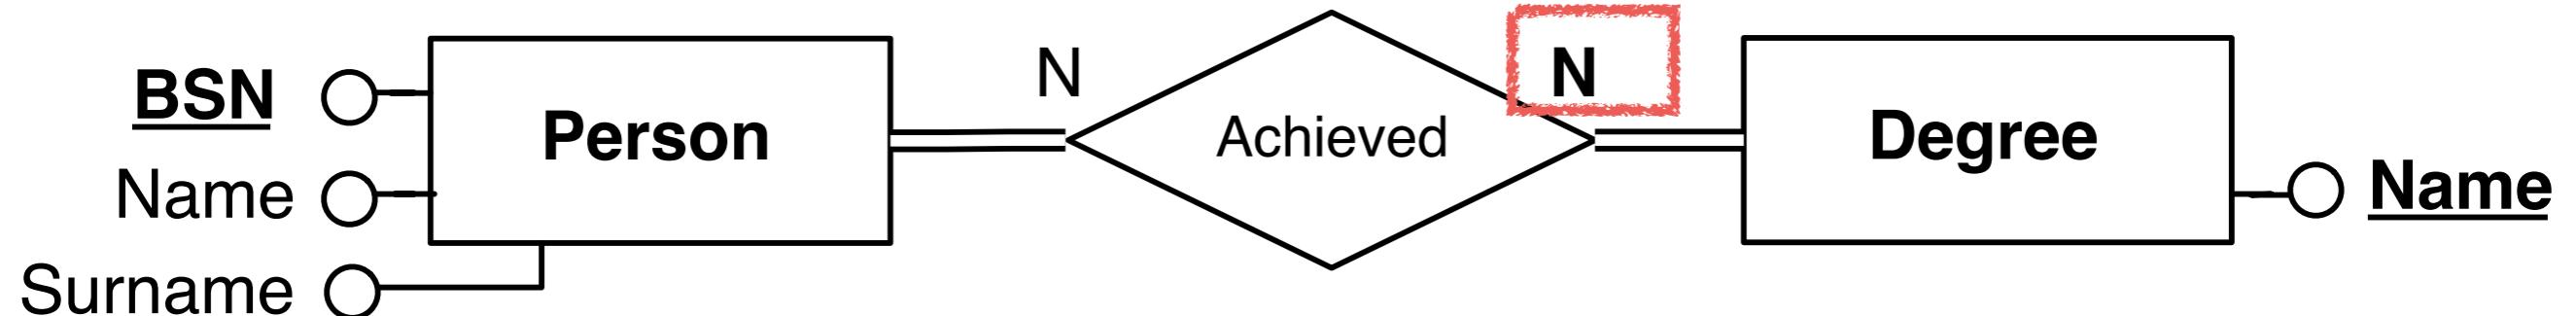
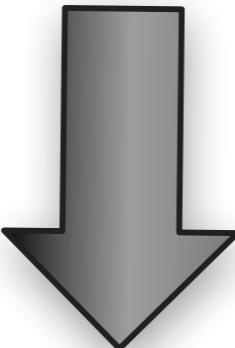
Removal of Multi-Value Attributes

- No representation in the relational model
- Multivalued attribute represented by a **new entity** associated to the original entity
 - Cardinalities of new relationships depend on cardinality of original attribute
 - But also on the nature of the attribute!

Example Schema: Multi-Value / I

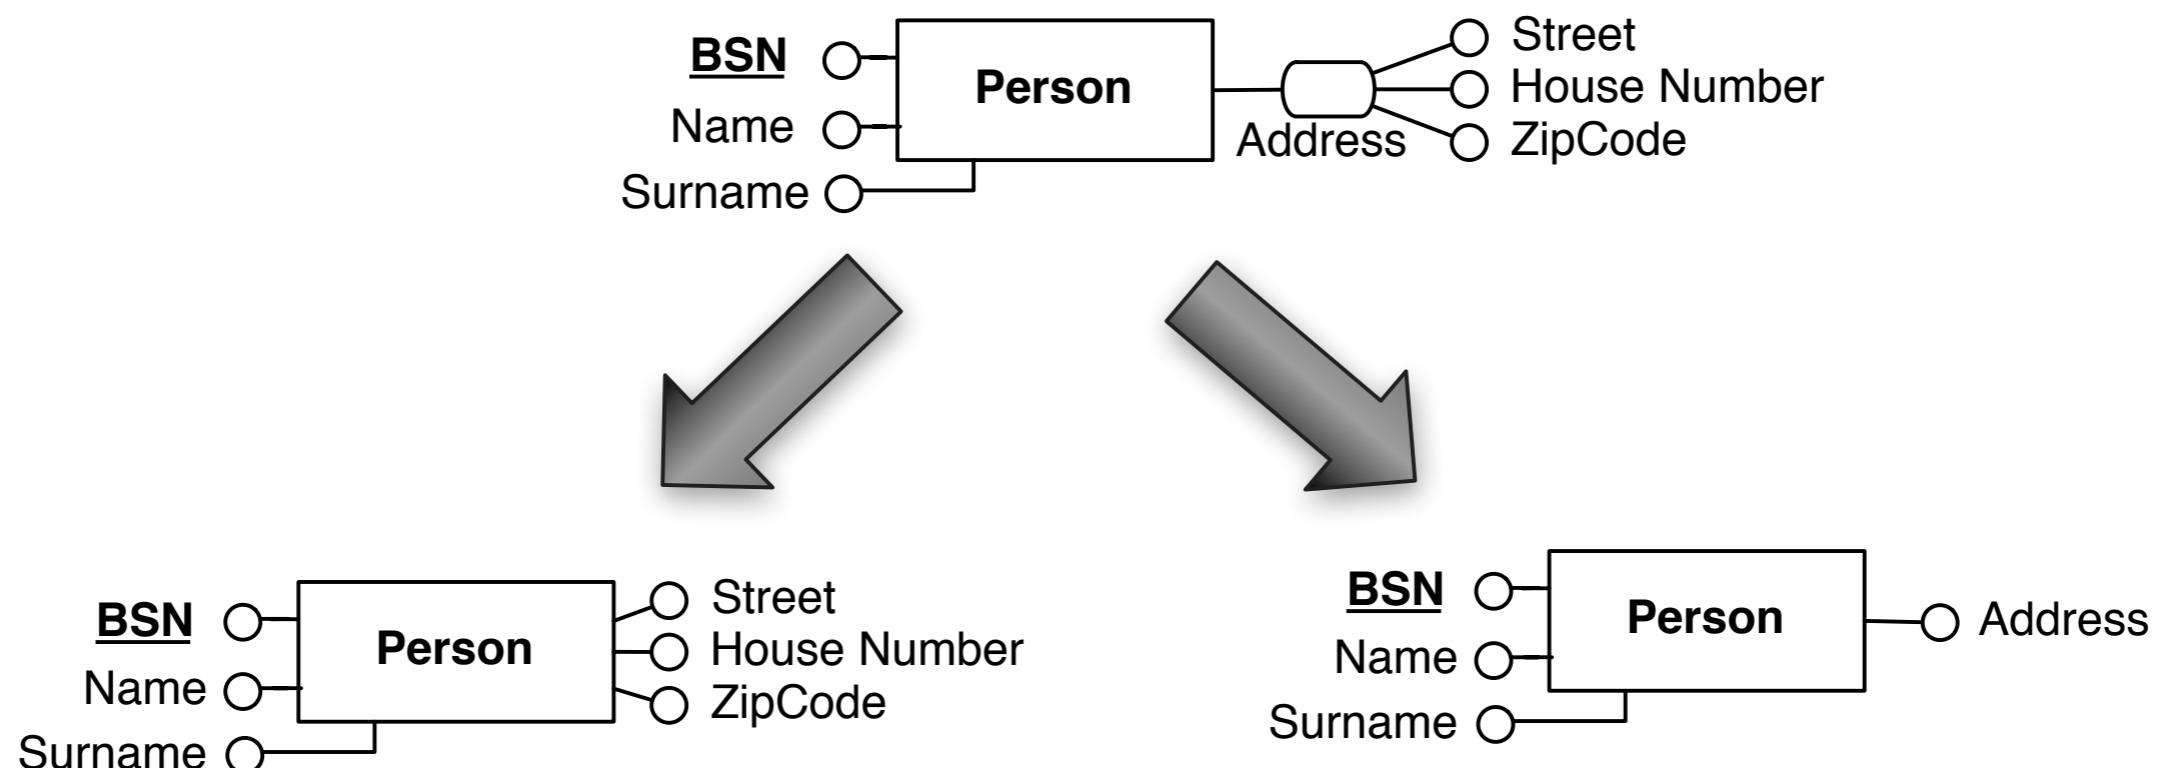


Example Schema: Multi-Value /2



Removal of Complex Attributes

- No representation in the relational model
- Two alternatives:
 - Collapse sub-attributes
 - Better if need to access the single attributes
 - Or composite attribute is candidate identifier
 - Introduce a new attribute which concatenates the others
 - Better if access to the whole information



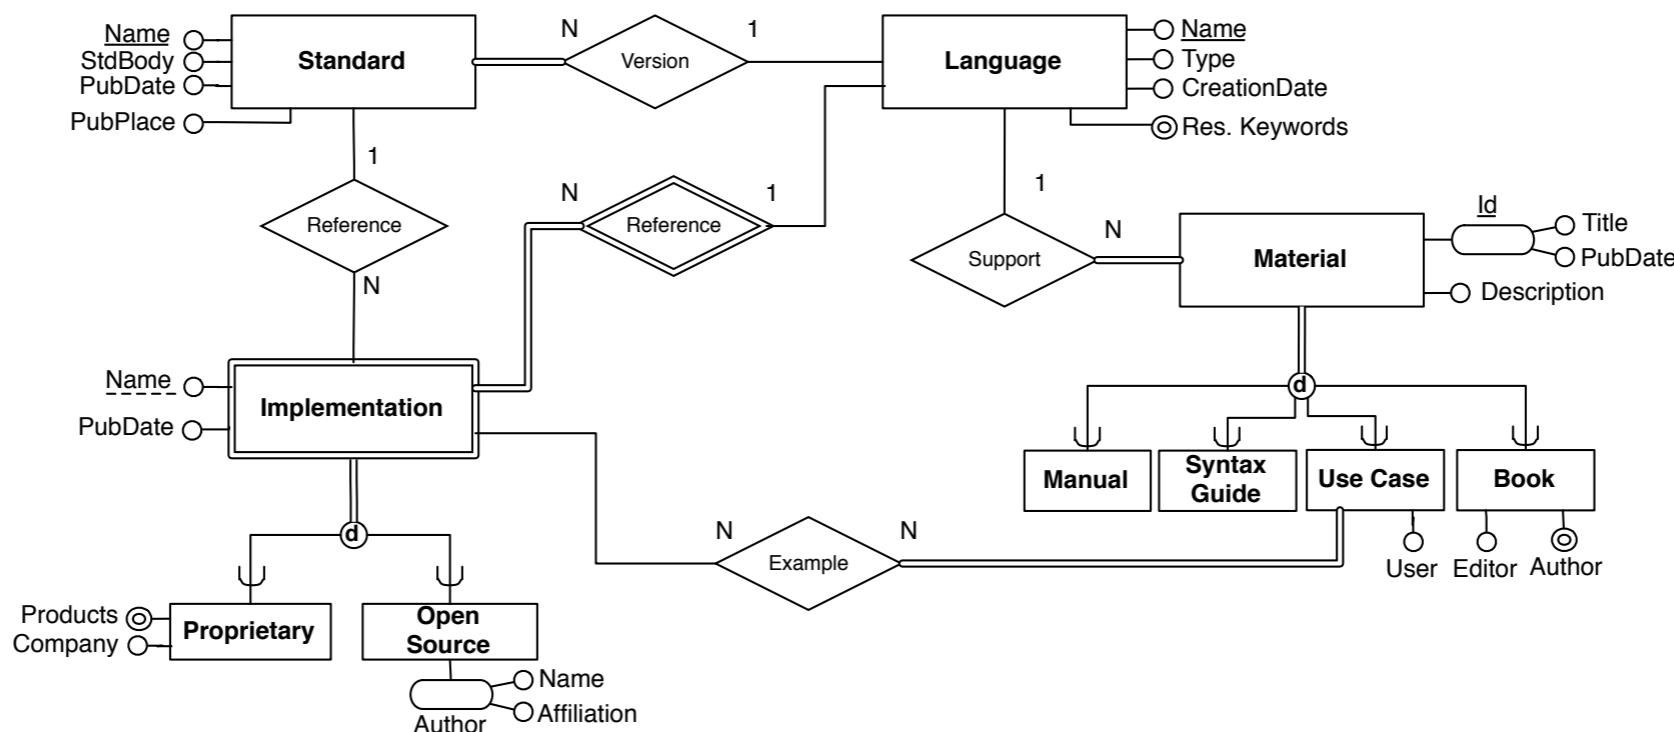
Selection of Primary Identifiers

- Required to define **primary keys** for tables
- A good identifier
 - no NULL values
 - **One or few** attributes
 - **internal attribute**, better than external ones
 - Is often used by many access operations
- If none of the candidate identifiers satisfies the above requirements
 - introduce further attribute to the entity
 - special values (often called codes) generated solely for the purpose of identifying occurrences of the entity

ASQ Exercise (1 min)

A S Q

Consider the restructuring of the EER diagram in figure.
Which of the following statements about the *Title* and *PubDate* attributes of the *Material* entity is correct?



1

They are merged into a single *id* attribute

2

They are preserved as independent attributes, and an *Id* attribute is a new primary identifier

3

They are preserved as independent attributes, and together they become primary identifiers

4

They are discarded from the model

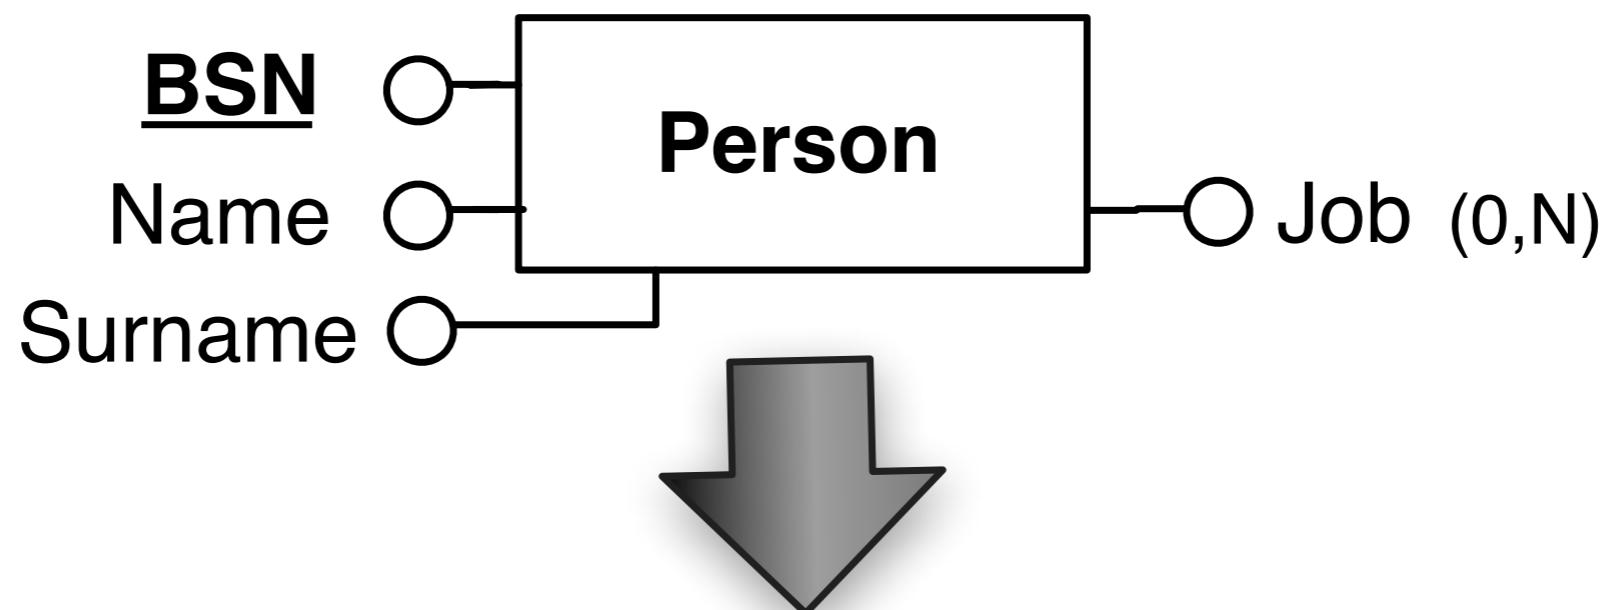
Translation into the Relational Model

Creation Of The Relational Model

- Transformation between different data models
- Starting from a (restructured) E-R schema, **an equivalent** relational schema is constructed
 - **Equivalent:** a schema capable of representing the same information
- Transformations
 - Each entity maps to a table having the same attributes
 - Mapping of relationships considers maximum cardinality

Example of Entity Transformation

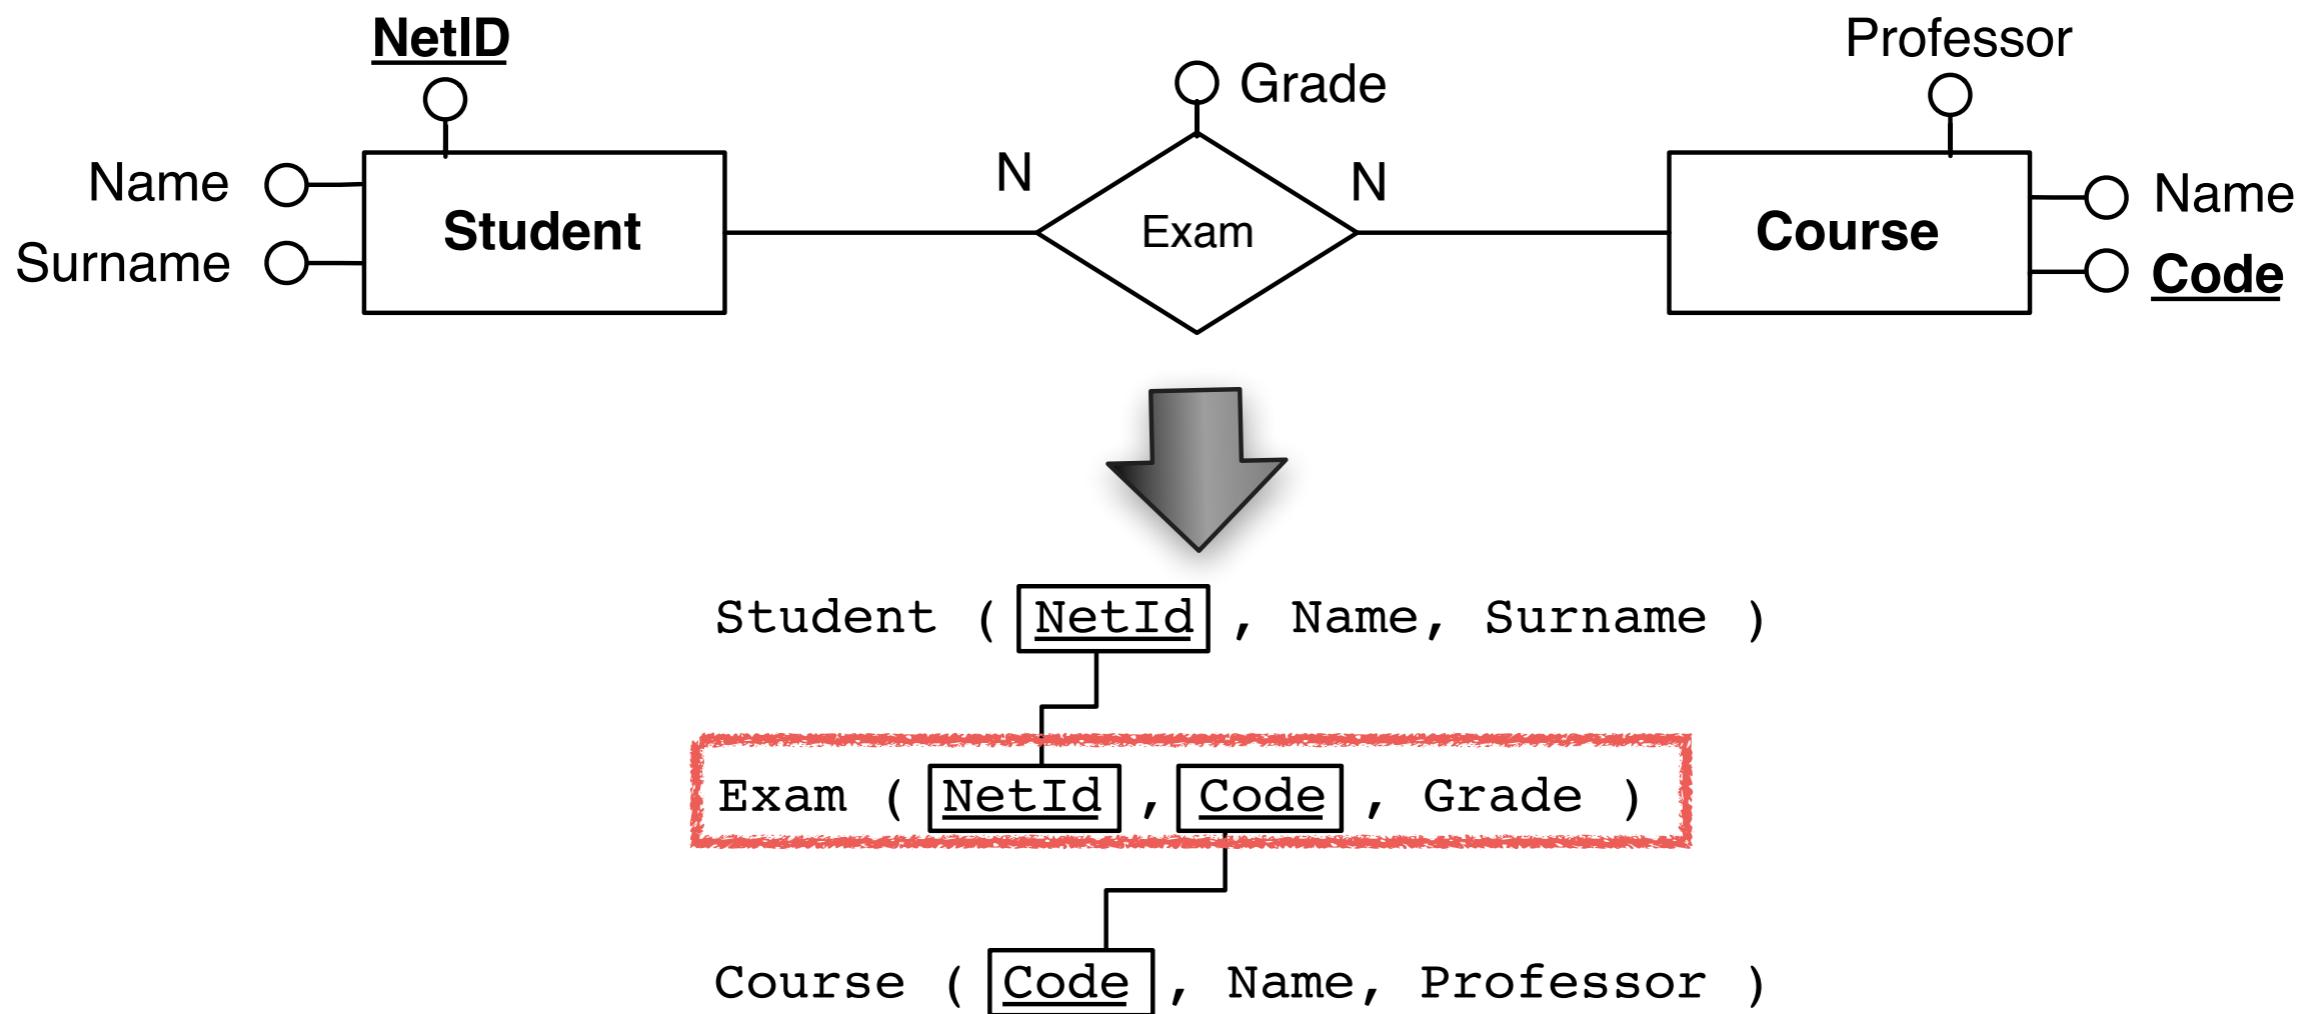
- Primary Key is underlined
- Optional attributes are denoted by an *



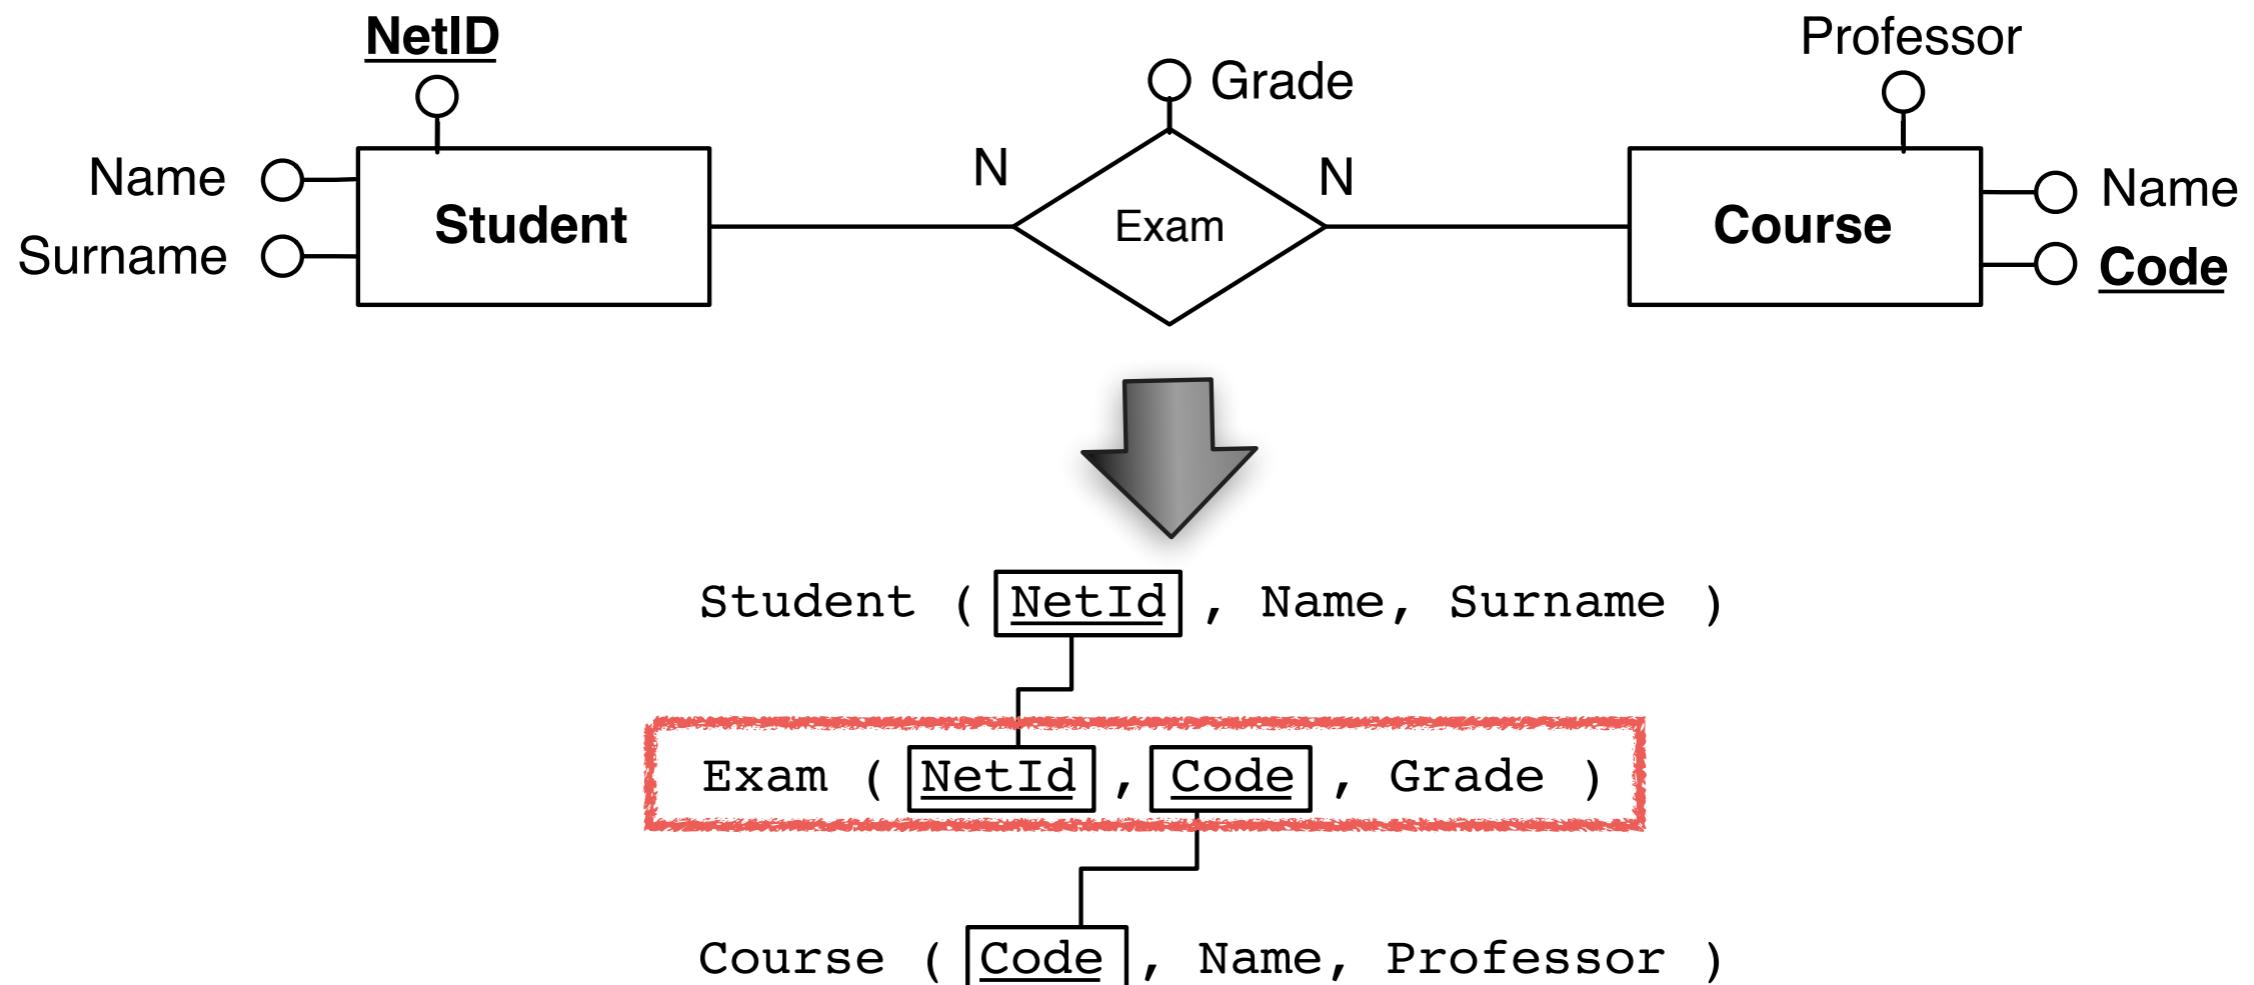
Person (BSN, Name, Surname, Job*)

Translation of Many to Many binary relationship

- Each many to many relationship maps to a **new table**
 - Primary key is the **combination** of the primary keys of the related entities
 - Name of attributes in the relationship table can be changed (in case of recursive relationships)

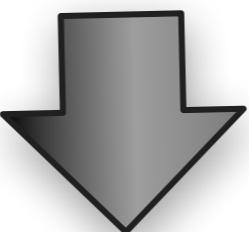
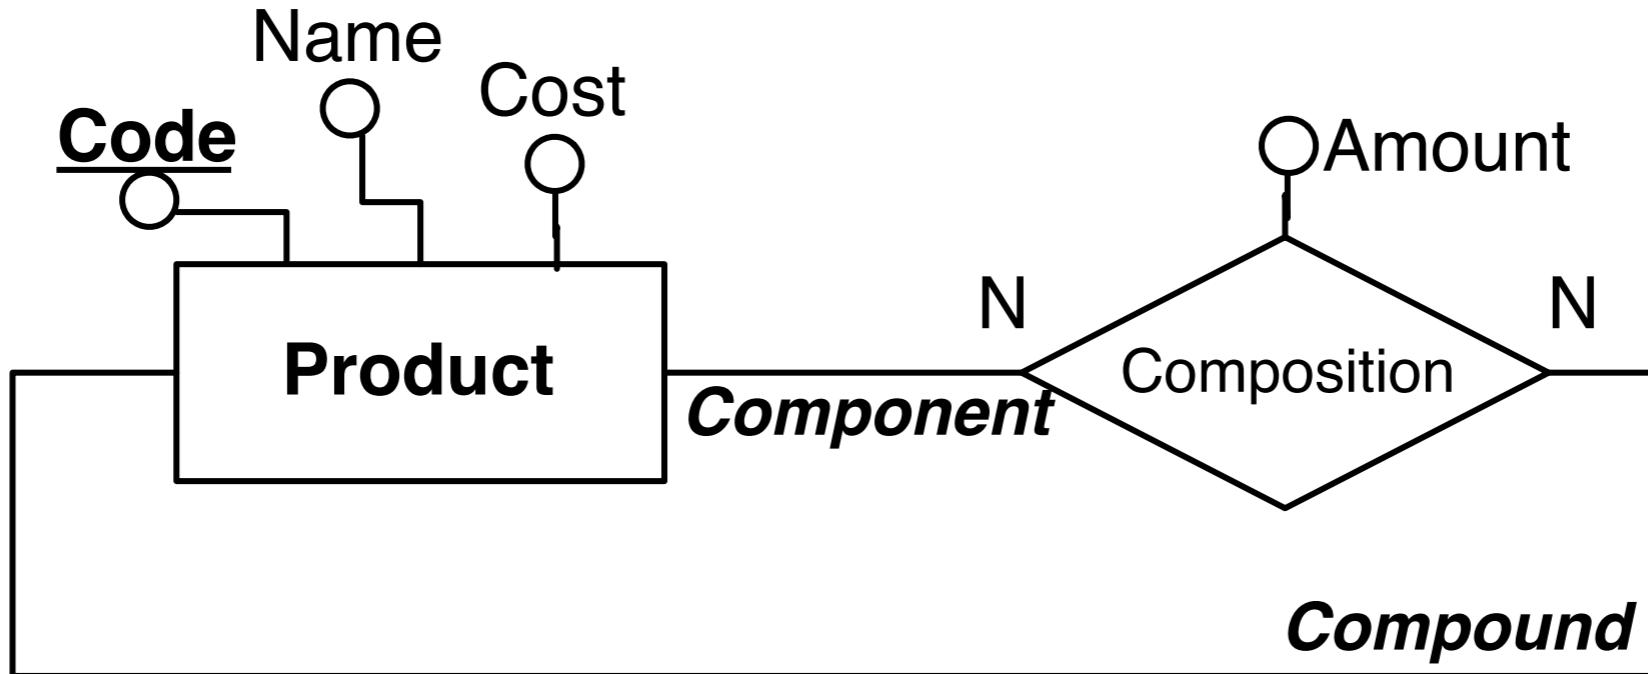


Example of Referential Integrity Constraint



```
CREATE table Exam
(
    NetId      CHARACTER(6) REFERENCES Student(NetId)
    Code       CHARACTER(6) REFERENCES Course(Code),
    Grade     DECIMAL (4) DEFAULT 0,
    PRIMARY KEY (NetId, Code)
```

Recursive Many To Many



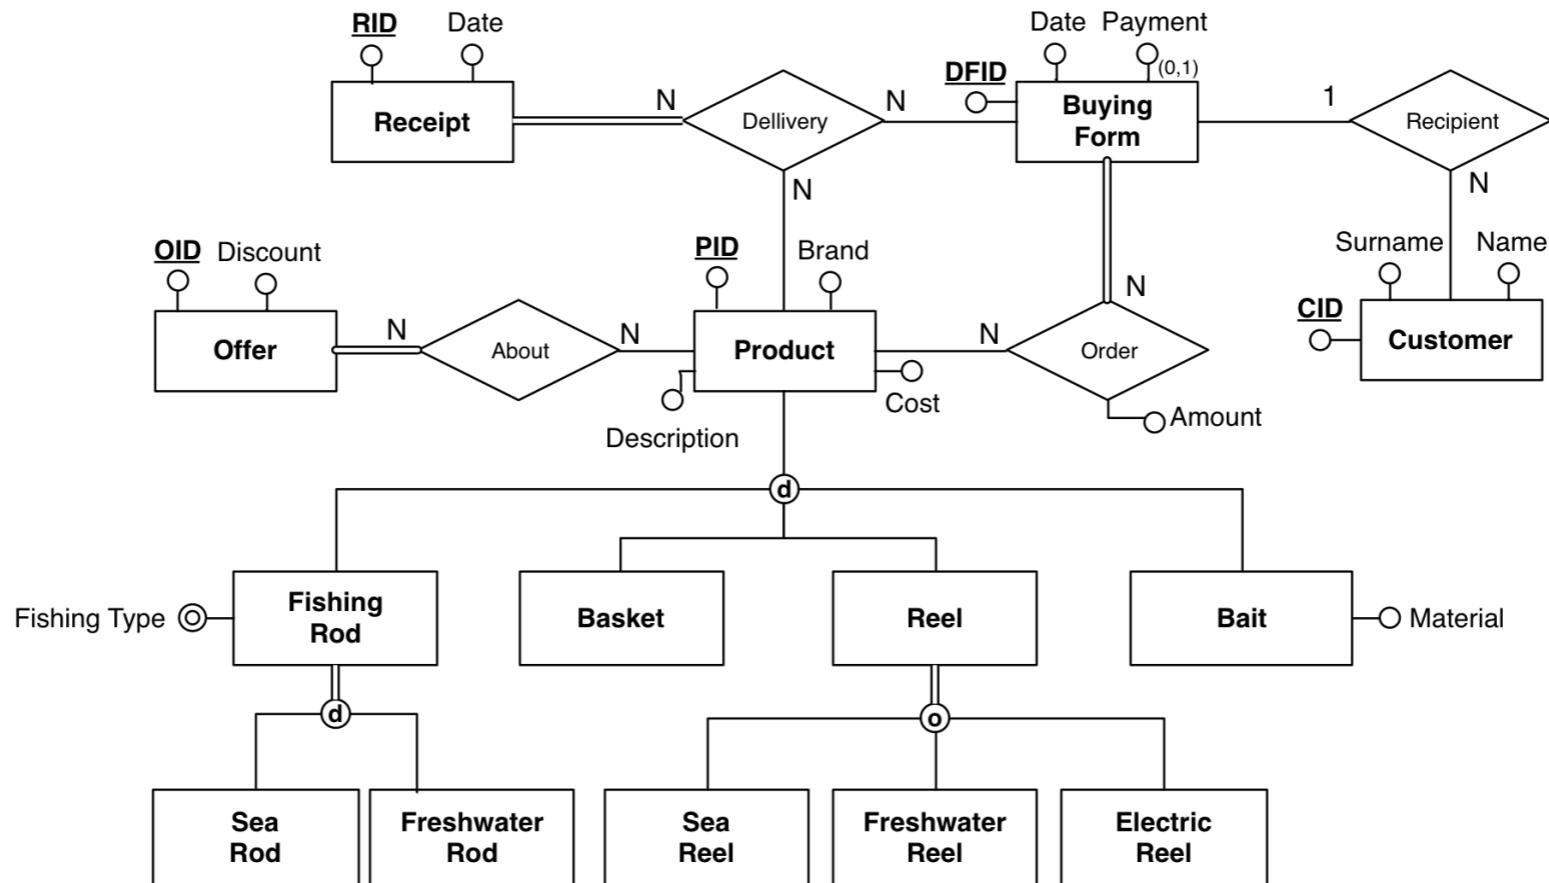
Product (Code , Name, Cost)

Composition (CodeComponent , CodeCompound , Quantity)

Consider the transformation of the EER diagram in figure.

Which of the following statements about the Order relationship is/are correct? (There can be more than one)

(1 min)



1

The resulting table will include 3 attributes

2

Given that the *Product* table has optional participation, the *PID* attribute not be part of the primary key of the resulting table

3

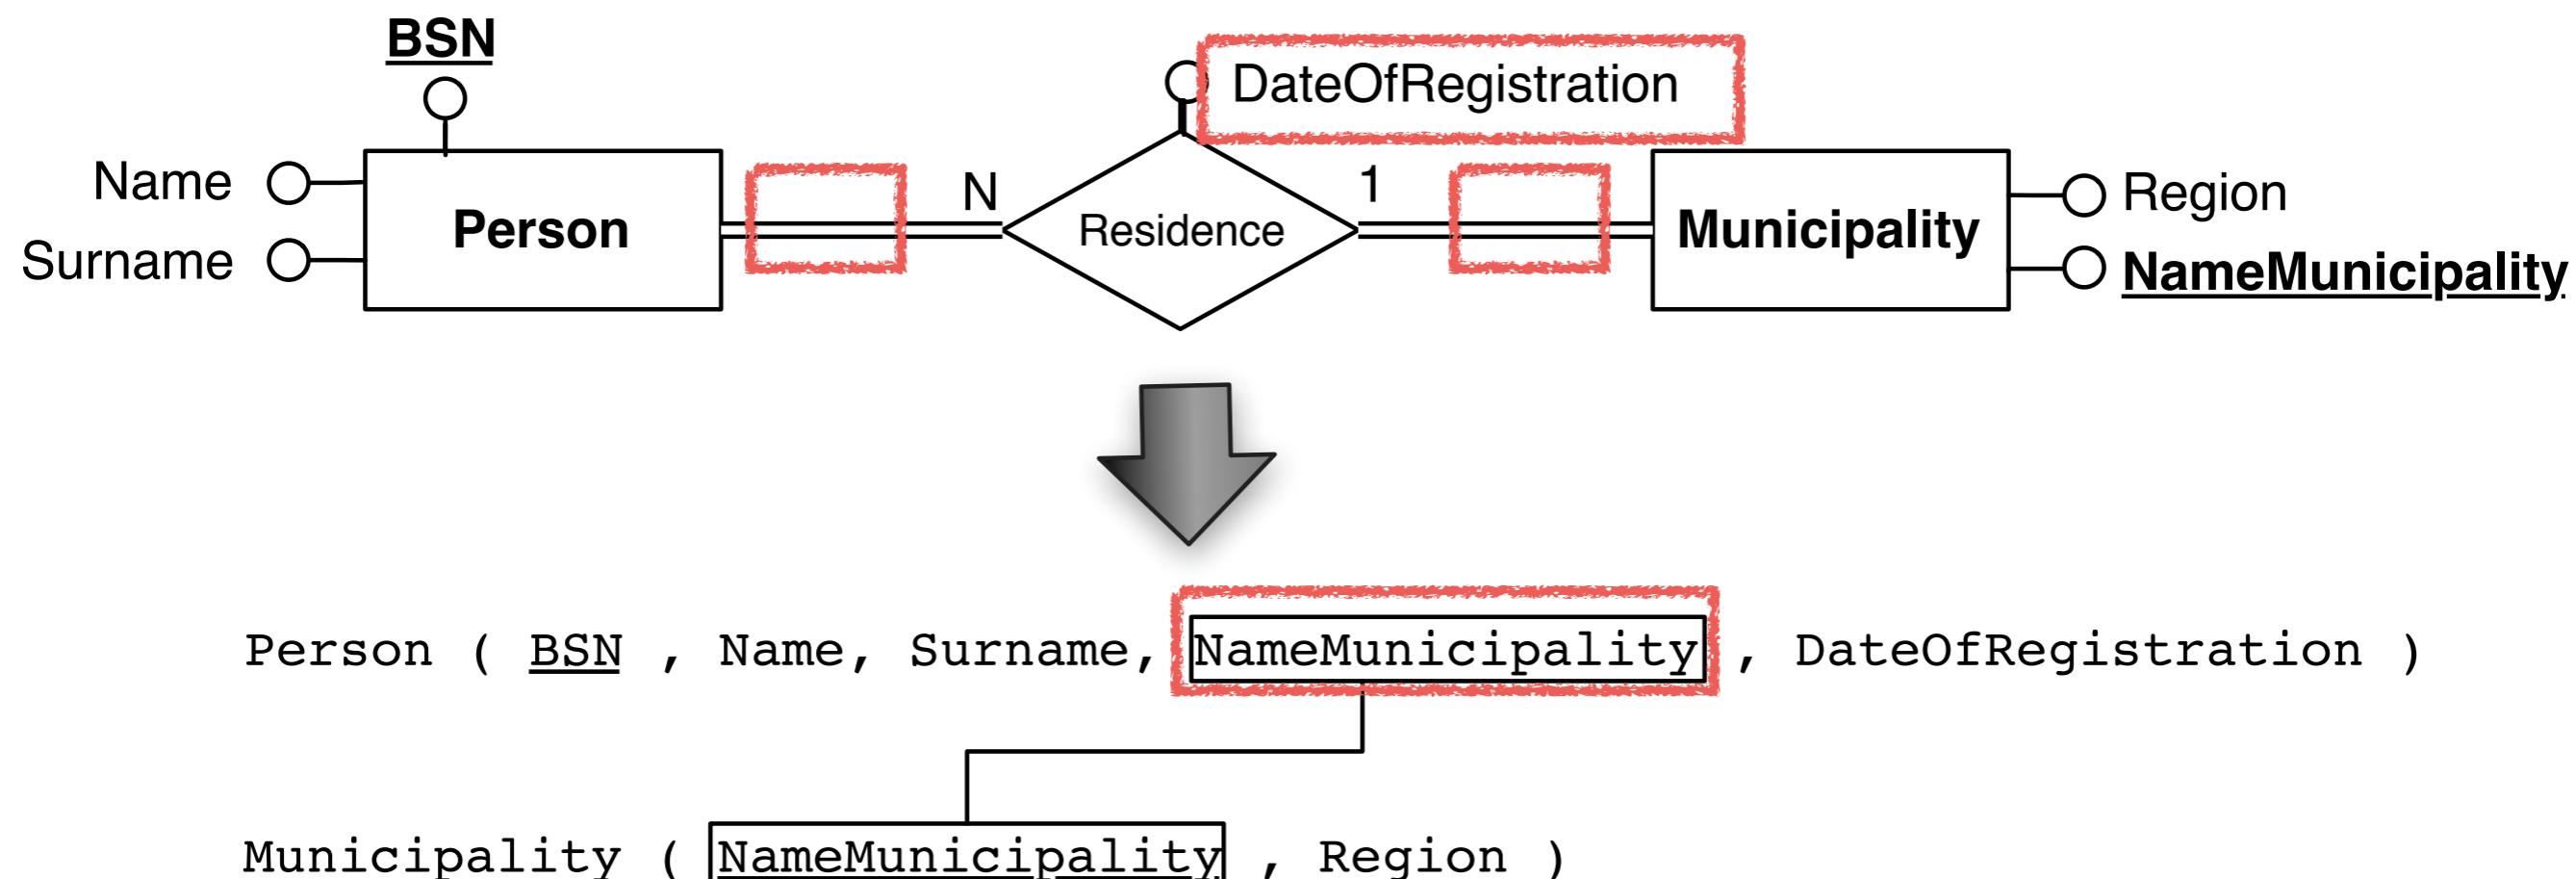
The *Amount* attribute will also be part of the primary key of the resulting table

4

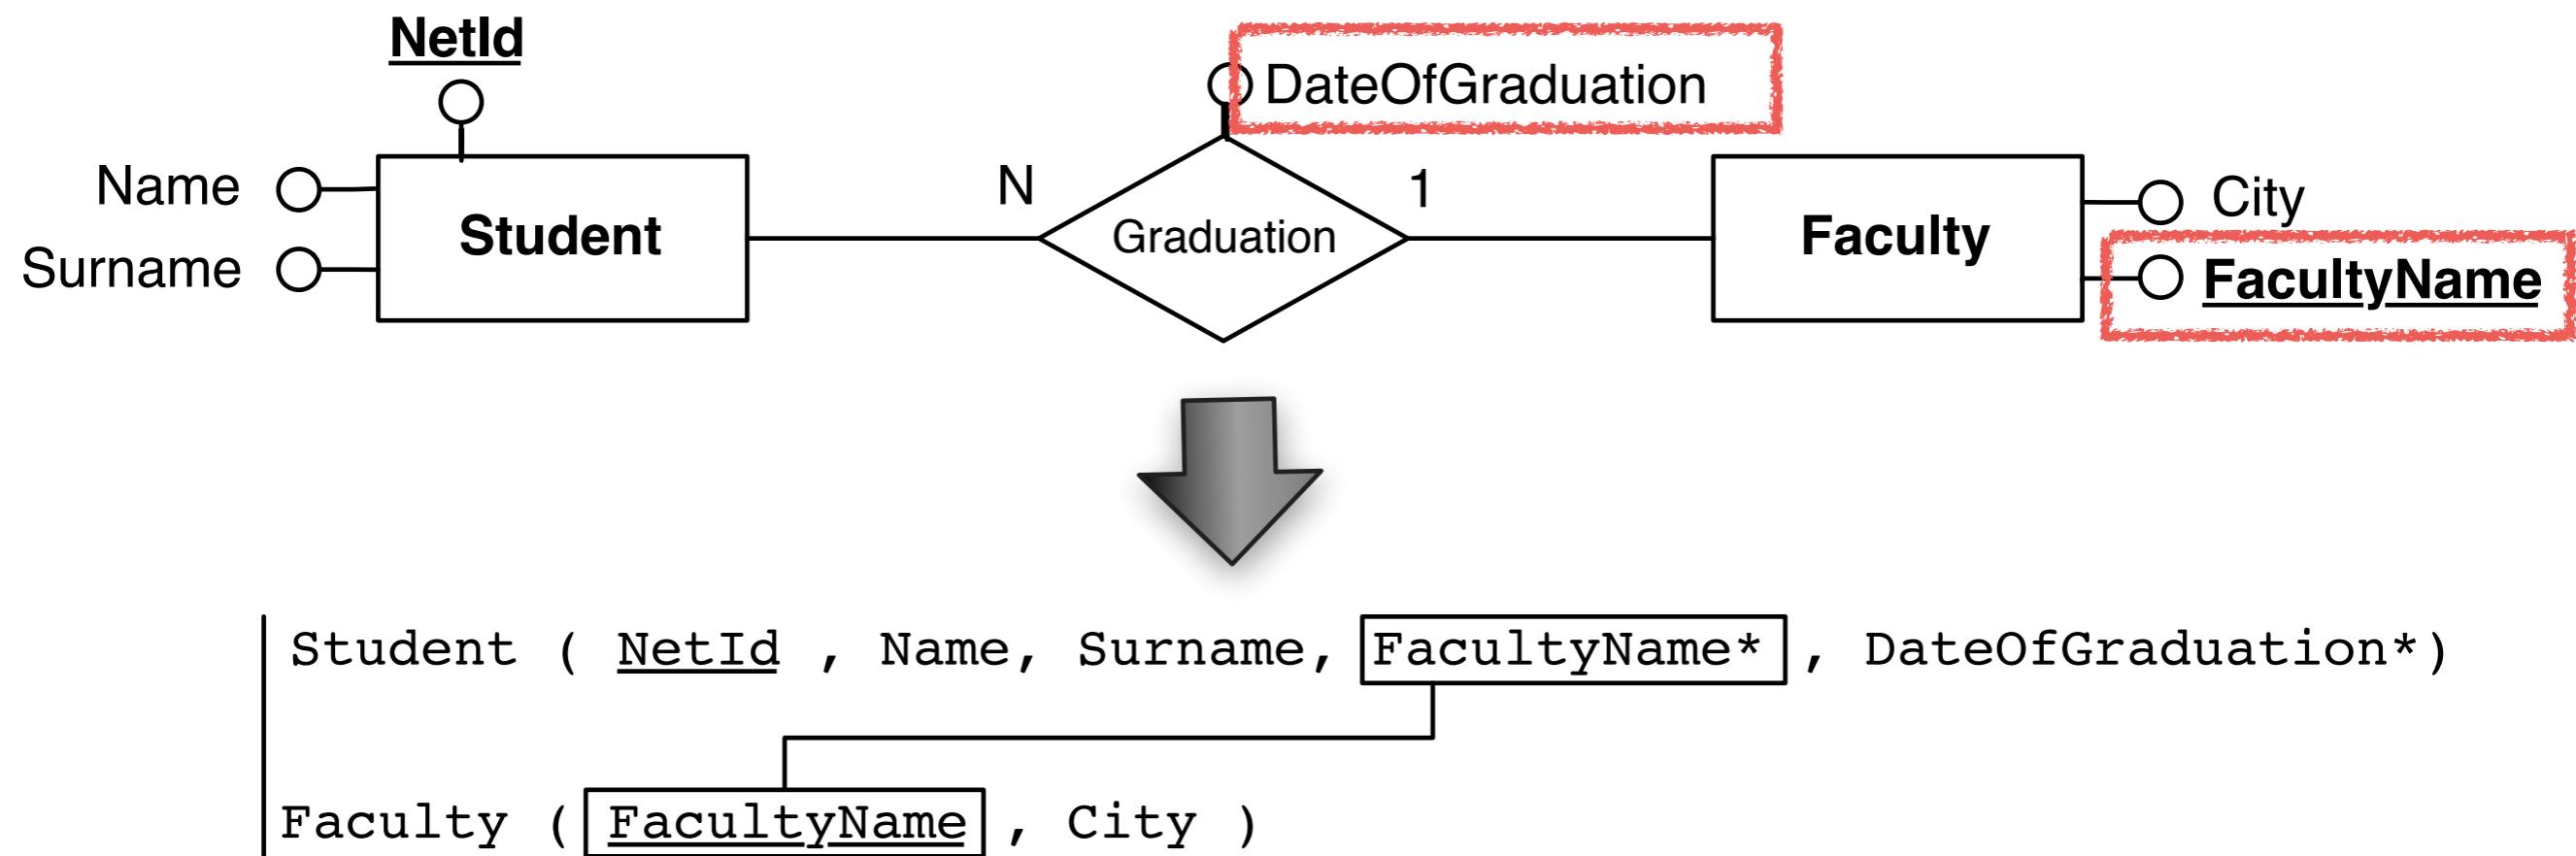
The resulting table will also include the *Date* and *Payment* attributes

Translation of One to Many binary relationship

- Two possible translations
 - Attributes (**foreign key**)
 - New table (**relationship relation**)



Example One To Many

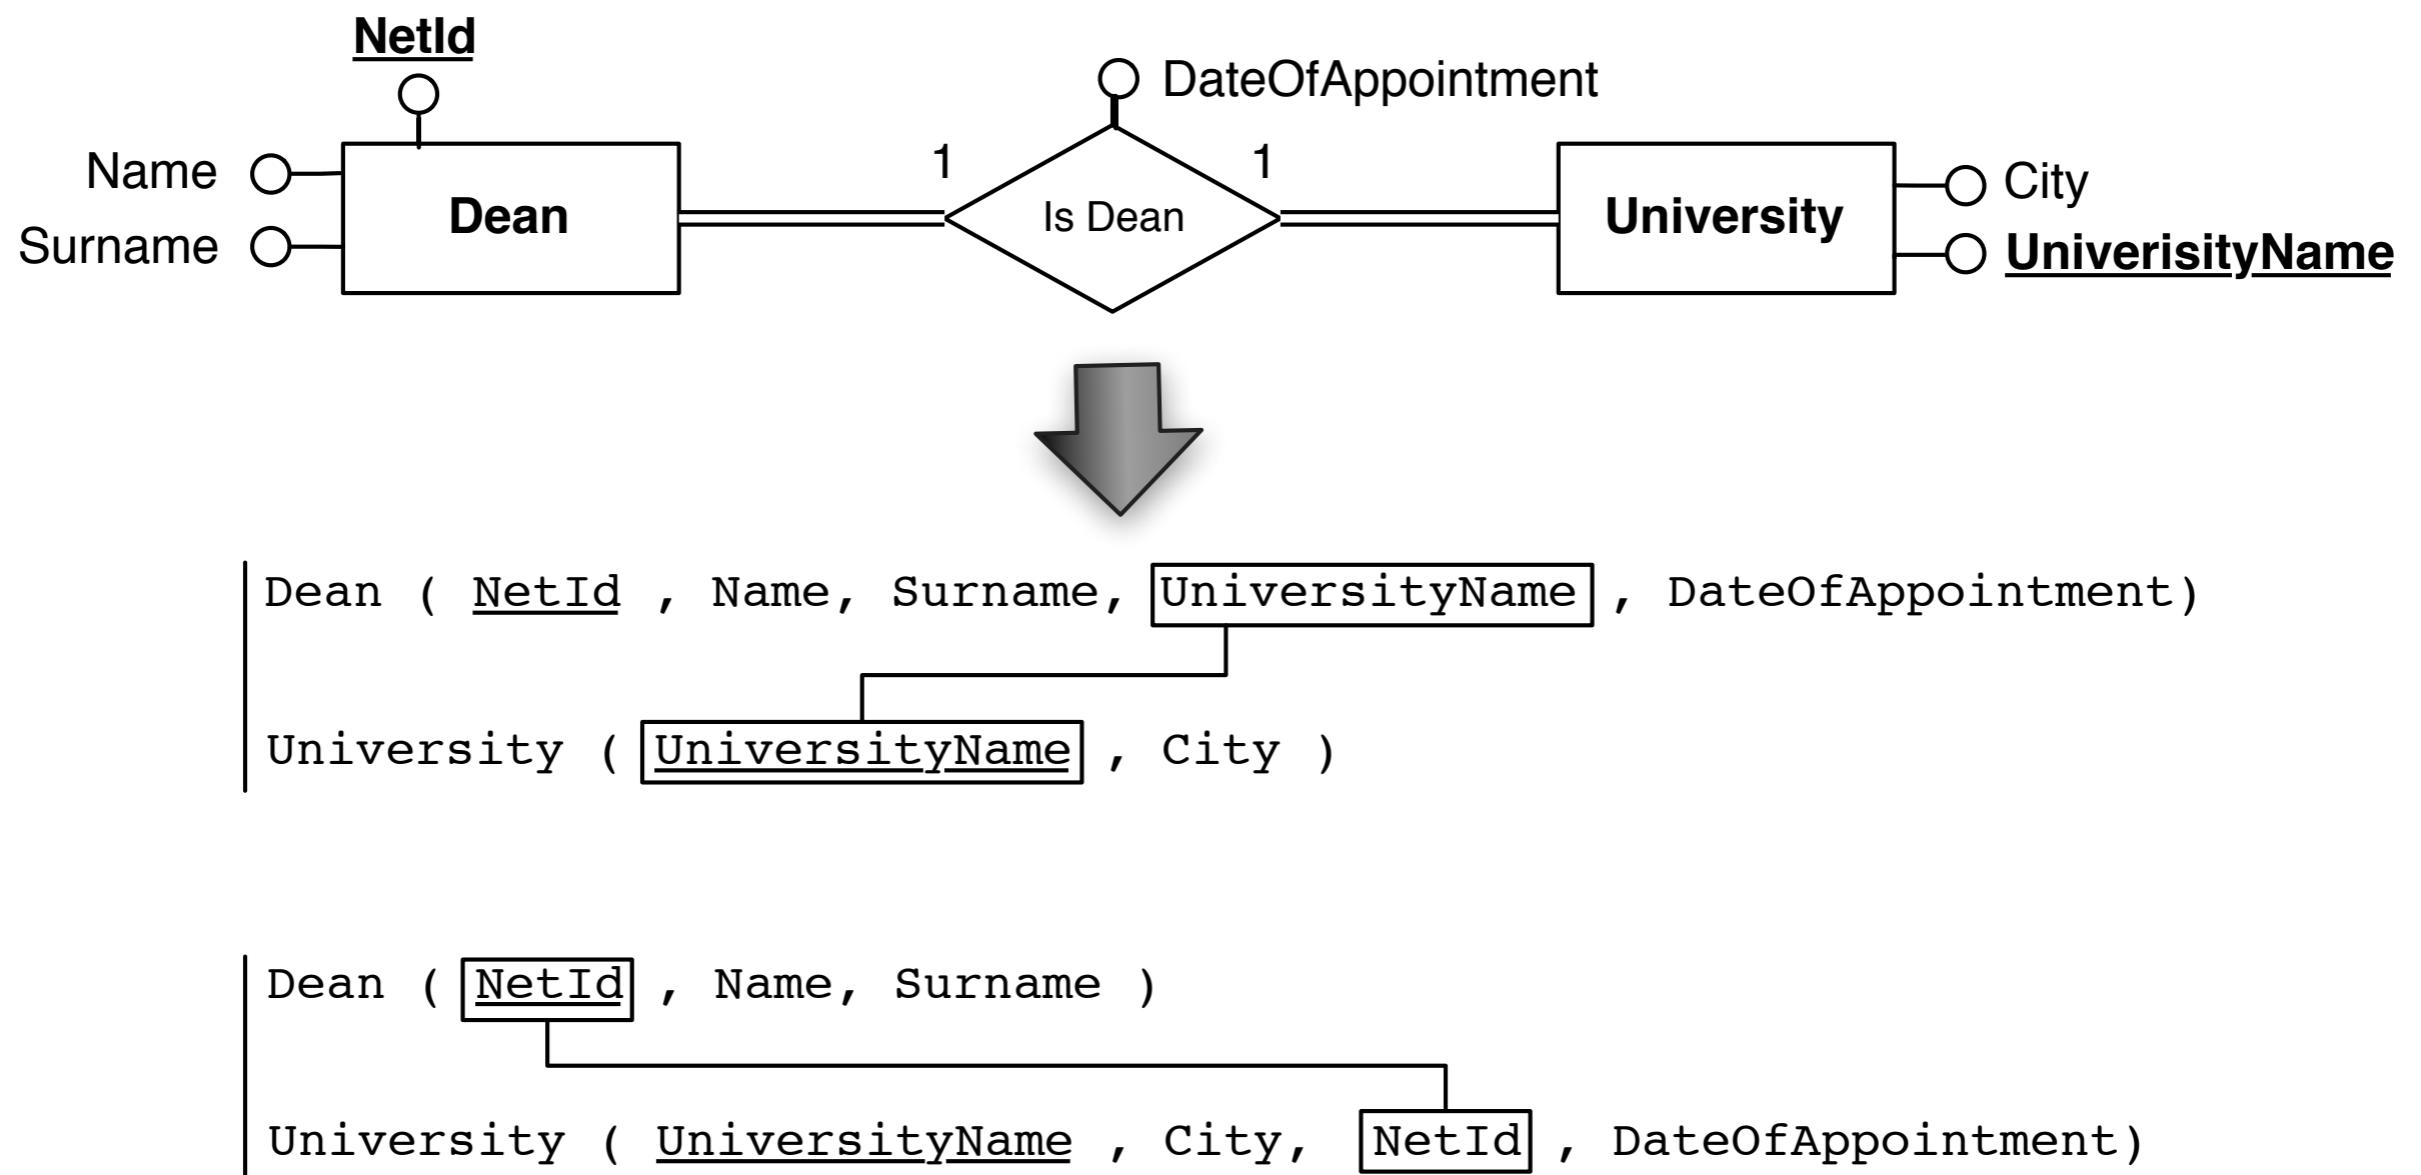


Translation of One to One binary relationship

- There can be several translations, according to the **minimum cardinality value**
 - Total participation on both roles
 - Total participation on one role
 - Optional participation on both roles

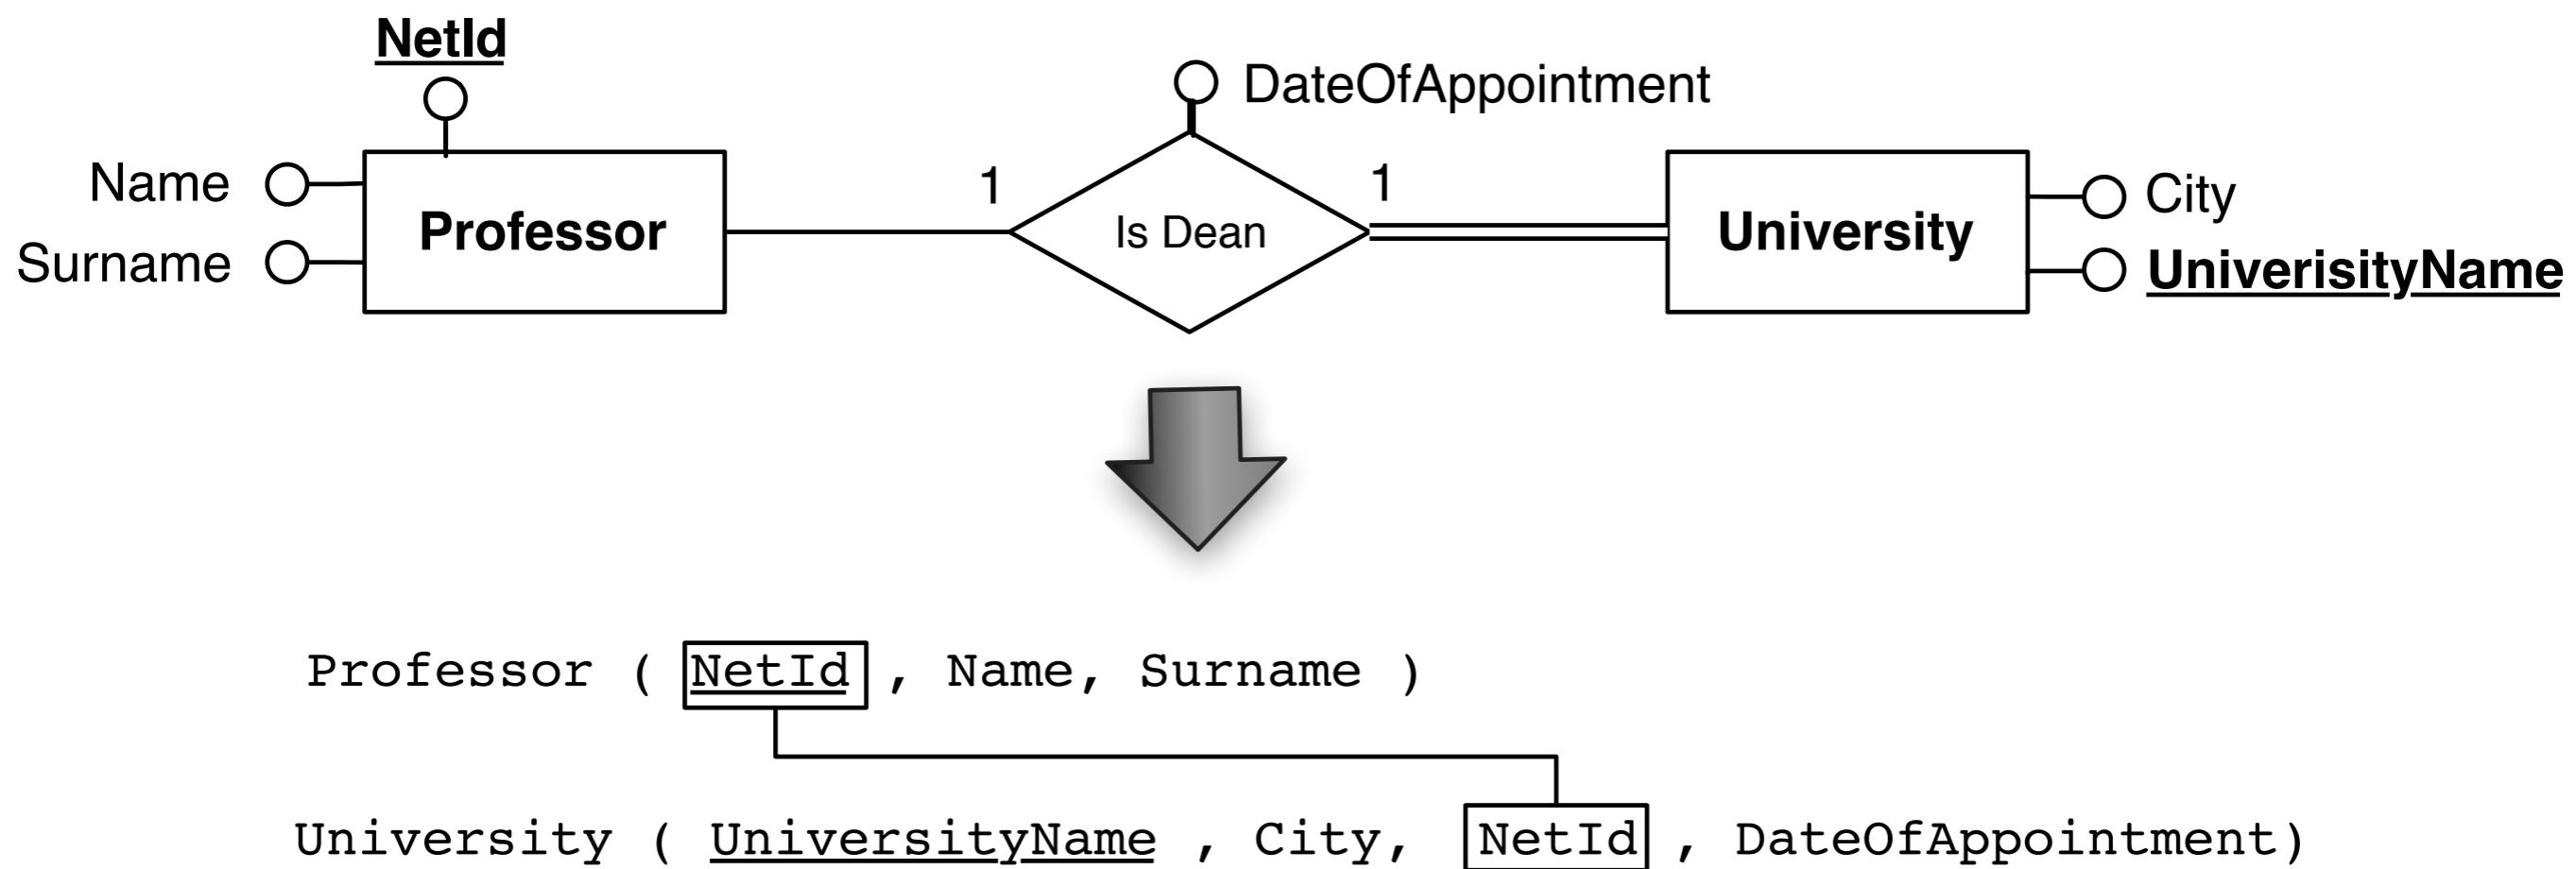
Example One To One /1

- Total participation on both roles: **Foreign key** approach
 - On one of the two sides

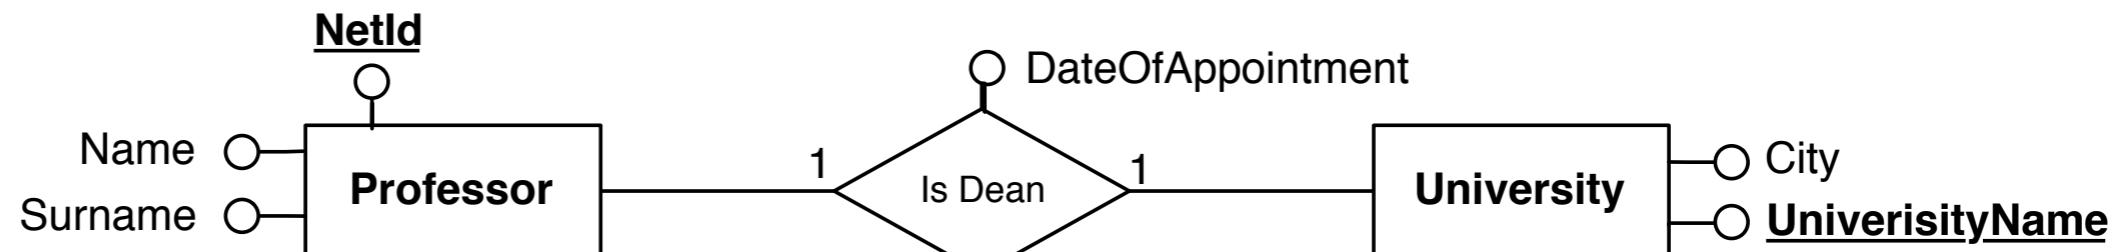


Example One To One /2

- Total participation on one role: **Foreign key** approach

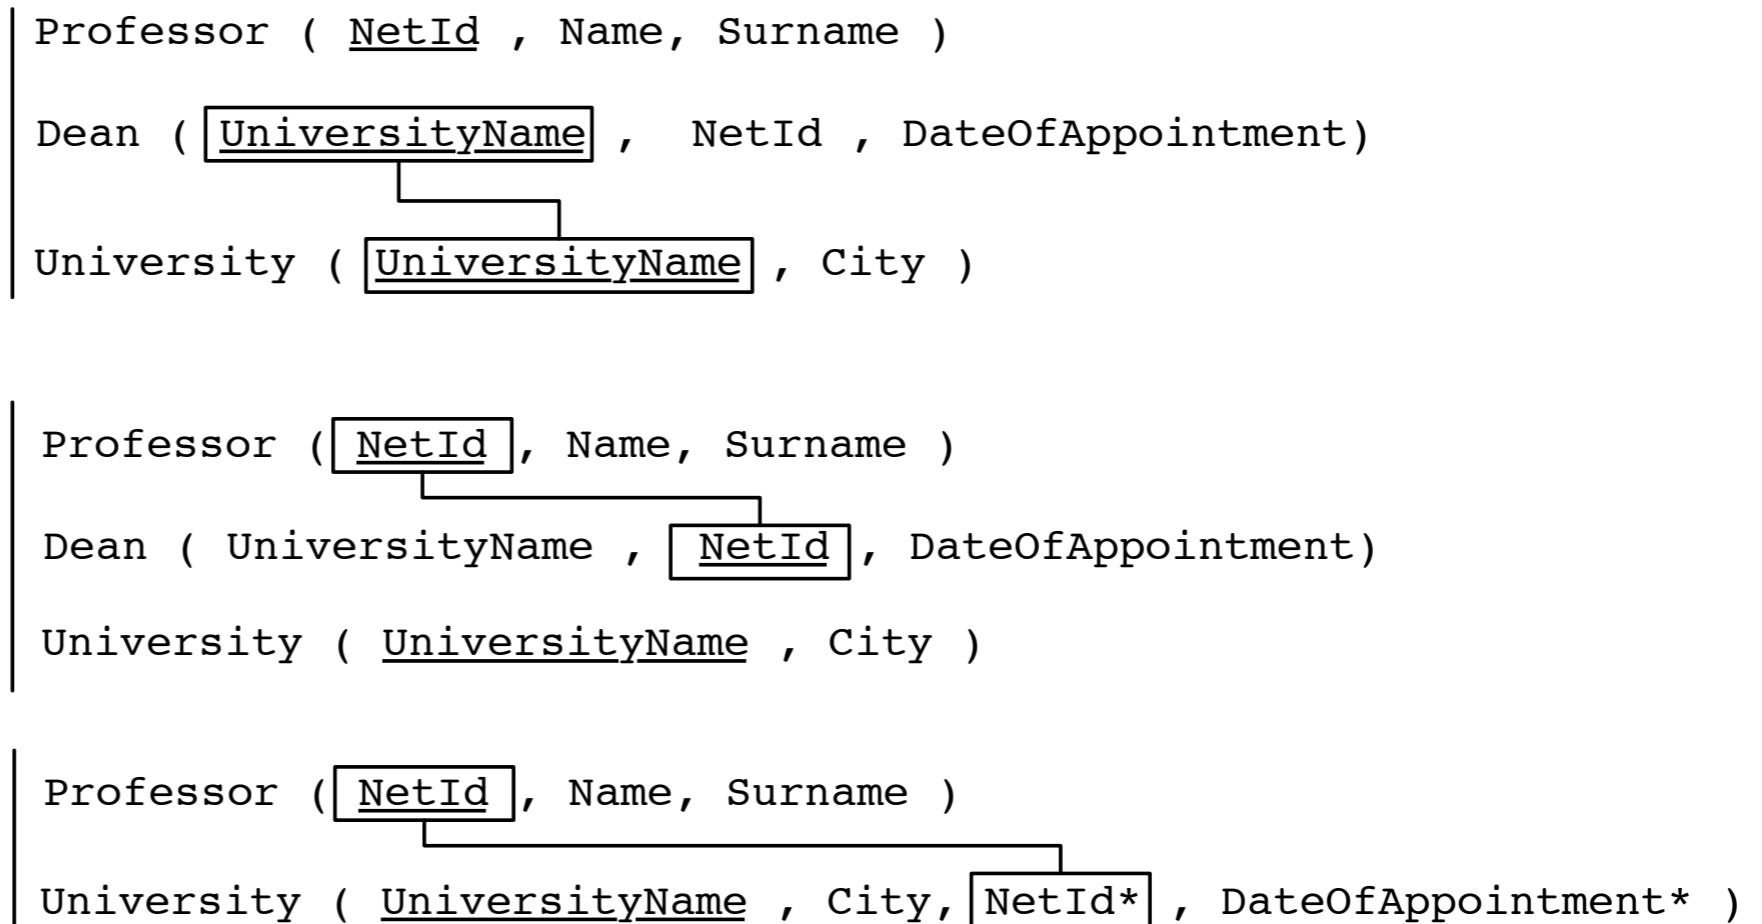


Example One To One /3



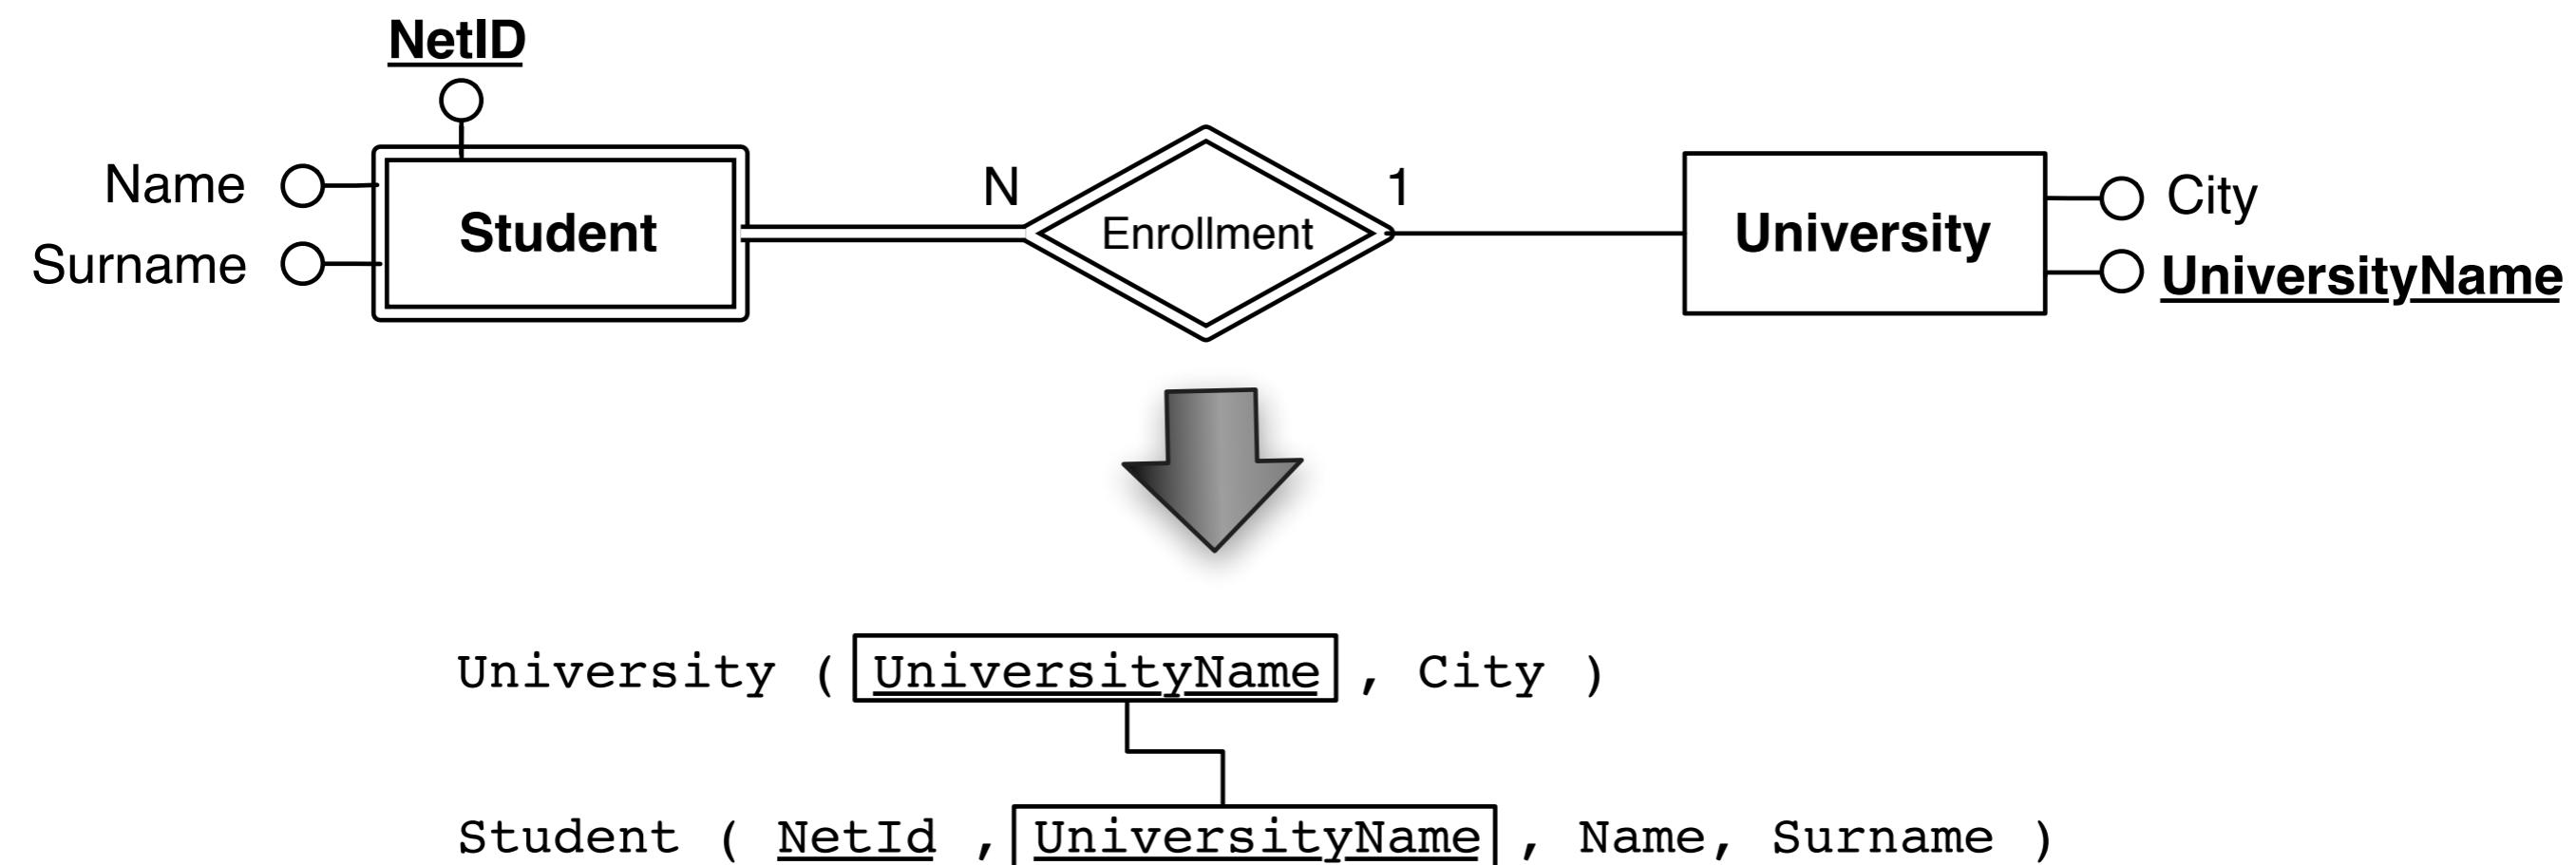
Optional participation on both roles

- **Foreign key** approach
- **Relationship relation** approach



Entity With External Identifier

- Relationship is represented with the identifier

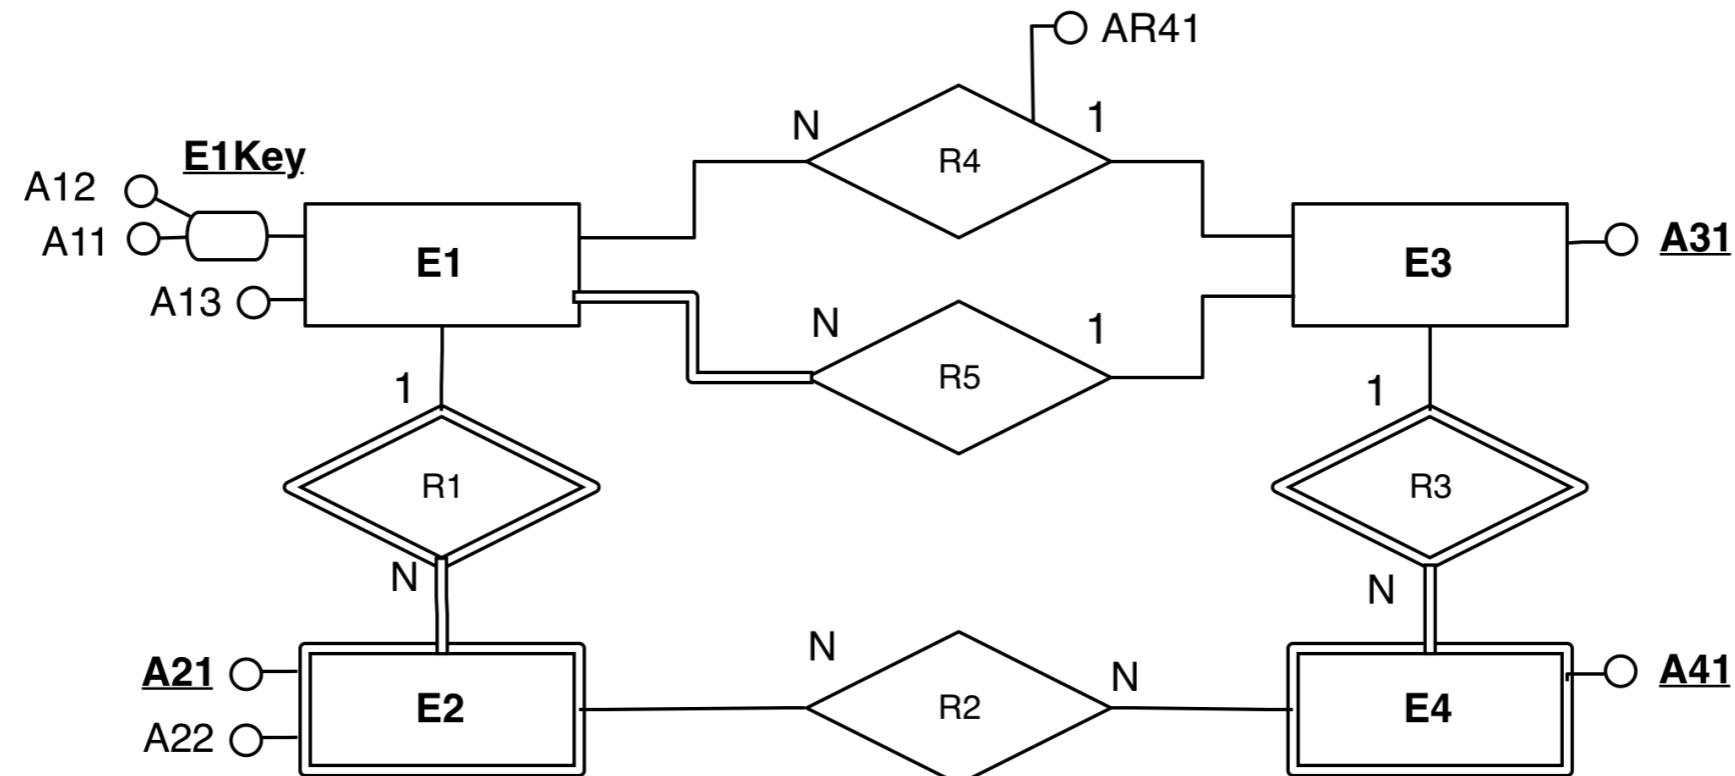


ASQ Exercise (1 min)

A S Q

Consider the EER diagram in figure.

Which is the minimum number of tables resulting from the transformation of the EER schema in Figure 1 into a logical schema?



1

5

2

6

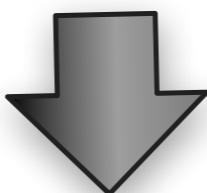
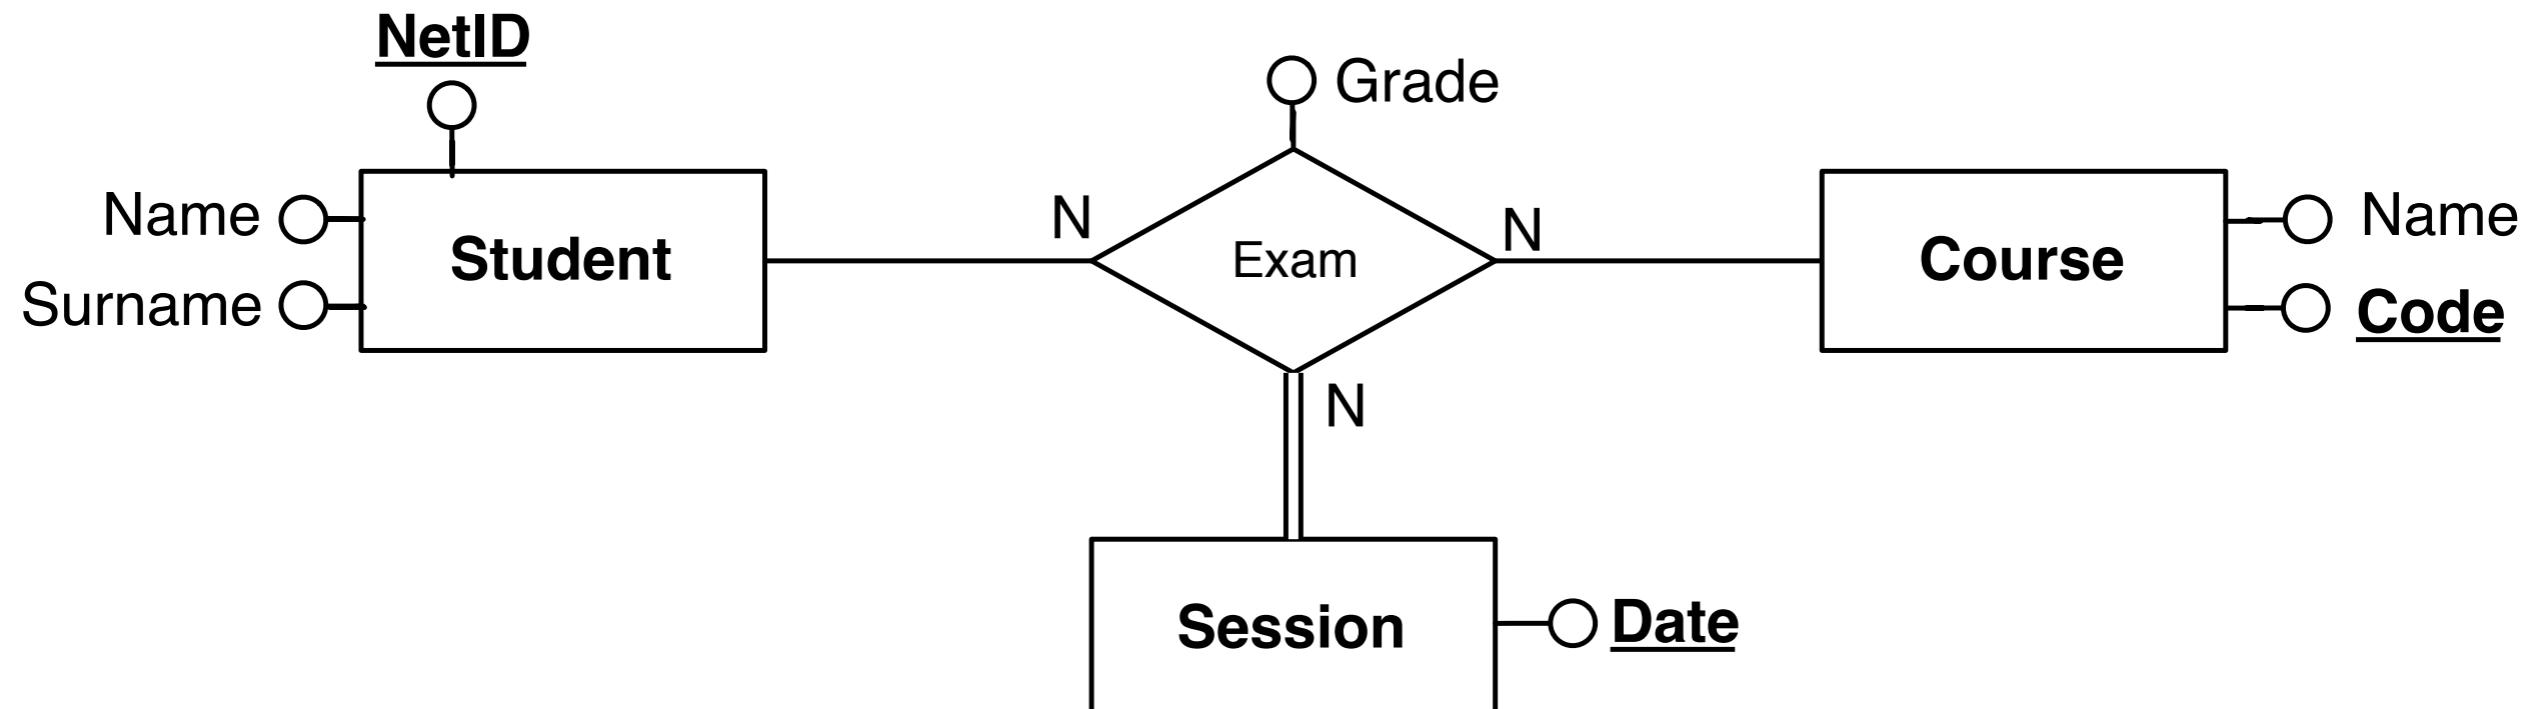
3

7

4

8

Ternary Relationships



Student (NetId , Name, Surname)

Exam (NetId , CourseName , Date , Grade)

Course (CourseName , Name)

Session (Date)

Readings

- **Database Book**
 - *Chapter 9 of the Database book*

End of Lecture